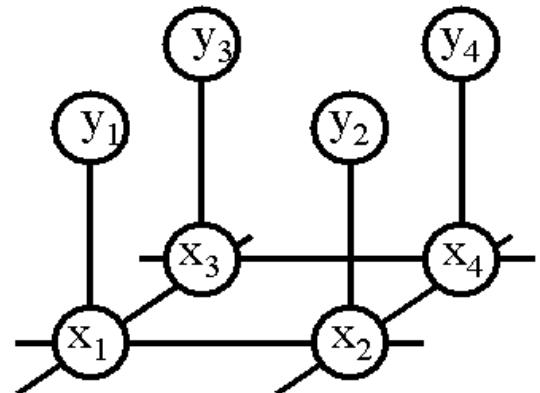
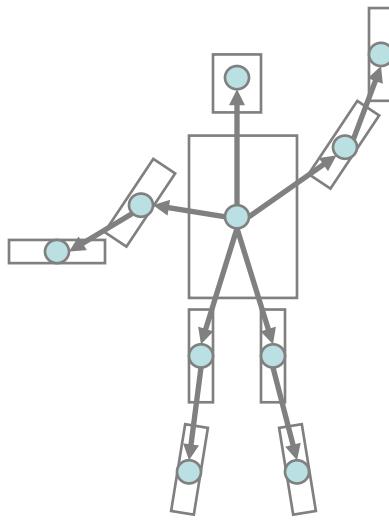
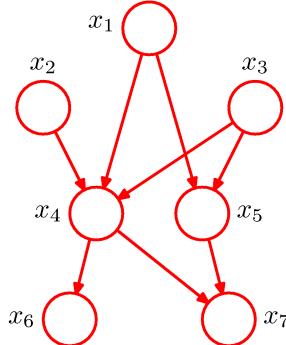


# Introduction to Supervised Learning



Lecture 10: Introduction to Structured Prediction

Iasonas Kokkinos

i.kokkinos@cs.ucl.ac.uk

University College London

# Tentative Course Schedule: generative & discriminative

- 1<sup>st</sup> week: Introduction, fundamental concepts
- 2<sup>nd</sup> week: Linear Regression
- 3<sup>rd</sup> week: Logistic Regression
- 4<sup>th</sup> week: Support Vector Machines
- 5<sup>th</sup> week: Ensemble Models (Adaboost, Random Forests)
- 6<sup>th</sup> week: Unsupervised learning (K-means, PCA, Sparse Coding)
- 7<sup>th</sup> week: Deep Learning (neural networks, backpropagation, SGD)
- 8<sup>th</sup> week: Probabilistic modelling (hidden variable models, EM)
- 9<sup>th</sup> week: Intro to Structured Prediction (Random Fields, Graphical Models)
- 10<sup>th</sup> week: review and applications



# Lecture outline

Introduction

Graphs and computation

Graphical models

Chain-structured graphical models

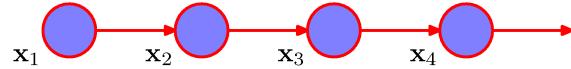
Tree-structured graphical models

Loopy graphical models

# Problem 1: measuring the length of a queue



**while only complaining to your closest neighbors**



## Problem 2: two queues



# Problem N: N queues



# Lecture outline



Introduction

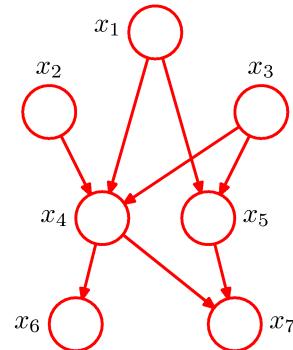
Computation on graphs

Graphical models

Chain-structured graphical models

Tree-structured graphical models

Loopy graphical models



# Graphical models

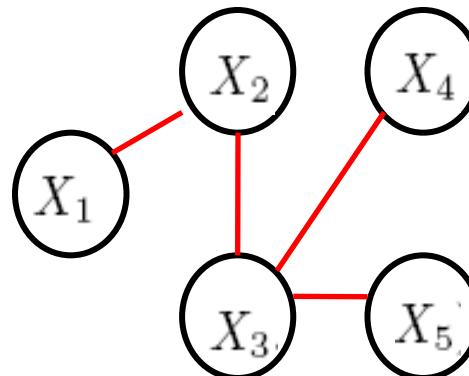
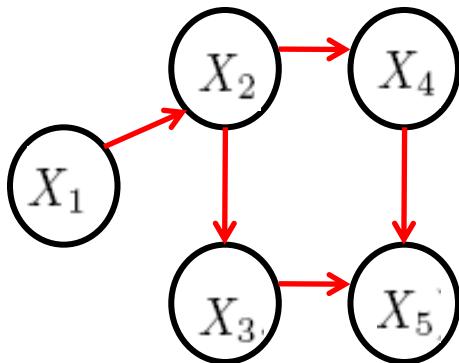
Blend of graph theory and probability

Graph nodes: random variables

Graph edges: relationships

Directed graphs: Bayesian Networks

Undirected graphs: Markov Random Fields



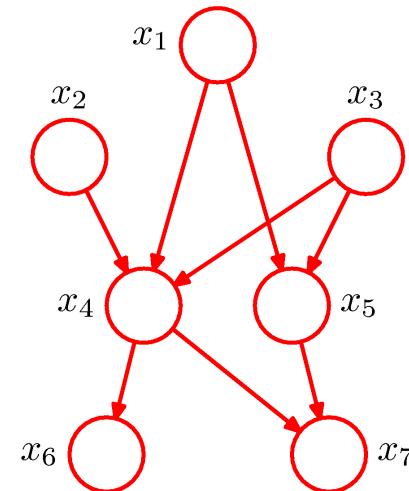
# Bayesian Networks

Directed Acyclic Graph  $G(V, E)$

Nodes -  $V$  : random variables

Edges -  $E$  : dependencies

Leave from 'parents', arrive at children



Distribution of each child conditioned on parent:  $P(X_v | X_{\pi_v})$

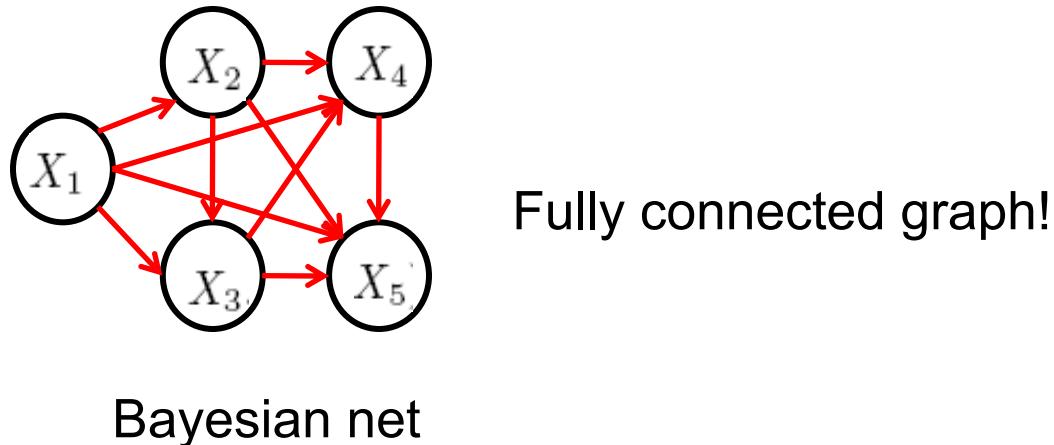
$$P(X_v | \emptyset) = P(X_v)$$

Joint probability distribution:  $P(X) = \prod_v P(X_v | X_{\pi_v})$

# Conditional independencies

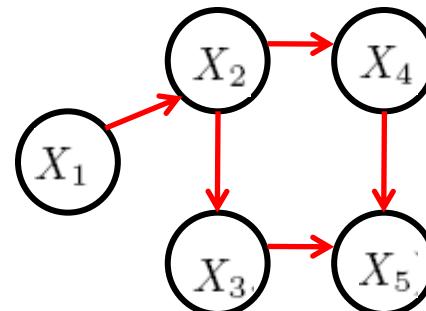
We can trivially obtain a bayesian network for a probability distribution

$$P(X_1, X_2, X_3, X_4, X_5) = P(X_1)P(X_2|X_1)P(X_3|X_2, X_1)P(X_4|X_3, X_2, X_1)P(X_5|X_4, X_3, X_2, X_1)$$



Non-trivial cases: some conditional independence

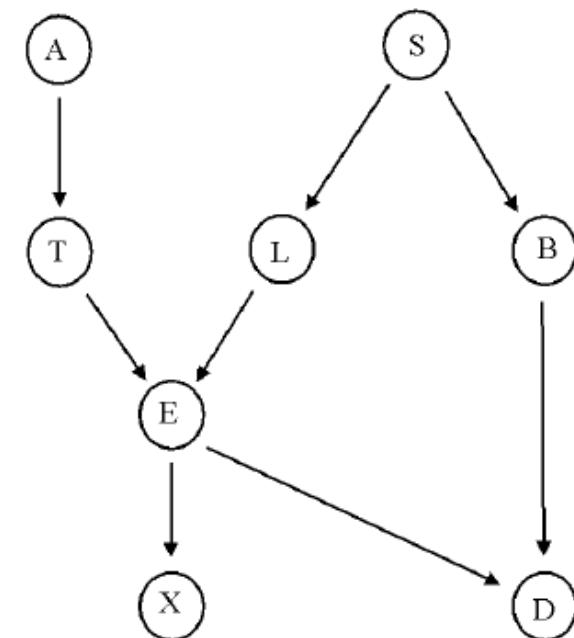
$$\begin{aligned} P(X_3|X_2, X_1) &= P(X_3|X_2) \\ P(X_4|X_3, X_2, X_1) &= P(X_4|X_2) \\ P(X_5|X_1, X_2, X_3, X_4, X_5) &= P(X_5|X_3, X_4) \end{aligned}$$



Main interest: exploit independencies for training & inference

# A Bayesian Network-based Expert System

- ‘Asia Network’
  - A: trip to Asia
  - T: tuberculosis
  - S: smoking
  - L: Lung Cancer
  - B: Bronchitis
  - E: Tuberculosis/Lung Cancer
  - X-ray results
  - D: Dyspnea



- Given: X-rays, Dyspnea, patient went to Asia, patient smokes
- Wanted: posterior probability of Bronchitis

# Detection vs. Tracking



**t=1**



**t=2**

...



**t=20**



**t=21**

# Detection vs. Tracking



- $t=1$   $t=2$  Detection: each frame independently

$t=20$

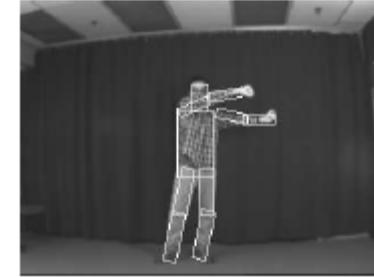
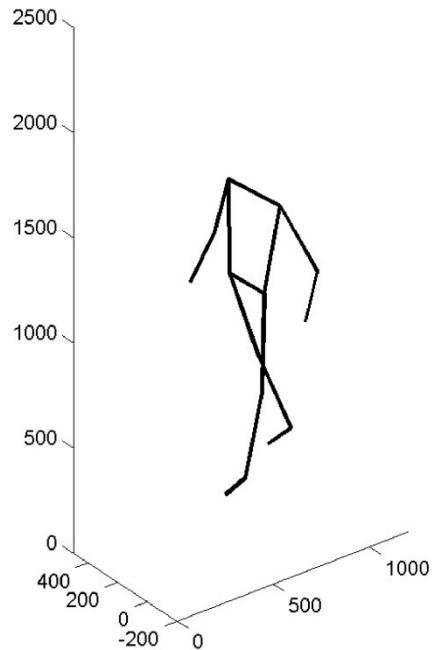
$t=21$

# Detection vs. Tracking



- $t=1$   $t=2$   $t=20$   $t=21$   
Tracking with *dynamics*: combine image measurements with our expectation of the object's motion.
  - Restrict search for the object
  - Measurement noise is reduced by trajectory smoothness.

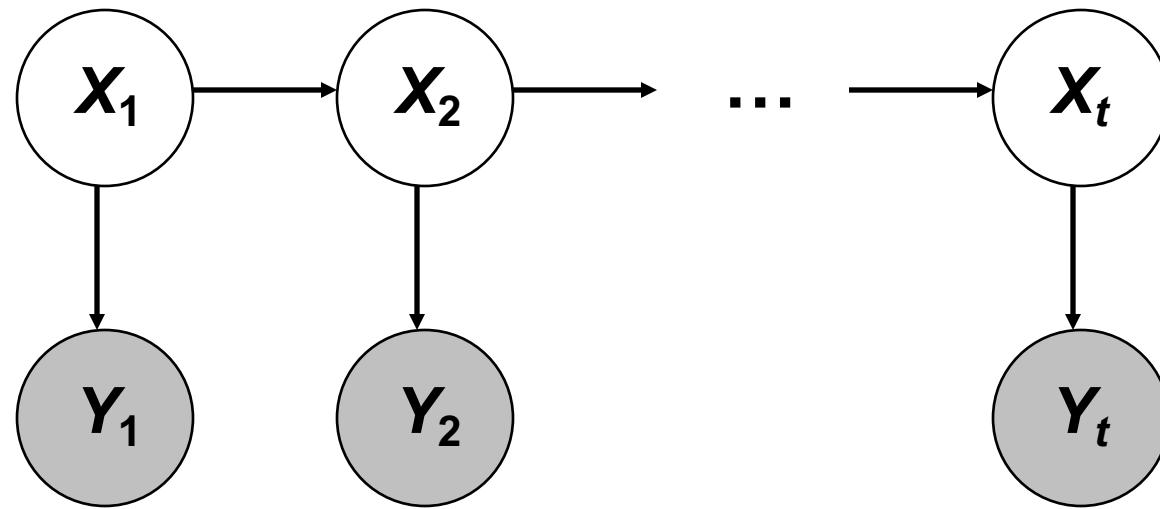
# 'State space'



- X: System state (parameters of interest)
- Y: Measurement (what we observe)

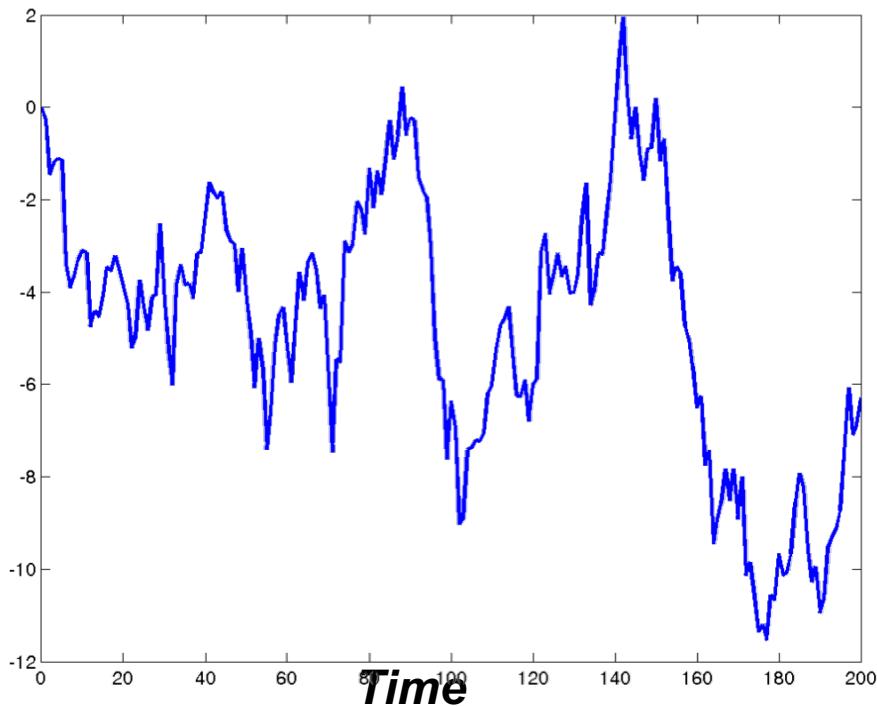
# General Model for Tracking

- The moving object of interest is characterized by an underlying state  $X$
- State  $X$  gives rise to *measurements* or *observations*  $Y$
- At each time  $t$ , the state changes to  $X_t$  and we get a new observation  $Y_t$



# Simple Linear Dynamics

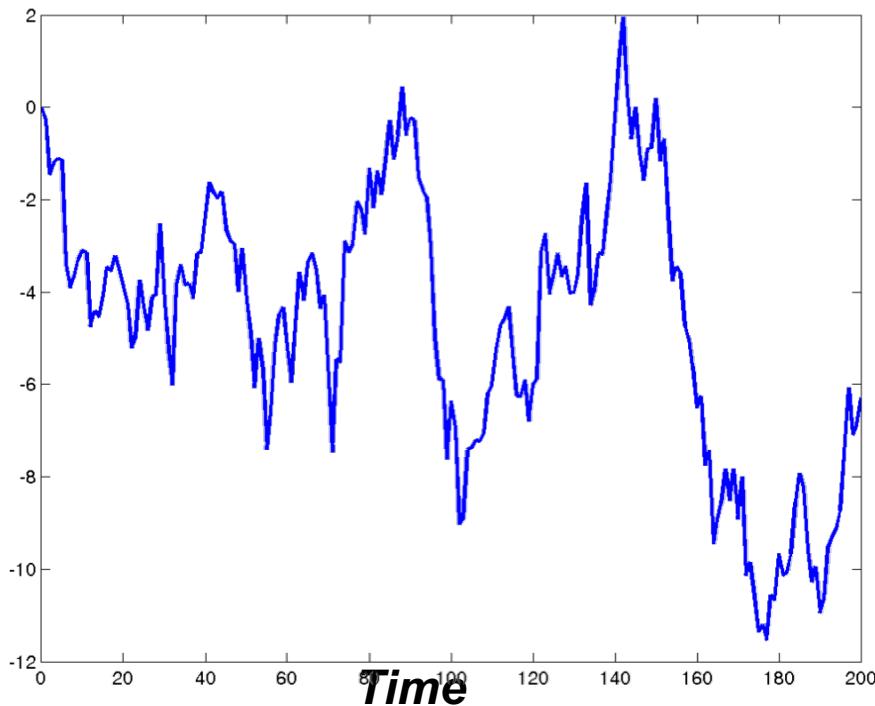
*Brownian Motion*



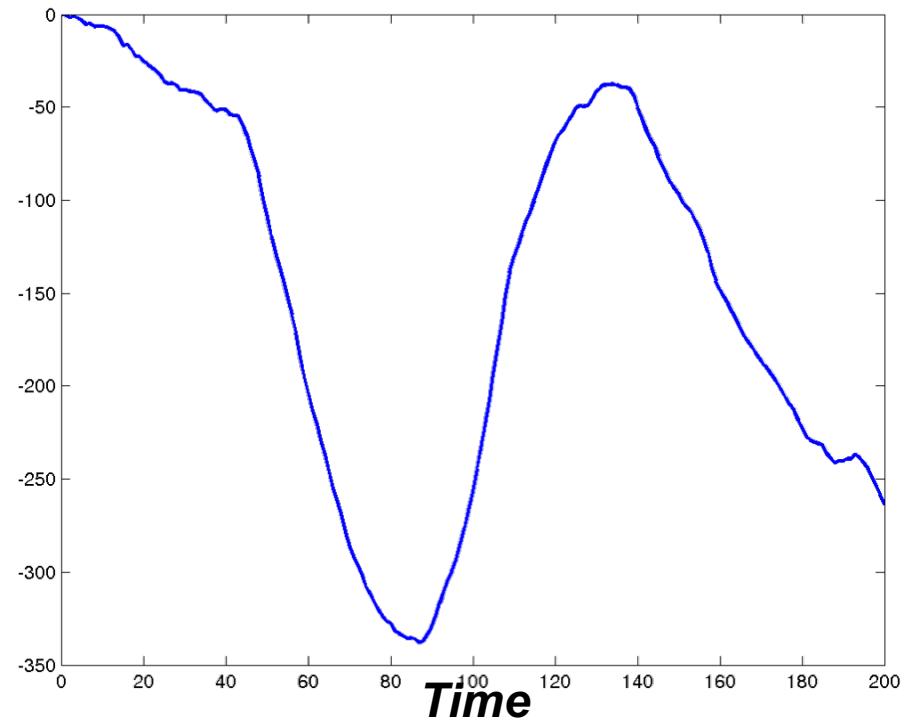
$$x_{t+1} = x_t + w_t$$

# Simple Linear Dynamics

*Brownian Motion*



*Constant Velocity*



$$x_{t+1} = x_t + w_t$$

$$\begin{bmatrix} x_{t+1} \\ \delta_{t+1} \end{bmatrix} = \begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} x_t \\ \delta_t \end{bmatrix} + w_t$$

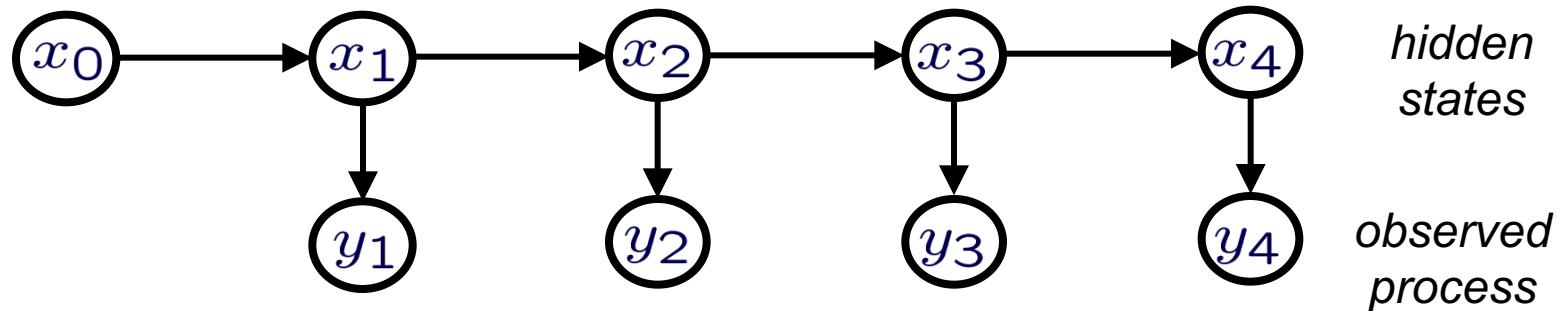
# Linear Dynamical Systems

- Noise in State, noise in Observations

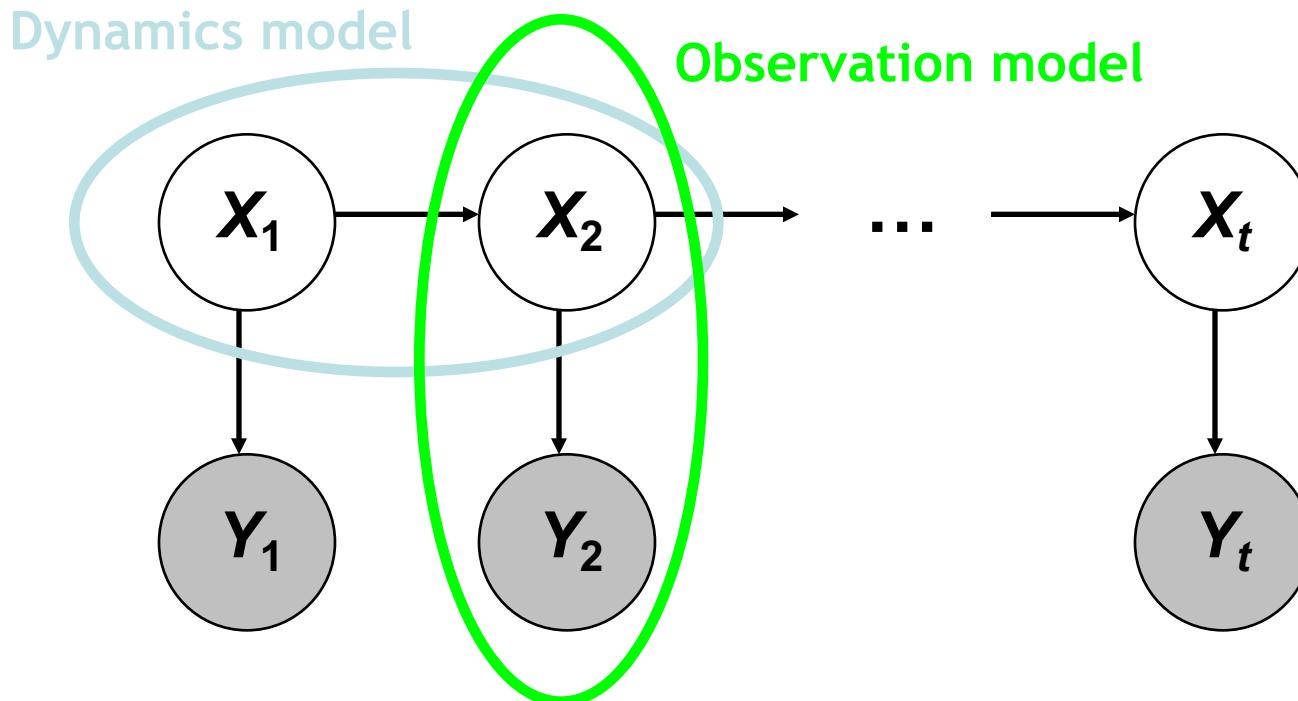
$$x_t = Ax_{t-1} + v_t, \quad v_t \sim N(0, \Sigma)$$

$$y_t = Cx_t + w_t, \quad w_t \sim N(0, Q)$$

$$x_1 = \mu_0 + u, \quad u \sim N(0, V)$$



# Markov Assumptions



$$P(X_t | X_{t-1}, X_{t-2}, \dots, X_1) = P(X_t | X_{t-1})$$

$$P(Y_t | X_1, \dots, X_T, Y_1, \dots, Y_{t-1}, Y_{t+1}, \dots, Y_T) = P(Y_t | X_t)$$

--

# Tracking as Induction

- Base case:
  - Prior prediction of state (in absence of any evidence):  $P(X_0)$
  - At the first frame, *correct* this given the value of  $Y_0=y_0$

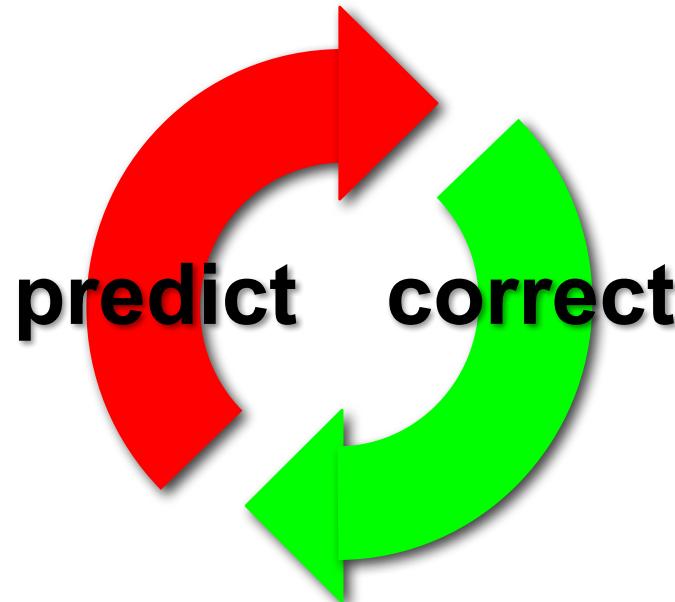
$$P(X_0 | Y_0 = y_0) = \frac{P(y_0 | X_0)P(X_0)}{P(y_0)} \propto P(y_0 | X_0)P(X_0)$$

**Posterior prob.  
of state given  
measurement**

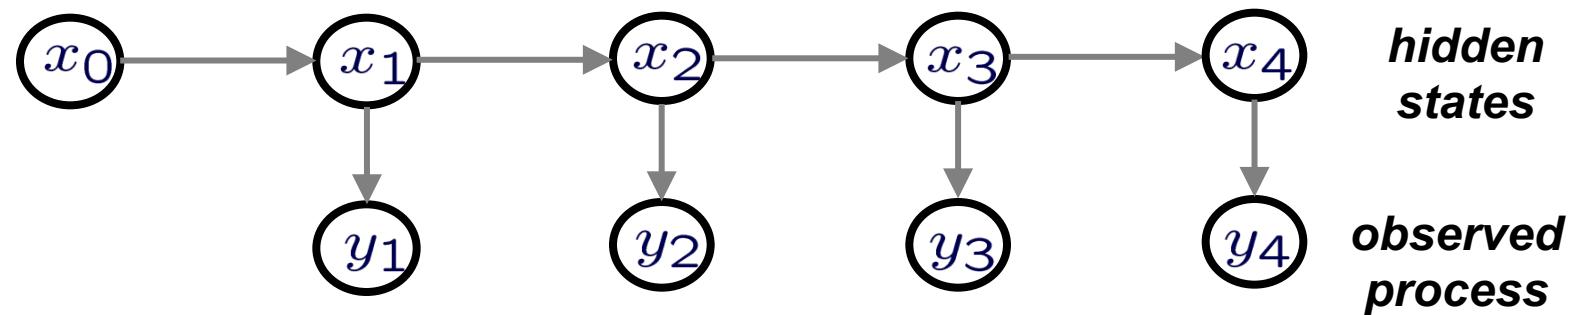
**Likelihood of measurement      Prior of the state**

# Tracking as Induction

- Base case:
  - Assume we have initial prior that predicts state in absence of any evidence:  $P(X_0)$
  - At the first frame, *correct* this given the value of  $Y_0=y_0$
- Given corrected estimate for frame  $t$ :
  - Predict for frame  $t+1$
  - Correct for frame  $t+1$



# Hidden Markov Models



- Dependent sequence of observations
- Conditioned on state sequence, observations become independent

Parameters

Observation probabilities  $P(y_t|x_t)$

$$\pi(y_t = k|x_t = m) = n_{k,m} \quad P(y_t|x_t = k) = N(y_t; \mu_k, \Sigma_k)$$

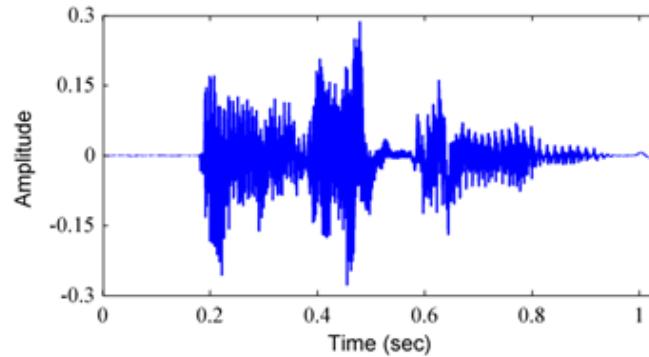
discrete    continuous

Transition probabilities  $\pi(x_{t+1} = l|x_t = m) = \alpha_{l,m}$

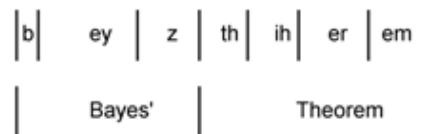
Starting probabilities  $\pi(x_0 = m) = \beta_m$

# HMMs & Speech Recognition

- Input: Air pressure signal

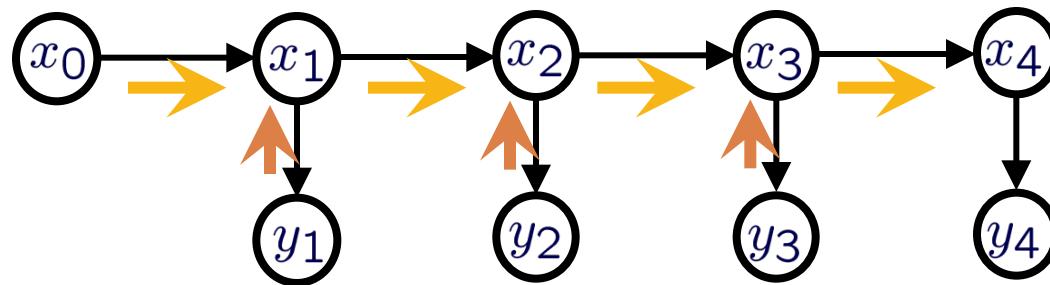


- Output: words/phrases
- Probabilistic model: signal given word
- Hidden Variables:
  - phonemes (/ae/, /ey/, /z/...)
  - phone subunits
- Hidden variables: not independent
  - Show, shoe, she, sell, sea,
  - Sv?
  - (Sven, svelte,...)
  - Sb?

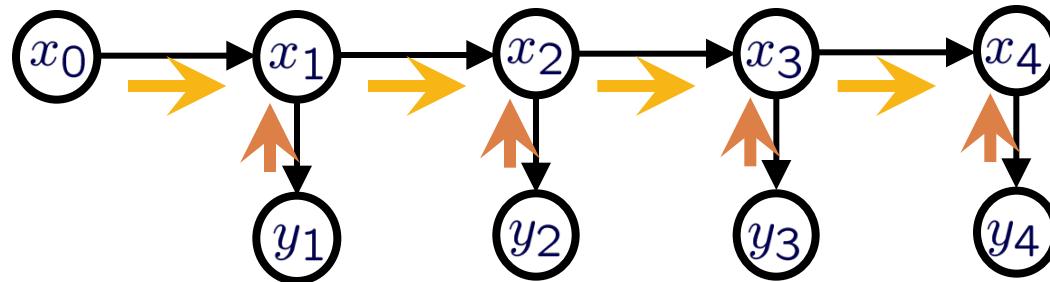


# Inference Tasks

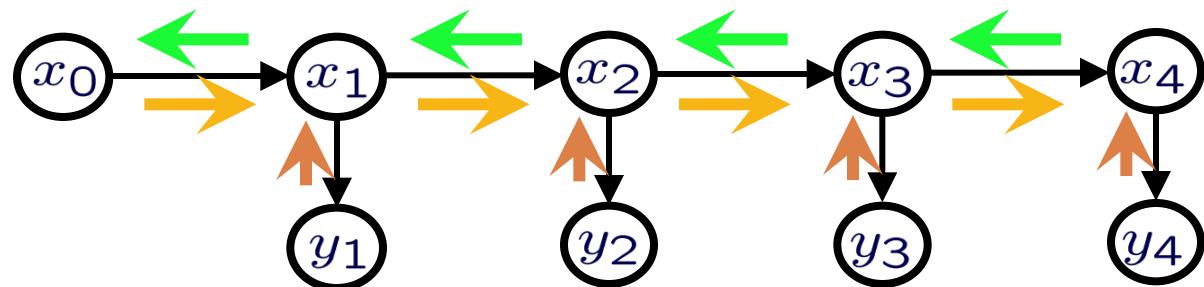
- Prediction



- Filtering  
(Kalman Filter)



- Smoothing



# Lecture outline



Introduction

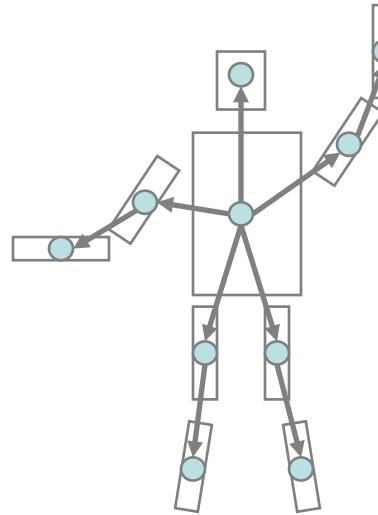
Chain-structured graphical models

Tree-structured graphical models

Pictorial Structures

Max/Sum-Product algorithm

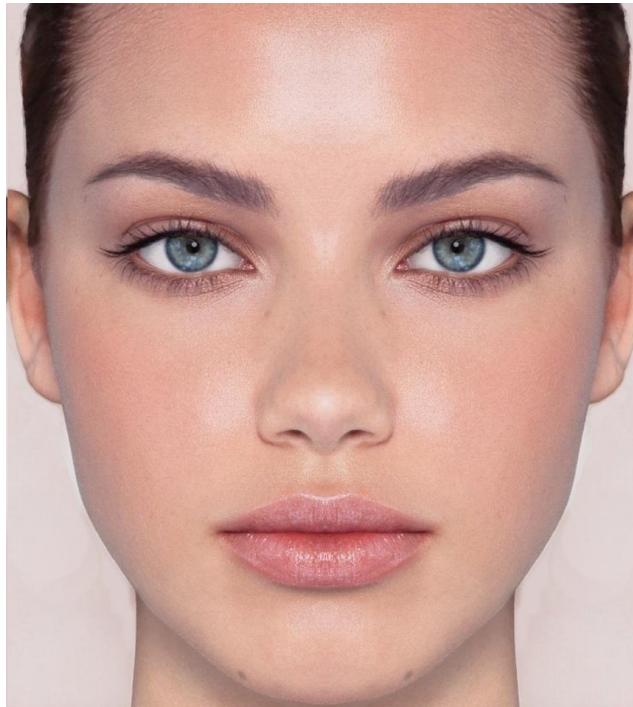
Loopy Graphical Models



# Problem N: N queues

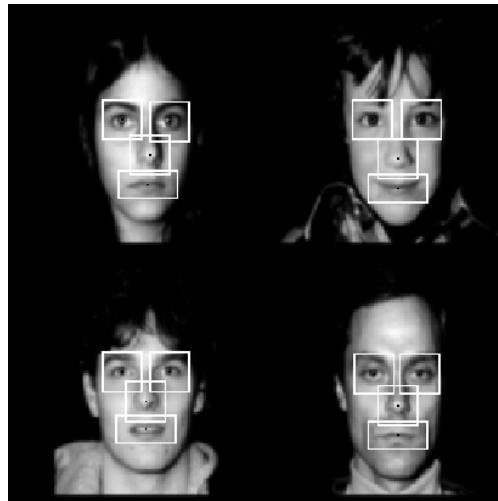


# Part-based Models

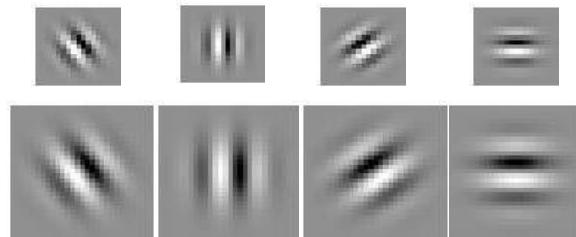


# Unary terms (appearance model)

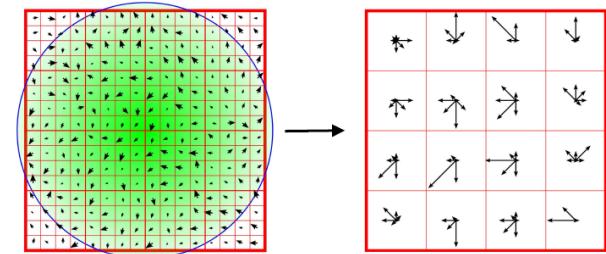
- Consider a patch of the image around the location of the part
- Construct statistical model for appearance at part location

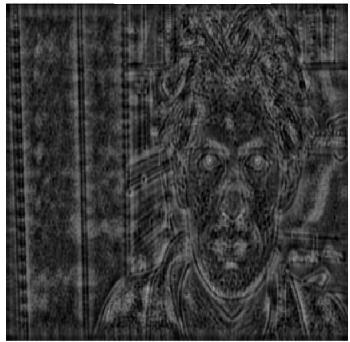


Filterbank responses



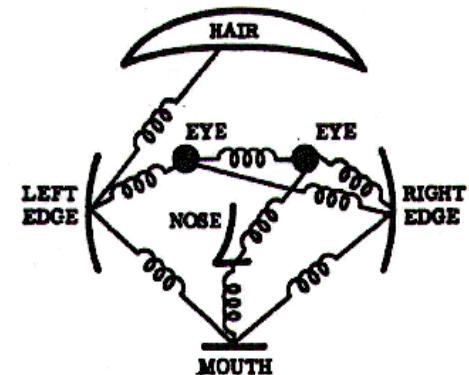
SIFT features



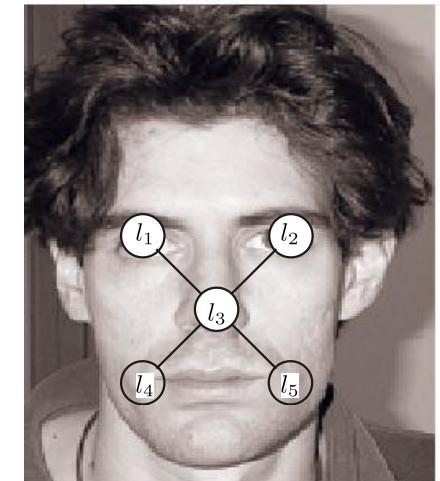
$\Phi_1(l)$  $\Phi_2(l)$  $\Phi_3(l)$  $\Phi_4(l)$  $\Phi_5(l)$ **Left eye****Right eye****Nose****Left mouth****Right mouth**

# Pairwise terms (deformation model)

- Deformable objects as physical systems
  - Masses: driven towards good image locations
  - Springs: enforce relations among parts



- Star models (tree-structured models)
  - Pick single part as root (e.g. the nose)
  - Model other parts (eyes, mouth corners) w.r.t. root
- Probability of configuration  $L = (l_1, l_2, l_3, l_4, l_5)$



$$P(L|I) \propto P(I|L)P(L) = \prod_i^{\text{Unary Terms}} \Phi_i(I_i|l_i) \prod_{(i,j) \in \mathcal{C}} \Psi_{i,j}(l_i, l_j)$$

**Unary Terms**      **Pairwise Terms**

- How can we maximize over L?



# Lecture outline

Introduction

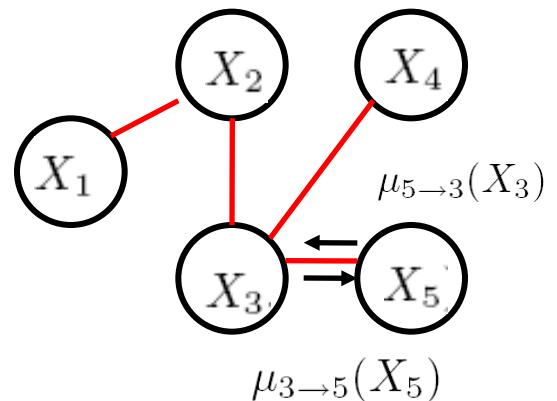
Chain-structured graphical models

Tree-structured graphical models

Pictorial Structures

Max/Sum-Product algorithm

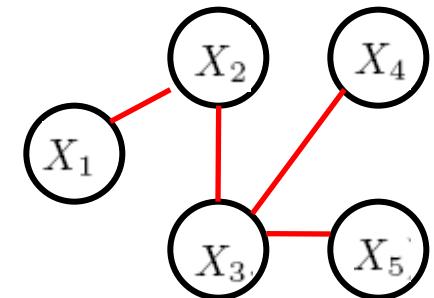
Loopy Graphical Models



# Computation on tree-structured Graphs

- Assume we want to find the most likely value of  $X_1$

$$\begin{aligned} P(X) &= \frac{1}{Z} \prod_{i=1}^5 \Phi_i(X_i) \prod_{(i,j) \in \mathcal{C}} \Psi_{i,j}(X_i, X_j) \\ \mathcal{C} &= \{(1,2), (2,3), (3,4), (3,5)\} \end{aligned}$$



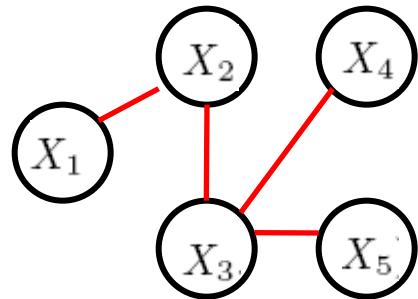
- Brute-force maximization:

$$\max P(X_1) = \max_{X_2, X_3, X_4, X_5} P(X_1, X_2, X_3, X_4, X_5)$$

$$|X_1| = K \rightarrow K^4$$

- Exploit factorization

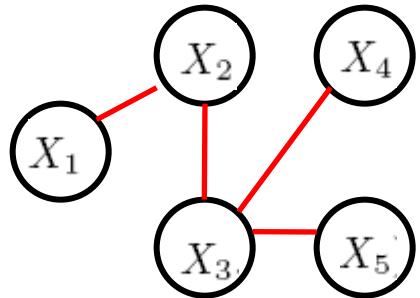
# Exploiting the factorization



$$P(X) = \frac{1}{Z} \prod_{i=1}^5 \Phi_i(X_i) \prod_{(i,j) \in \mathcal{C}} \Psi_{i,j}(X_i, X_j)$$
$$\mathcal{C} = \{(1,2), (2,3), (3,4), (3,5)\}$$

$\max P(X_1)$

# Exploiting the factorization



$$P(X) = \frac{1}{Z} \prod_{i=1}^5 \Phi_i(X_i) \prod_{(i,j) \in \mathcal{C}} \Psi_{i,j}(X_i, X_j)$$

$$\mathcal{C} = \{(1,2), (2,3), (3,4), (3,5)\}$$

$$\begin{aligned}
 & \max P(X_1) \\
 = & \max_{X_2, X_3, X_4, X_5} \Phi(X_1) \Phi(X_2) \Phi(X_3) \Phi(X_4) \Phi(X_5) \Psi(X_1, X_2) \Psi(X_2, X_3) \Psi(X_3, X_4) \Psi(X_3, X_5) \\
 = & \Phi(X_1) \max_{X_2} \Phi(X_2) \Psi(X_1, X_2) \max_{X_3, X_4} \Phi(X_3) \Psi(X_2, X_3) \Phi(X_4) \Psi(X_3, X_4) \underbrace{\max_{X_5} \Phi(X_5)}_{\mu_5(X_3)} \Psi(X_3, X_5) \\
 = & \Phi(X_1) \max_{X_2} \Phi(X_2) \Psi(X_1, X_2) \max_{X_3} \Phi(X_3) \Psi(X_2, X_3) \mu_5(X_3) \underbrace{\max_{X_4} \Phi(X_4)}_{\mu_4(X_3)} \Psi(X_3, X_4) \\
 = & \Phi(X_1) \max_{X_2} \Phi(X_2) \Psi(X_1, X_2) \underbrace{\max_{X_3} \Phi(X_3)}_{\mu_3(X_2)} \Psi(X_2, X_3) \mu_5(X_3) \mu_4(X_3) \\
 = & \Phi(X_1) \underbrace{\max_{X_2} \Phi(X_2)}_{\mu_2(X_1)} \Psi(X_1, X_2) \mu_3(X_2)
 \end{aligned}$$

# Max-Product algorithm

- Distributed computation on a graph
  - ‘message-passing’

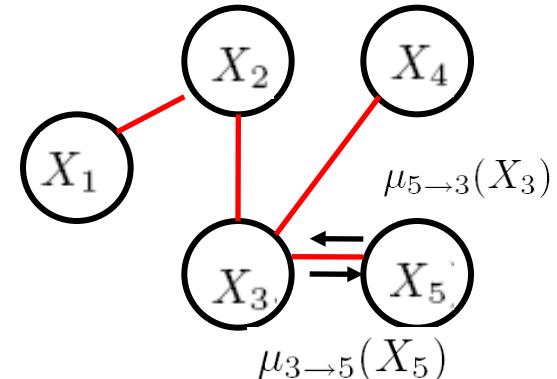
- Message from node i to node j

$$\mu_{i \rightarrow j}(X_j) = \max_{X_i} \Phi_i(X_i) \Psi_{i,j}(X_i, X_j) \prod_{k \in \mathcal{N}(i) \setminus j} \mu_{k \rightarrow j}(X_i)$$

- ‘Given all I know from the rest, this is where you should be’

- Beliefs:
$$B_j(X_j) = \Phi_j(X_j) \prod_{i \in \mathcal{N}(j)} \mu_{i \rightarrow j}(X_j)$$

- At convergence
$$B_j(X_j) = \max P(X_j)$$



# Sum-Product algorithm

- Distributed computation on a graph
  - ‘message-passing’

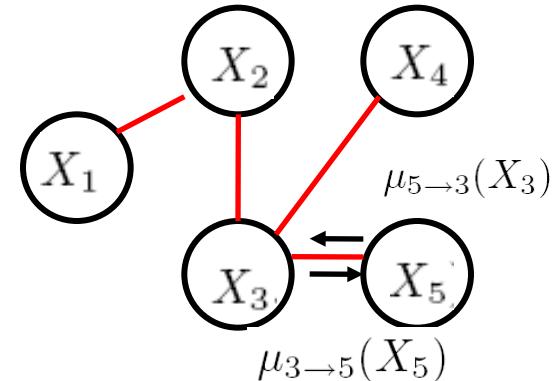
- Message from node i to node j

$$\mu_{i \rightarrow j}(X_j) = \sum_{X_i} \psi_i(X_i) \psi_{i,j}(X_i, X_j) \prod_{k \in \{\mathcal{N}(i) \setminus j\}} \mu_{k \rightarrow i}(X_i)$$

- ‘Given all I know from the rest, this is where you should be’

- Beliefs:
$$B_j(X_j) = \Phi_j(X_j) \prod_{i \in \mathcal{N}(j)} \mu_{i \rightarrow j}(X_j)$$

- At convergence
$$P(X_i) = B(X_i)$$



# Max/Sum-Product algorithms

- Asynchronous inference
- Local computations
- Guaranteed optimality for tree-structured graphs
- Equivalent algorithms:
  - Tracking: Kalman Filtering (Sum-Product)
  - HMMs: Alpha-Beta (Sum-Product)
  - Coding/HMMs: Viterbi Algorithm (Max-Product)
  - Bayesian Nets: Belief Propagation (Sum-Product)
- Max Product: Dynamic Programming

# Deformable Part Models (DPMs)

$$\mathbf{S}(\mathbf{x}) = \sum_{p=1}^P U_p(x_p) + B_p(x, x_p)$$

Local appearance

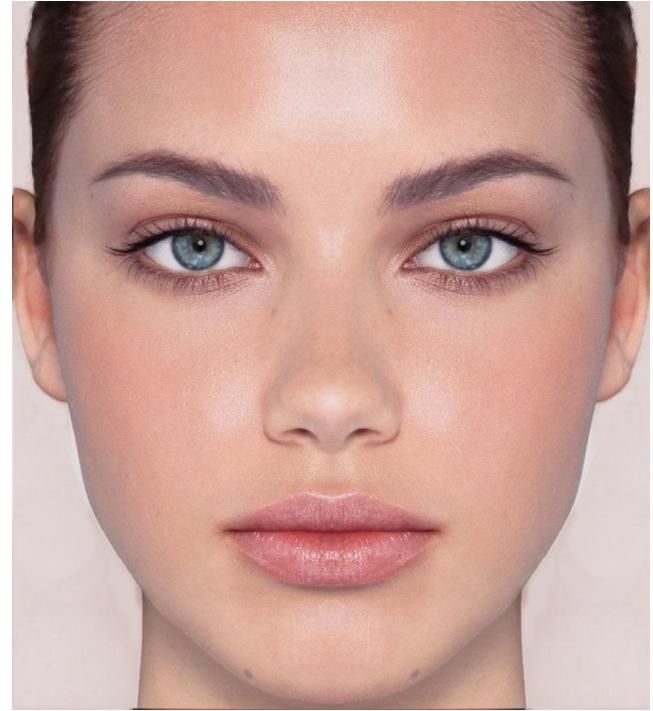
$$U_p(x_p) = \langle w_p, H(x_p) \rangle$$

Pairwise compatibility

$$B_p(x, x_p) = -(h - h_p - \hat{h}_p)^2 \eta - (v - v_p - \hat{v}_p)^2 \nu$$

**Object score**

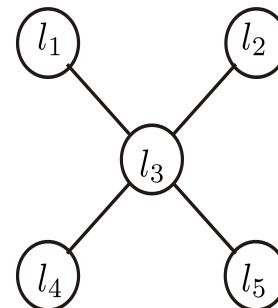
$$S(x) = \sum_{p=1}^P \max_{x'} [U_p(x') + B_p(x, x')]$$



# Pictorial Structures

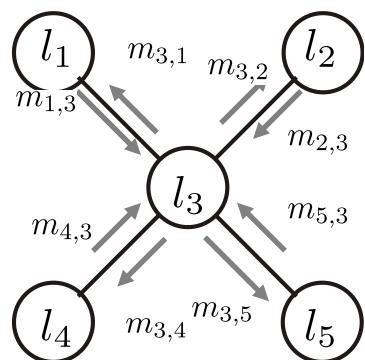
- Graphical model
  - Nodes: part locations
  - Edges: dependencies - relative locations
- Object localization: maximization with respect to  $L = (l_1, l_2, l_3, l_4, l_5)$ 
$$P(L|I) \propto P(I|L)P(L) = \prod_i \Phi_i(I_i|l_i) \prod_{(i,j) \in \mathcal{C}} \Psi_{i,j}(l_i, l_j)$$

**Unary terms**                    **Pairwise terms**
- Message Passing:
$$\mu_{i \rightarrow j}(X_j) = \max_{X_i} \Phi_i(X_i) \Psi_{i,j}(X_i, X_j) \prod_{k \in \mathcal{N}(i) \setminus j} \mu_{k \rightarrow j}(X_i)$$



$$\mu_{i \rightarrow j}(X_j) = \max_{X_i} \Phi_i(X_i) \Psi_{i,j}(X_i, X_j) \prod_{k \in \mathcal{N}(i) \setminus j} \mu_{k \rightarrow j}(X_i)$$

$$B_j(X_j) = \max P(X_j)$$



$$\Phi_1(l)$$



$$\Phi_2(l)$$



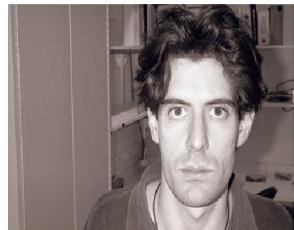
$$\Phi_3(l)$$



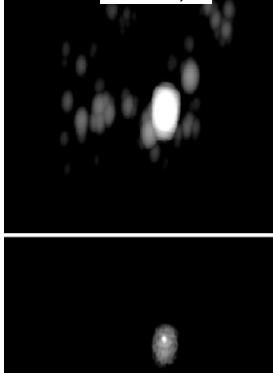
$$\Phi_4(l)$$



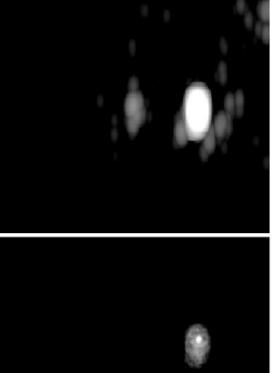
$$\Phi_5(l)$$



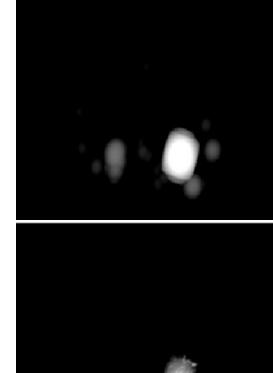
$$m_{1,3}$$



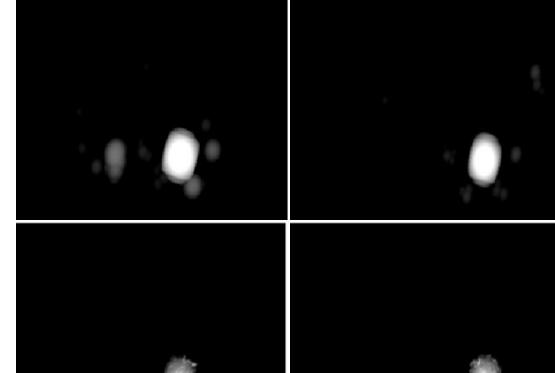
$$m_{2,3}$$



$$m_{3,4}$$



$$m_{3,5}$$



$$B_1(l)$$

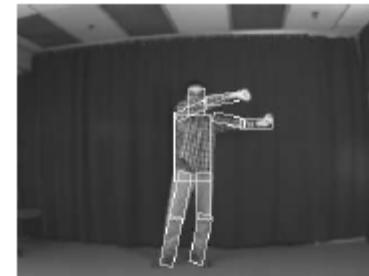
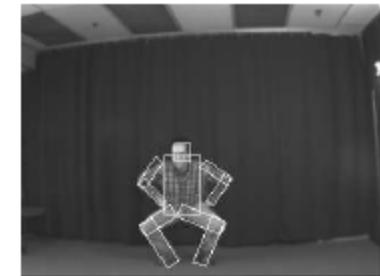
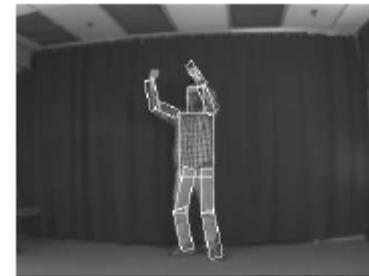
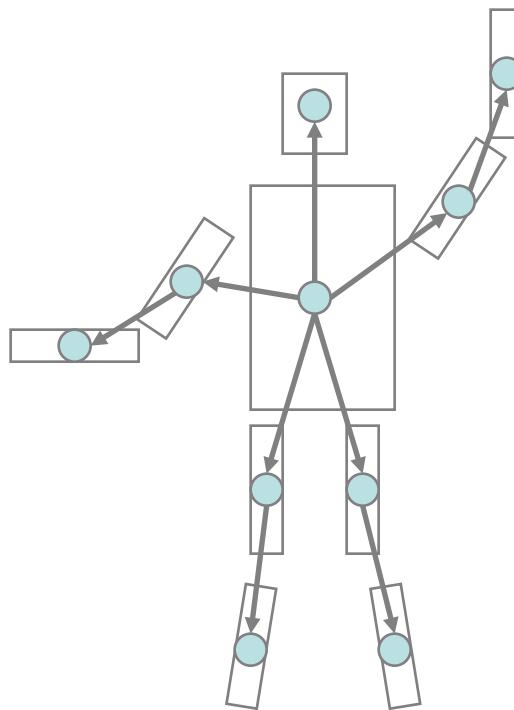
$$B_2(l)$$

$$B_4(l)$$

$$B_5(l)$$

$$B_3(l)$$

# Detection Results



# Part-based models + discriminative framework

**P. Felzenszwalb, D. McAllester, D. Ramanan`A Discriminatively Trained, Multiscale, Deformable Part Model' CVPR 2008**

- Linear classifier gives individual part cost
- Gaussian-like cost for spatial offsets
- Learn cost for parts and deformations
  - Discriminative framework



- State-of-the-art, 2008

# Lecture outline



Introduction

Chain-structured graphical models

Tree-structured graphical models

Pictorial Structures

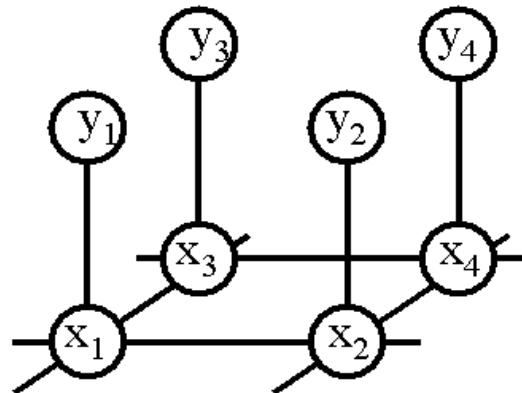
Max/Sum-Product algorithm

Loopy Graphical Models

Markov Random Fields

Graph Cuts

Gaussian Random Fields



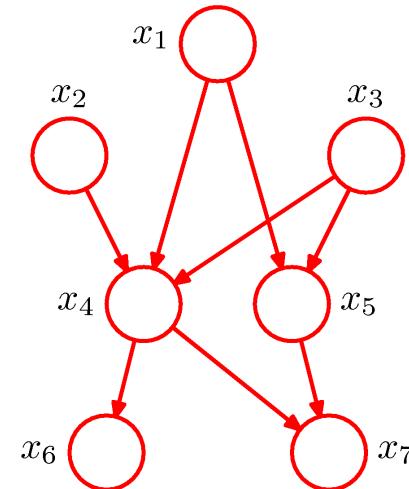
# Bayesian Networks

Directed Acyclic Graph  $G(V, E)$

Nodes -  $V$  : random variables

Edges -  $E$  : dependencies

Leave from 'parents', arrive at children

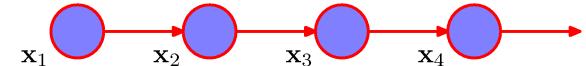


Distribution of each child conditioned on parent:  $P(X_v|X_{\pi_v})$

$$P(X_v|\emptyset) = P(X_v)$$

Joint probability distribution:  $P(X) = \prod_v P(X_v|X_{\pi_v})$

Bayesian network for Brownian motion:

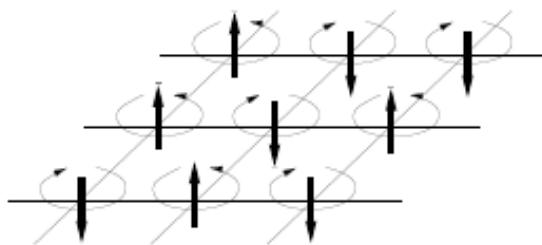


# Connection with statistical physics

Idea: image pixels are elements of a large physical system

Ising model: iron atoms on a lattice

Model spin orientations



Interactions among spin components

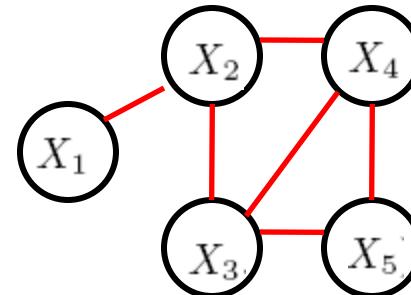
Nearby spins tend to be similar

Interactions among pixel and environment

External magnetization

# Markov Random Fields

- Undirected Graph  $G(V, E)$



- Maximal graph cliques (fully-connected subsets)

- Clique potentials

$$\mathcal{C} = \{(X_1, X_2), (X_2, X_3), (X_3, X_4, X_5)\}$$

$$\psi_c(x_c) > 0, \forall x_c.$$

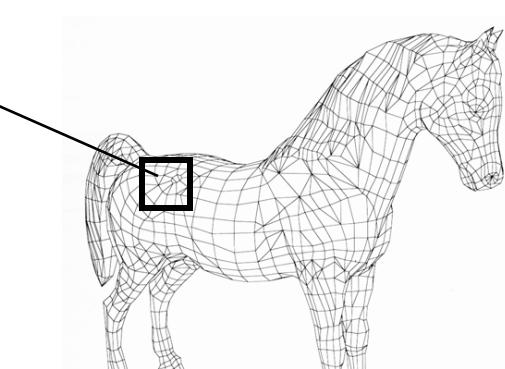
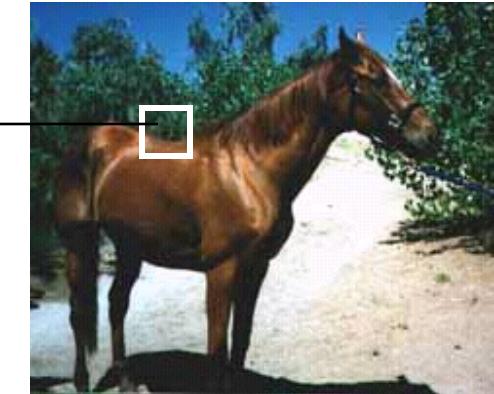
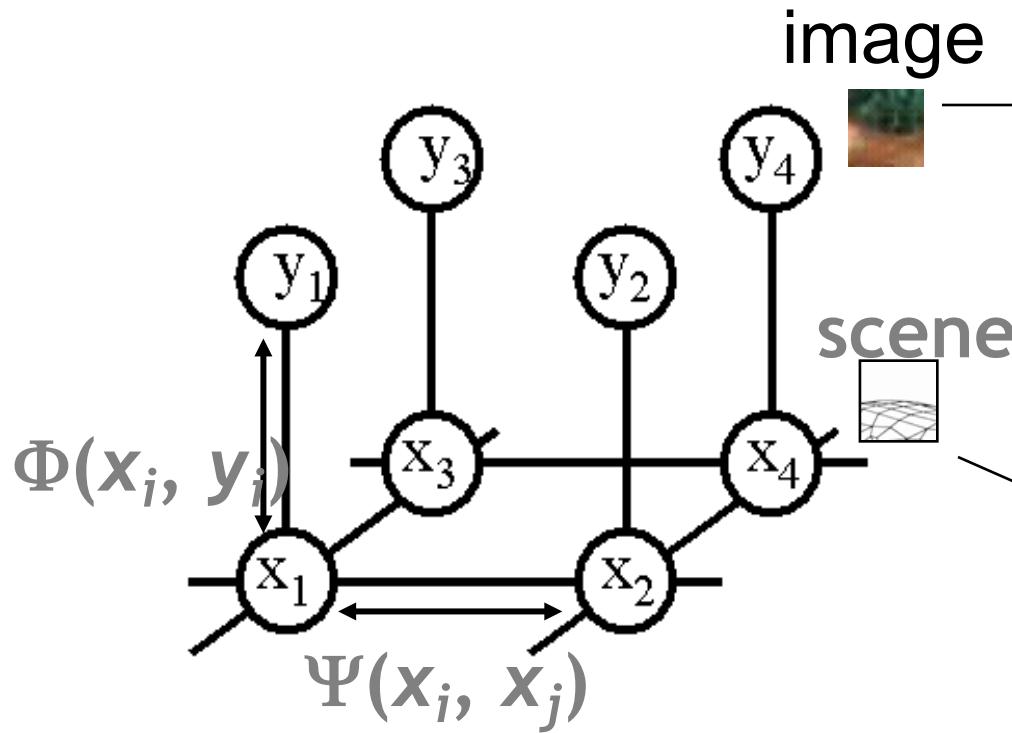
- Joint PDF

$$P(X) = \frac{1}{Z} \prod_{c \in \mathcal{C}} \psi_c(X_c)$$

$$Z = \sum_X \prod_{c \in \mathcal{C}} \psi_c(X_c)$$

# MRFs for Vision

$$P(X, Y) = \frac{1}{Z} \prod_i \Phi(Y_i, X_i) \prod_{i,j \in \mathcal{E}} \Psi(X_i, X_j)$$



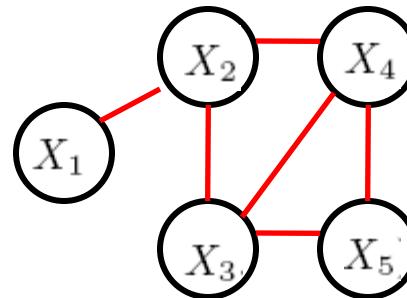
scene

# Network Joint Probability

$$P(X, Y) = \frac{1}{Z} \prod_i \Phi(Y_i, X_i) \prod_{i,j \in \mathcal{E}} \Psi(X_i, X_j)$$

# Markov Random Fields

- Undirected Graph  $G(V, E)$



- Maximal graph cliques (fully-connected subsets)

$$\mathcal{C} = \{(X_1, X_2), (X_2, X_3), (X_3, X_4, X_5)\}$$

- Clique potentials  $\psi_c(x_c) > 0, \forall x_c$

$$p(x_1, \dots, x_N) = \frac{1}{Z} \prod_{c \in \mathcal{C}} f_c\left(\{x_j | j \in c\}\right) = \frac{1}{Z} \prod_{c \in C} f_c(\mathbf{x}_c)$$

normalization term  
(partition function)

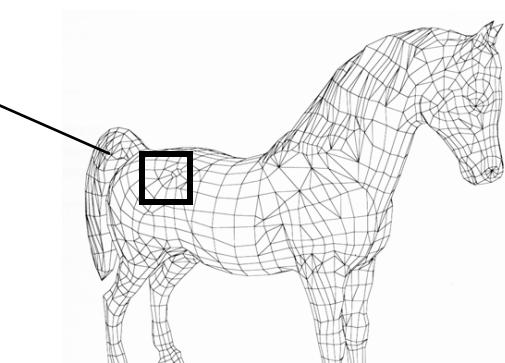
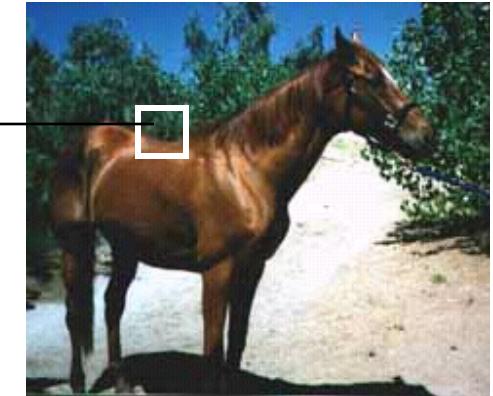
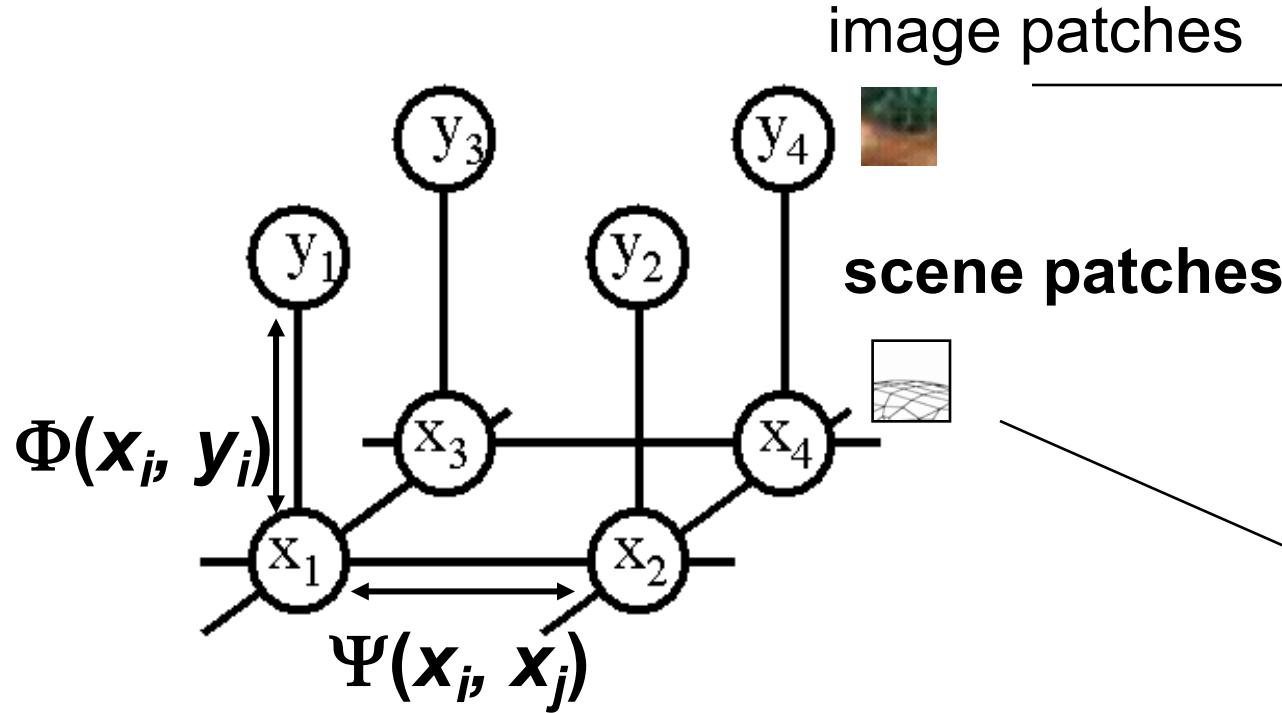
cliques of the graph

factors or potentials  
(non-negative, but  
unnormalized)

variables  
in clique  $c$

# MRFs for Vision

$$P(X, Y) = \frac{1}{Z} \prod_i \Phi(Y_i, X_i) \prod_{(i,j) \in C} \Psi(X_i, X_j)$$



$$P(X|Y) = ?$$

scene

# Network Joint Probability

$$P(X, Y) = \frac{1}{Z} \prod_i \Phi(Y_i, X_i) \prod_{i,j \in \mathcal{E}} \Psi(X_i, X_j)$$

**Scene**

**Image**

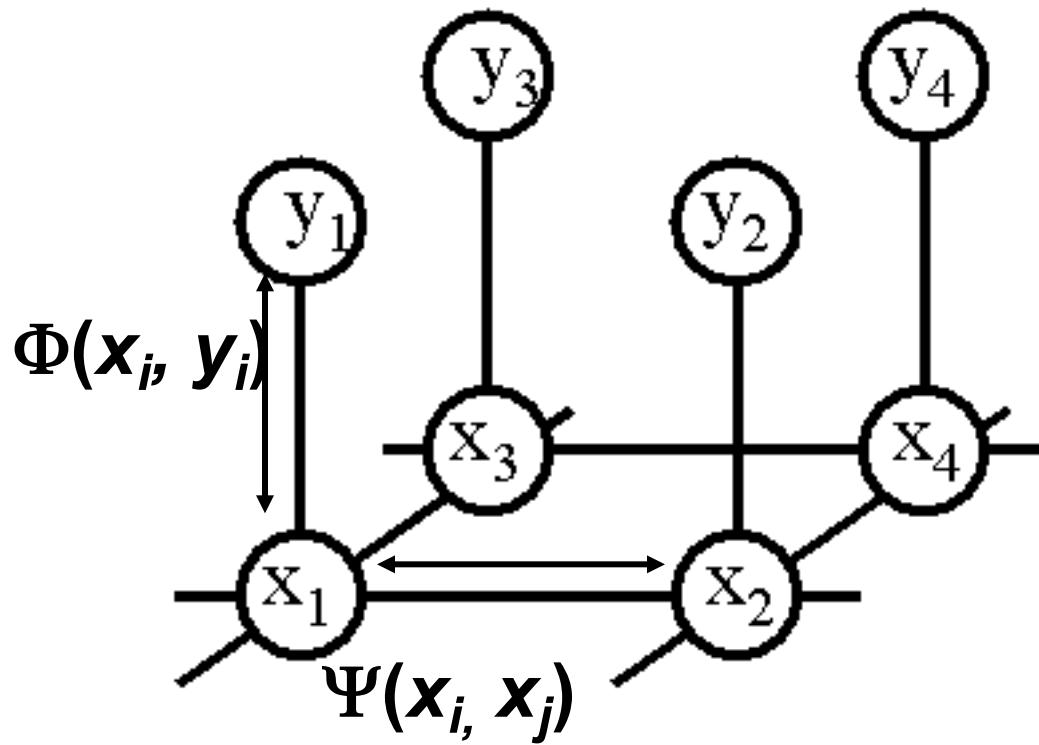
**Image-scene  
compatibility  
function**

**Local  
observations**

**Scene-scene  
compatibility  
function**

**Neighboring  
scene nodes**

# MRFs for Denoising



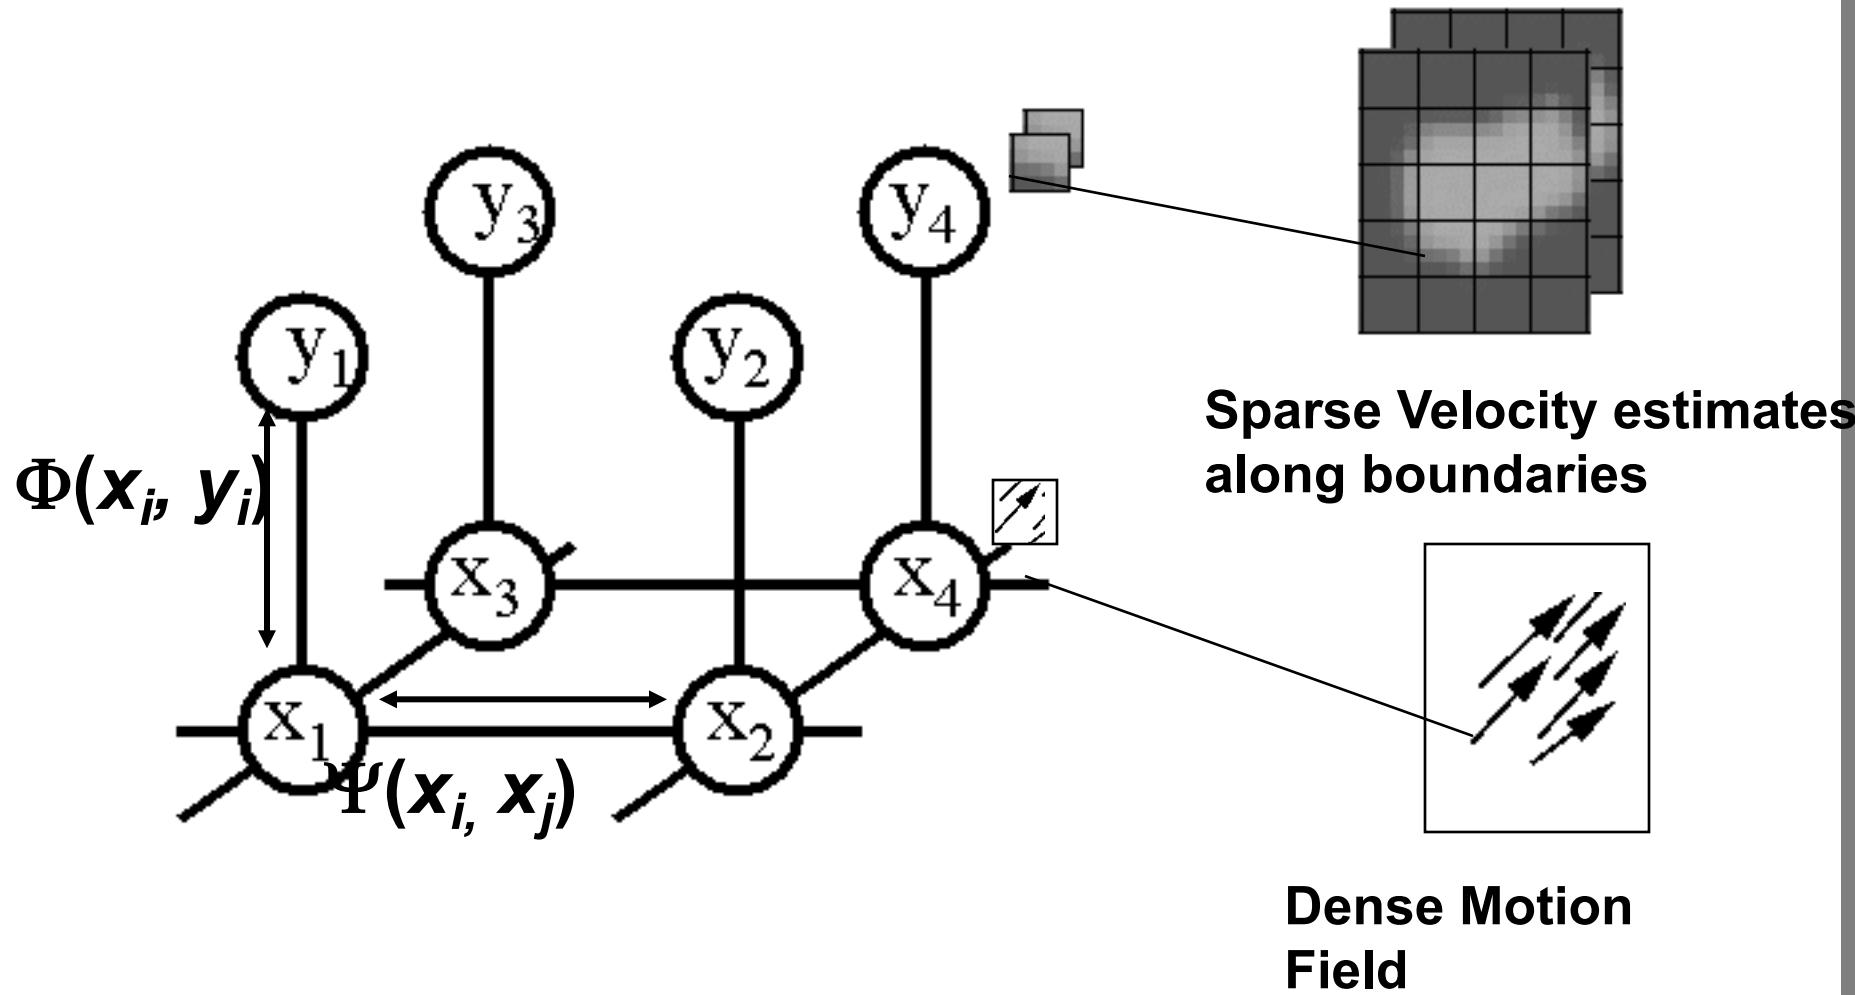
Bayes' Theorem

Noisy Pixel Intensities

Bayes' Theorem

Clean Image

# MRFs for Motion



# Lecture outline



Introduction

Chain-structured graphical models

Tree-structured graphical models

Pictorial Structures

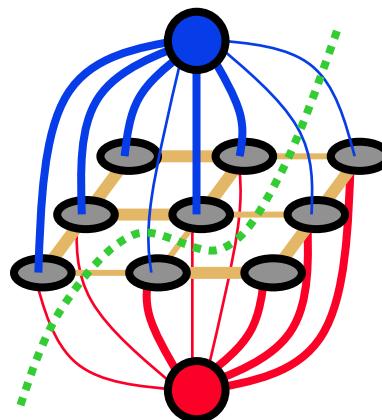
Max/Sum-Product algorithm

## Loopy Graphical Models

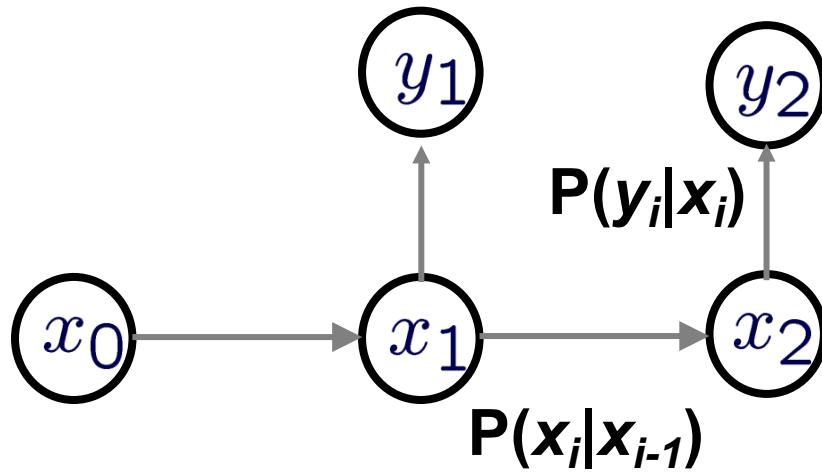
Markov Random Fields

Graph Cuts

Gaussian Random Fields



# Loops: Exact Computation becomes Hard!

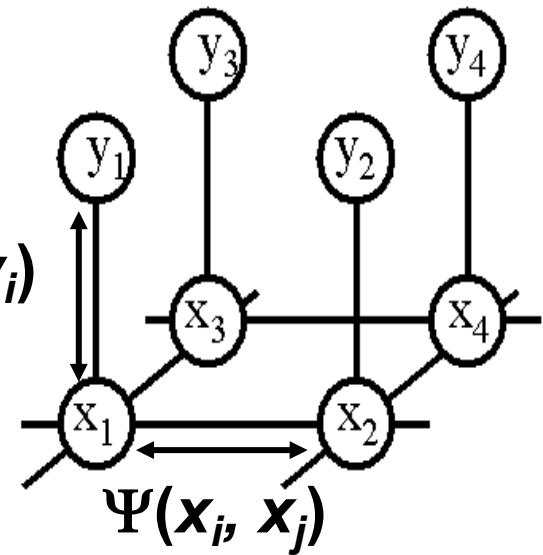


*observed  
process*

$$\Phi(x_i, y_i)$$

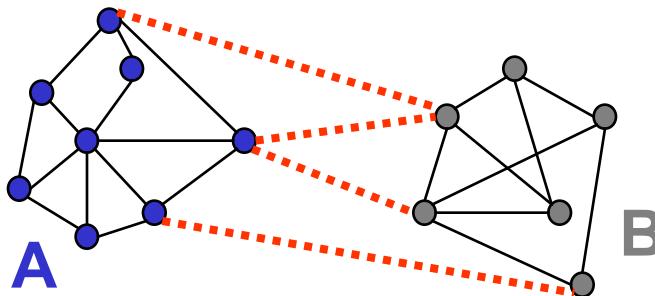
*hidden  
states*

Inference: Max/Sum Product (Polynomial)



NP-hard

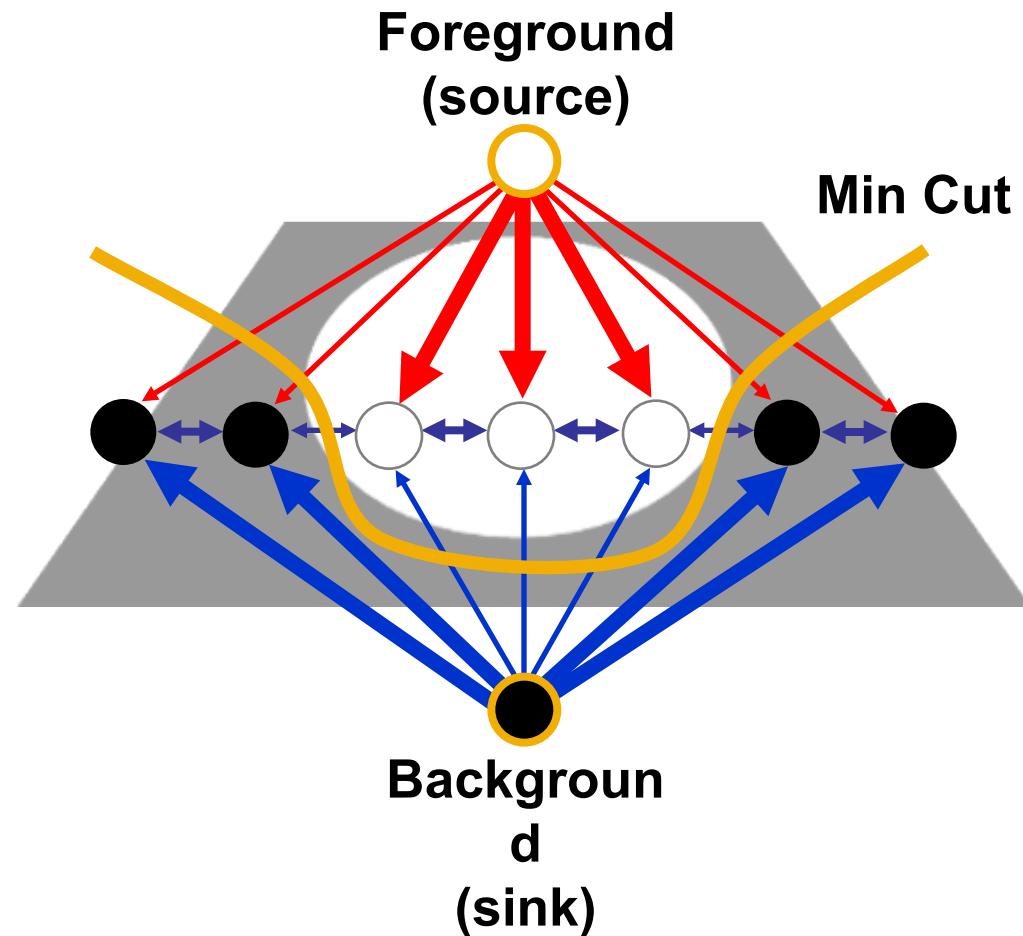
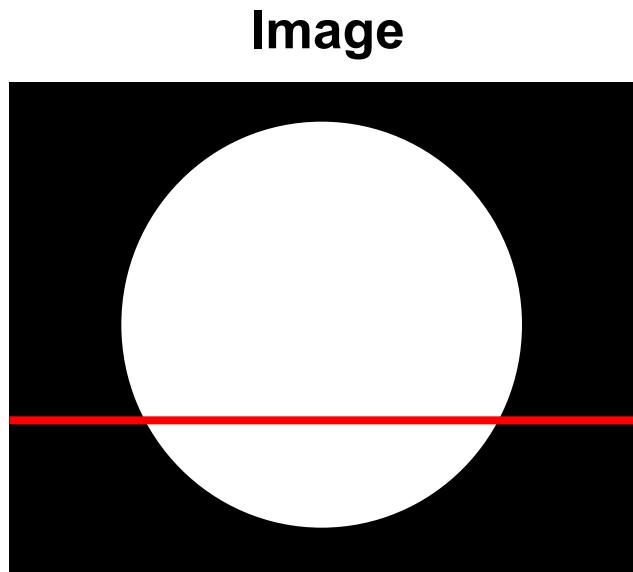
# Graph Cut



- Set of edges whose removal makes a graph disconnected
- Cost of a cut
  - Sum of weights of cut edges:
- A graph cut gives us a segmentation
  - What is a “good” graph cut and how do we find one?

$$cut(A, B) = \sum_{p \in A, q \in B} w_{p,q}$$

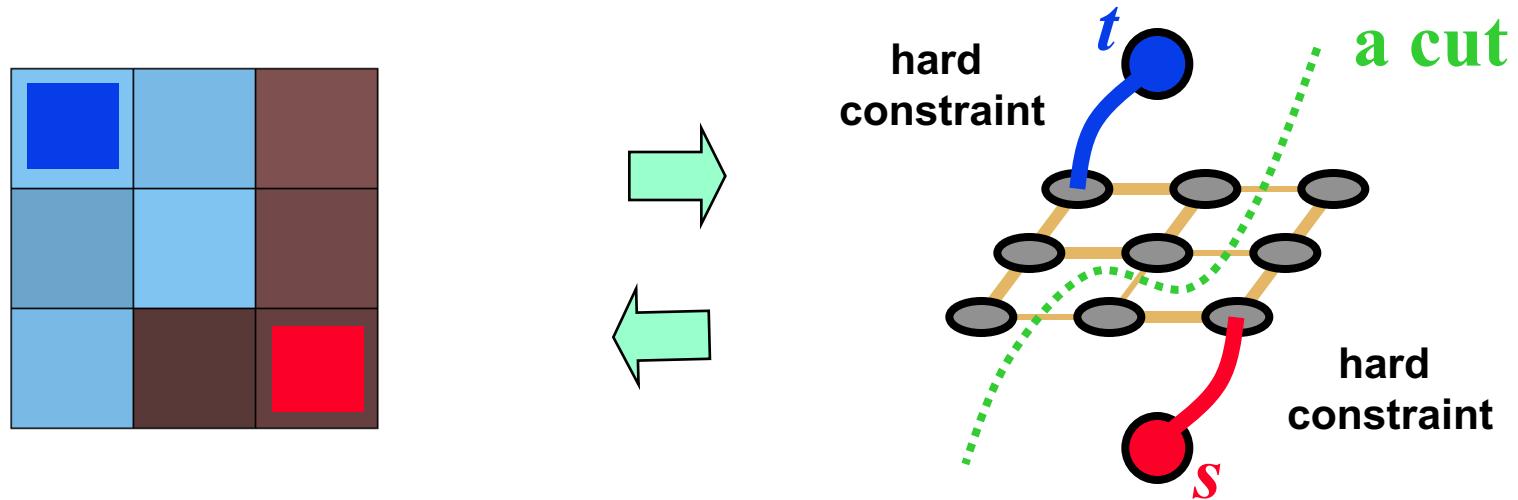
# Graph Cuts for Image Segmentation



- Segmentation by s-t mincut
  - *Cut*: separating source and sink
  - *Min Cut*: Global minimal energy in polynomial time (1MPixel/sec)

# Graph Cuts for Optimal Boundary Detection

- Idea: convert MRF into source-sink graph



**Minimum cost cut can be  
computed in polynomial  
time**

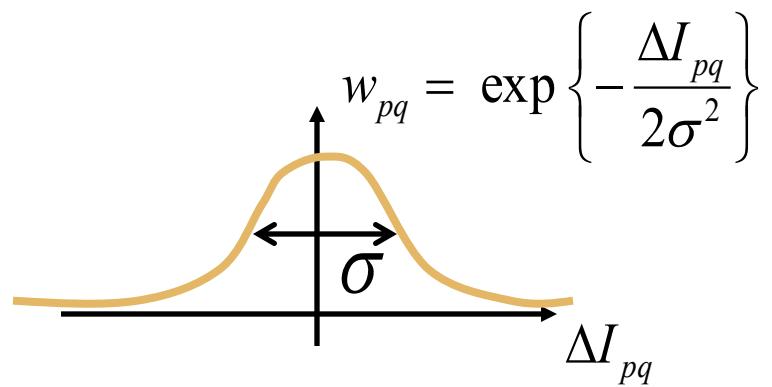
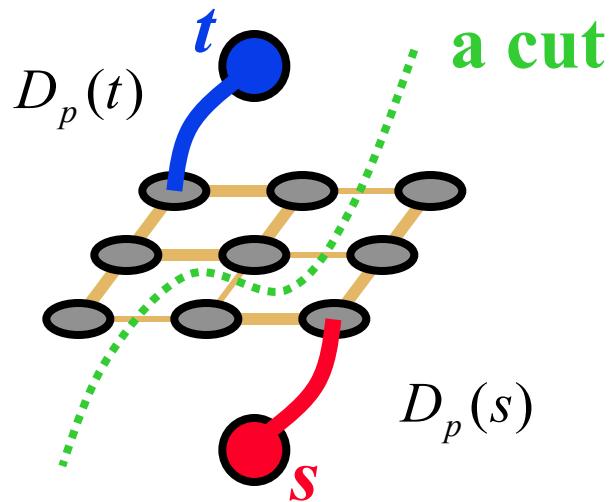
(max-flow/min-cut algorithms)

# Simple Example of Energy

$$E(L) = \sum_p D_p(L_p) + \sum_{pq \in N} w_{pq} \cdot \delta(L_p \neq L_q)$$

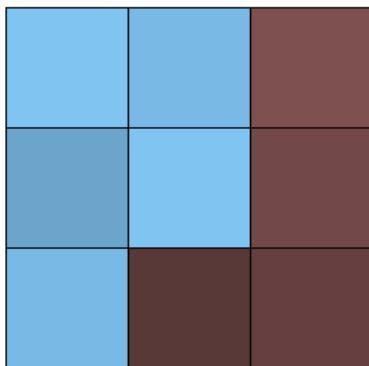
Regional term                      Boundary term

p                                    pq  $\in$  N  
t-links                            n-links

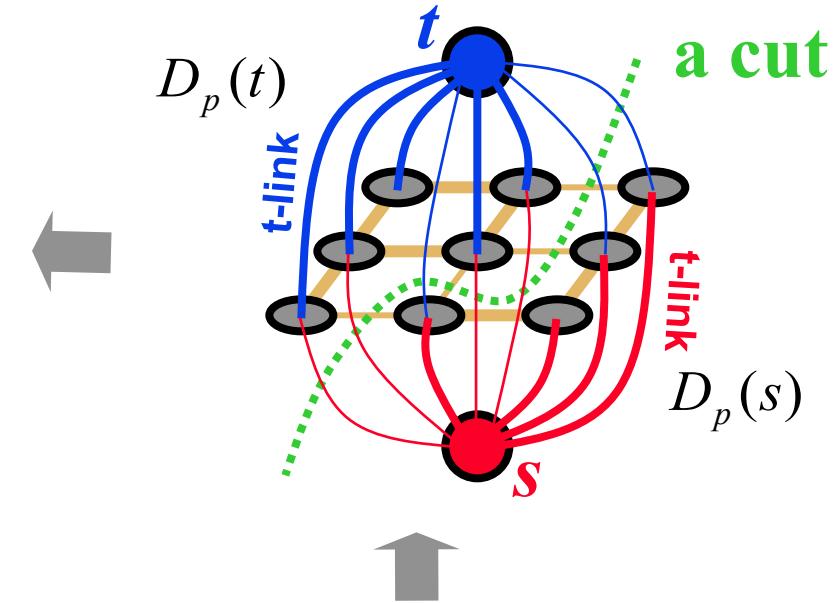


$L_p \in \{s, t\}$   
**(binary object segmentation)**

# Adding Regional Properties



“expected” intensities of  
**object and background**  
 $I^s$  and  $I^t$   
can be re-estimated



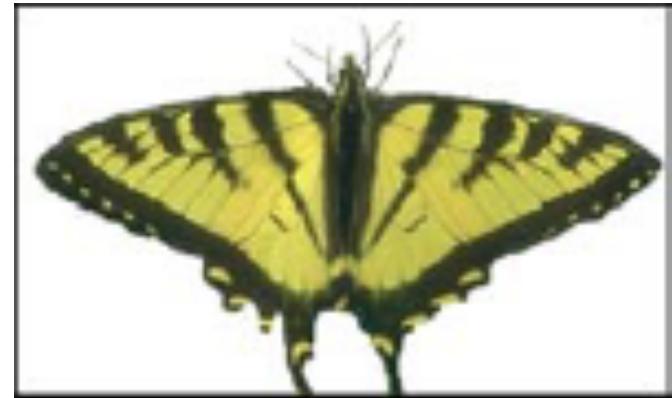
$$D_p(s) \propto \exp\left(-\|I_p - I^s\|^2 / 2\sigma^2\right)$$

$$D_p(t) \propto \exp\left(-\|I_p - I^t\|^2 / 2\sigma^2\right)$$

EM-style optimization

# Iterative learning of regional color-models

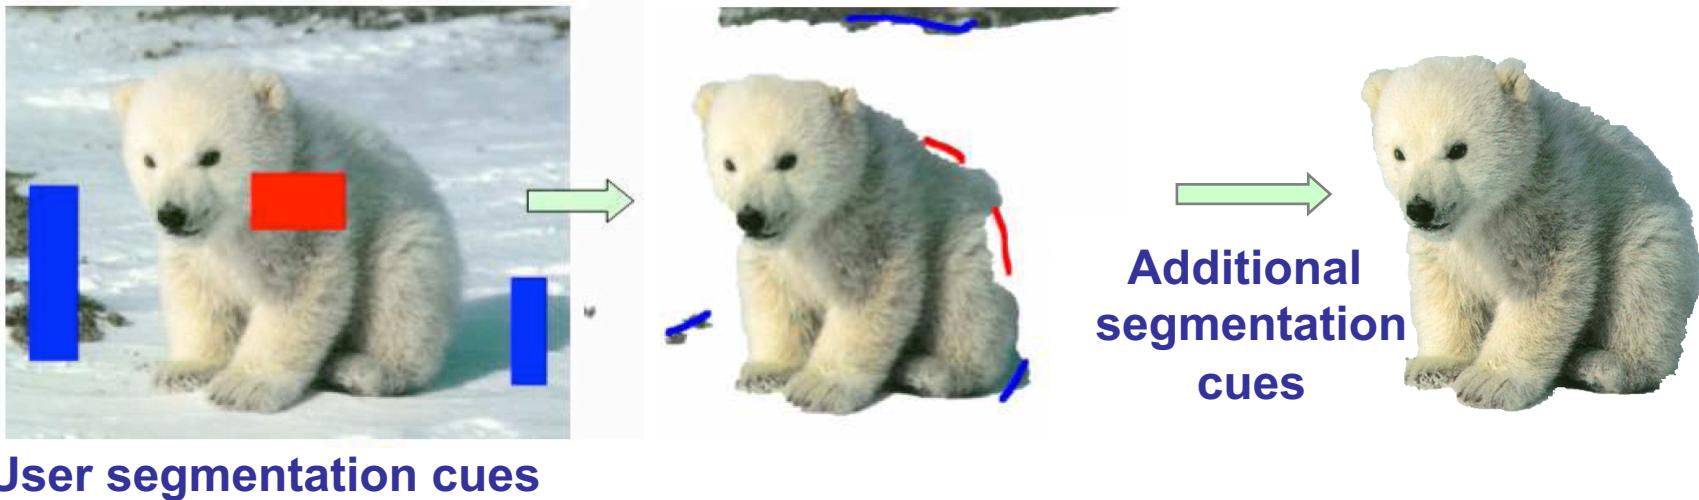
- GMMRF cuts (Blake et al., ECCV04)  
Region competition with Gaussian mixture distributions
- Grab-cut (Rother et al., SIGGRAPH 04)



**parametric regional model – Gaussian Mixture (GM)  
designed to guarantee convergence**

# GraphCut Applications: “GrabCut”

- Procedure
  - User marks foreground and background regions with a brush.
  - This is used to create an initial segmentation which can then be corrected by additional brush strokes.

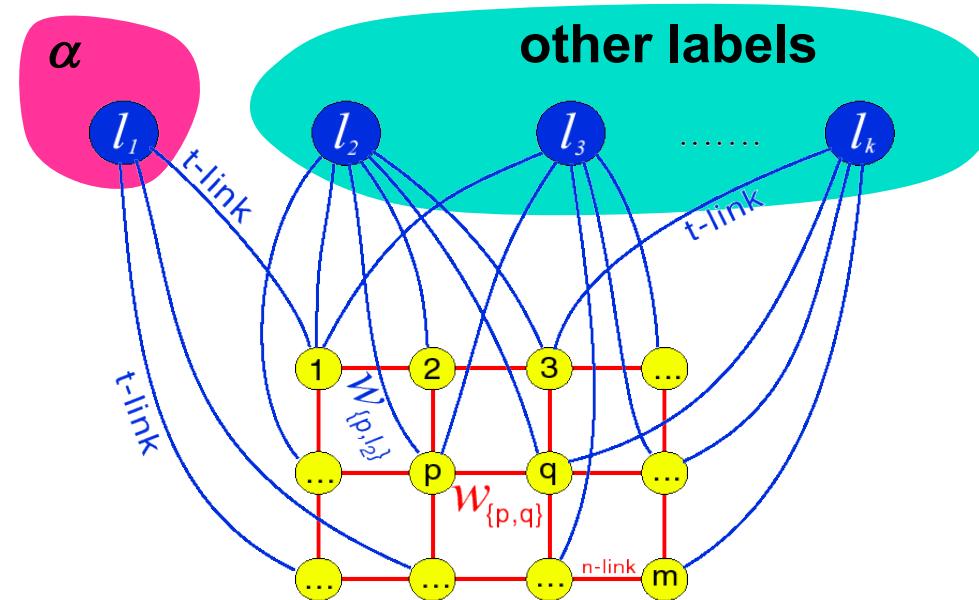


# GrabCut: Example Results



# More labels: $\alpha$ -Expansion

- Basic idea: k vs rest
  - Break multi-way cut computation into a sequence of binary s-t cuts.





# Lecture outline

Introduction

Chain-structured graphical models

Tree-structured graphical models

Pictorial Structures

Max/Sum-Product algorithm

## Loopy Graphical Models

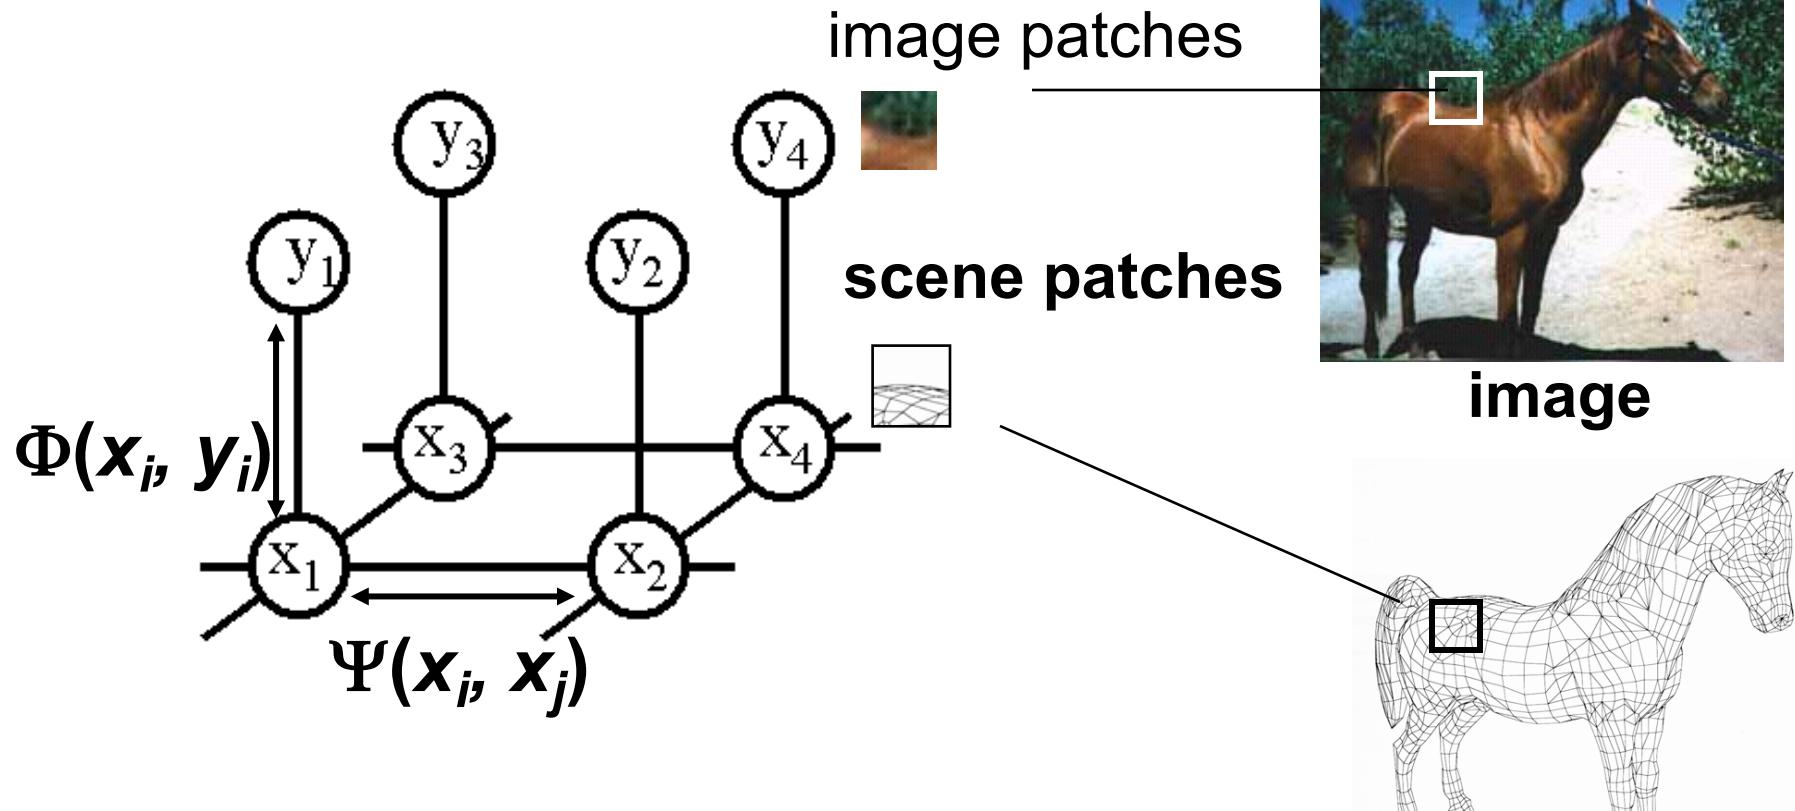
Markov Random Fields

Graph Cuts

Gaussian Random Fields

# Markov Random Fields in Vision

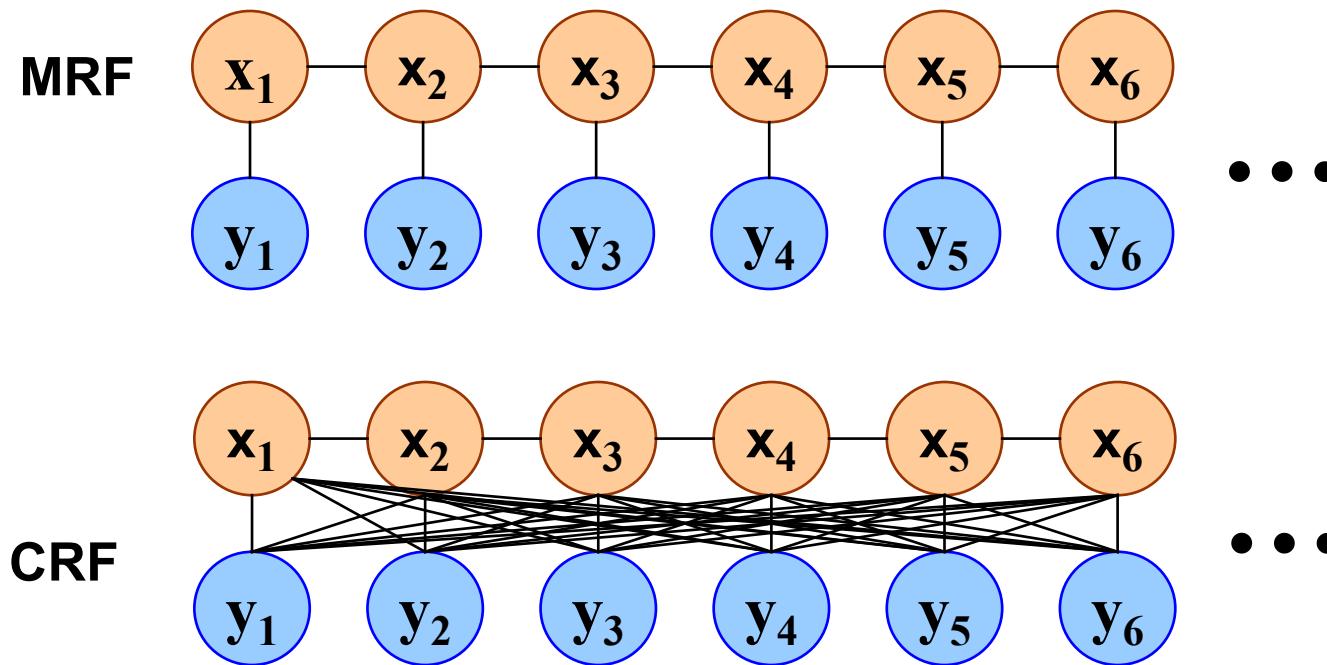
$$P(X, Y) = \frac{1}{Z} \prod_i \Phi(Y_i, X_i) \prod_{(i,j) \in C} \Psi(X_i, X_j)$$



$$P(X|Y) = ?$$

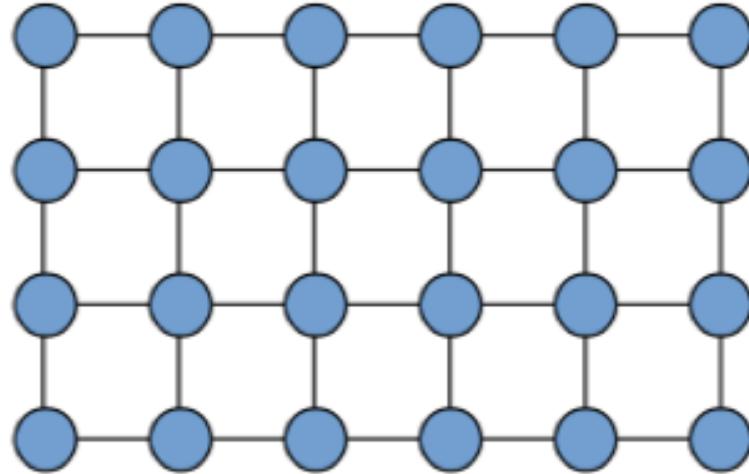
scene

# Conditional Random Fields



**CRFs: keep MRF tools, drop Bayesian aspect**

# Conditional Random Field

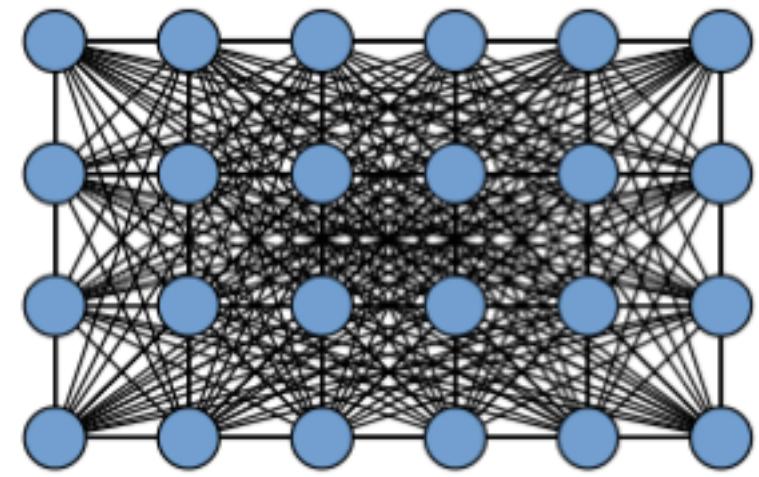
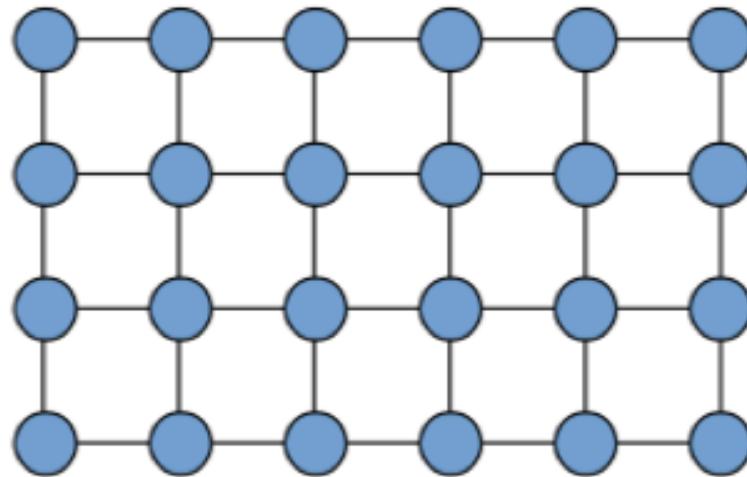


$$P(\mathbf{x}|\mathbf{y}) = \frac{1}{Z_{\mathbf{y}}} \exp(-E_{\mathbf{y}}(\mathbf{x}))$$

$$E_{\mathbf{y}}(\mathbf{x}) = \sum_i U_i(\mathbf{x}_i, \mathbf{y}) + \sum_{i,j \in \mathcal{E}} B_{i,j}(\mathbf{x}_i, \mathbf{x}_j, \mathbf{y})$$

**pre-softmax FCNN**      **Pairwise terms**

# Dense Conditional Random Field



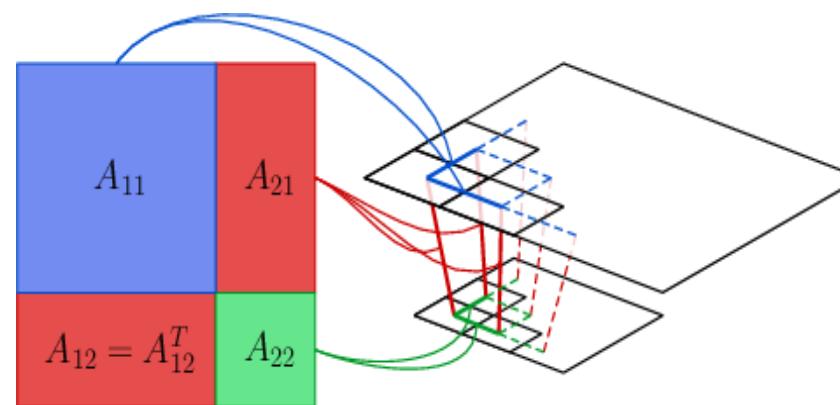
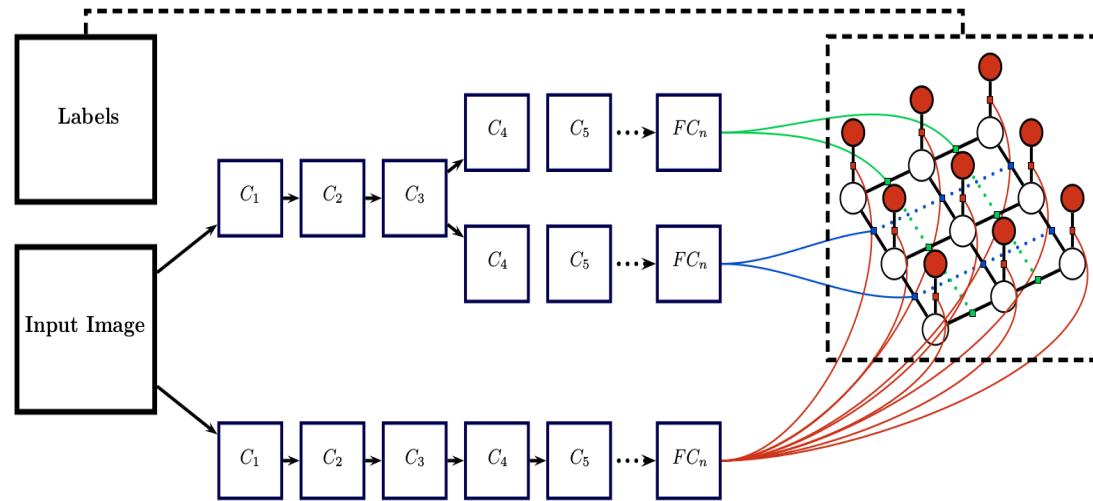
$$P(\mathbf{x}|\mathbf{y}) = \frac{1}{Z_{\mathbf{y}}} \exp(-E_{\mathbf{y}}(\mathbf{x}))$$

$$E_{\mathbf{y}}(\mathbf{x}) = \sum_i U_i(\mathbf{x}_i, \mathbf{y}) + \sum_{i,j \in \mathcal{E}} B_{i,j}(\mathbf{x}_i, \mathbf{x}_j, \mathbf{y})$$

**pre-softmax FCNN**      **Pairwise terms**

# Fast, Exact, and Multi-Scale Inference for FCNN-CRF

S. Chandra

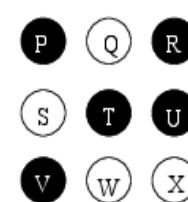


S. Chandra, I. Kokkinos, Fast, Exact and Multi-Scale Inference for Semantic Image Segmentation with Deep Gaussian CRFs, arXiv:1603.08358

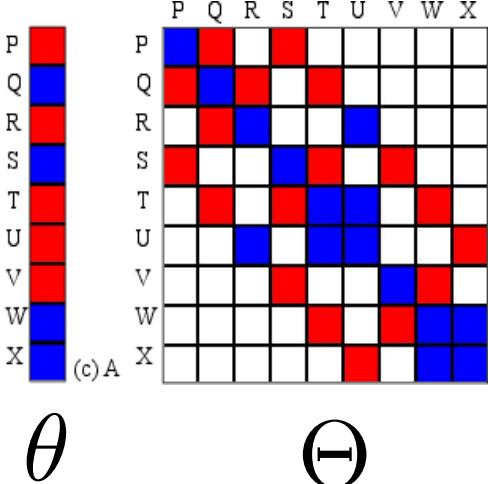
# Gaussian Random Fields: Random Fields for dummies

$$\pi(\mathbf{x}) = \frac{1}{Z} \exp \left( -\frac{1}{2} \mathbf{x}^T \Theta \mathbf{x} + \theta^T \mathbf{x} \right)$$

$$\Theta \mathbf{x}^* = \theta$$



$$\mathbf{x}^*$$



**Maximum-A-Posteriori inference =  
Minimum Mean-Squared Error inference =  
solution of linear system**

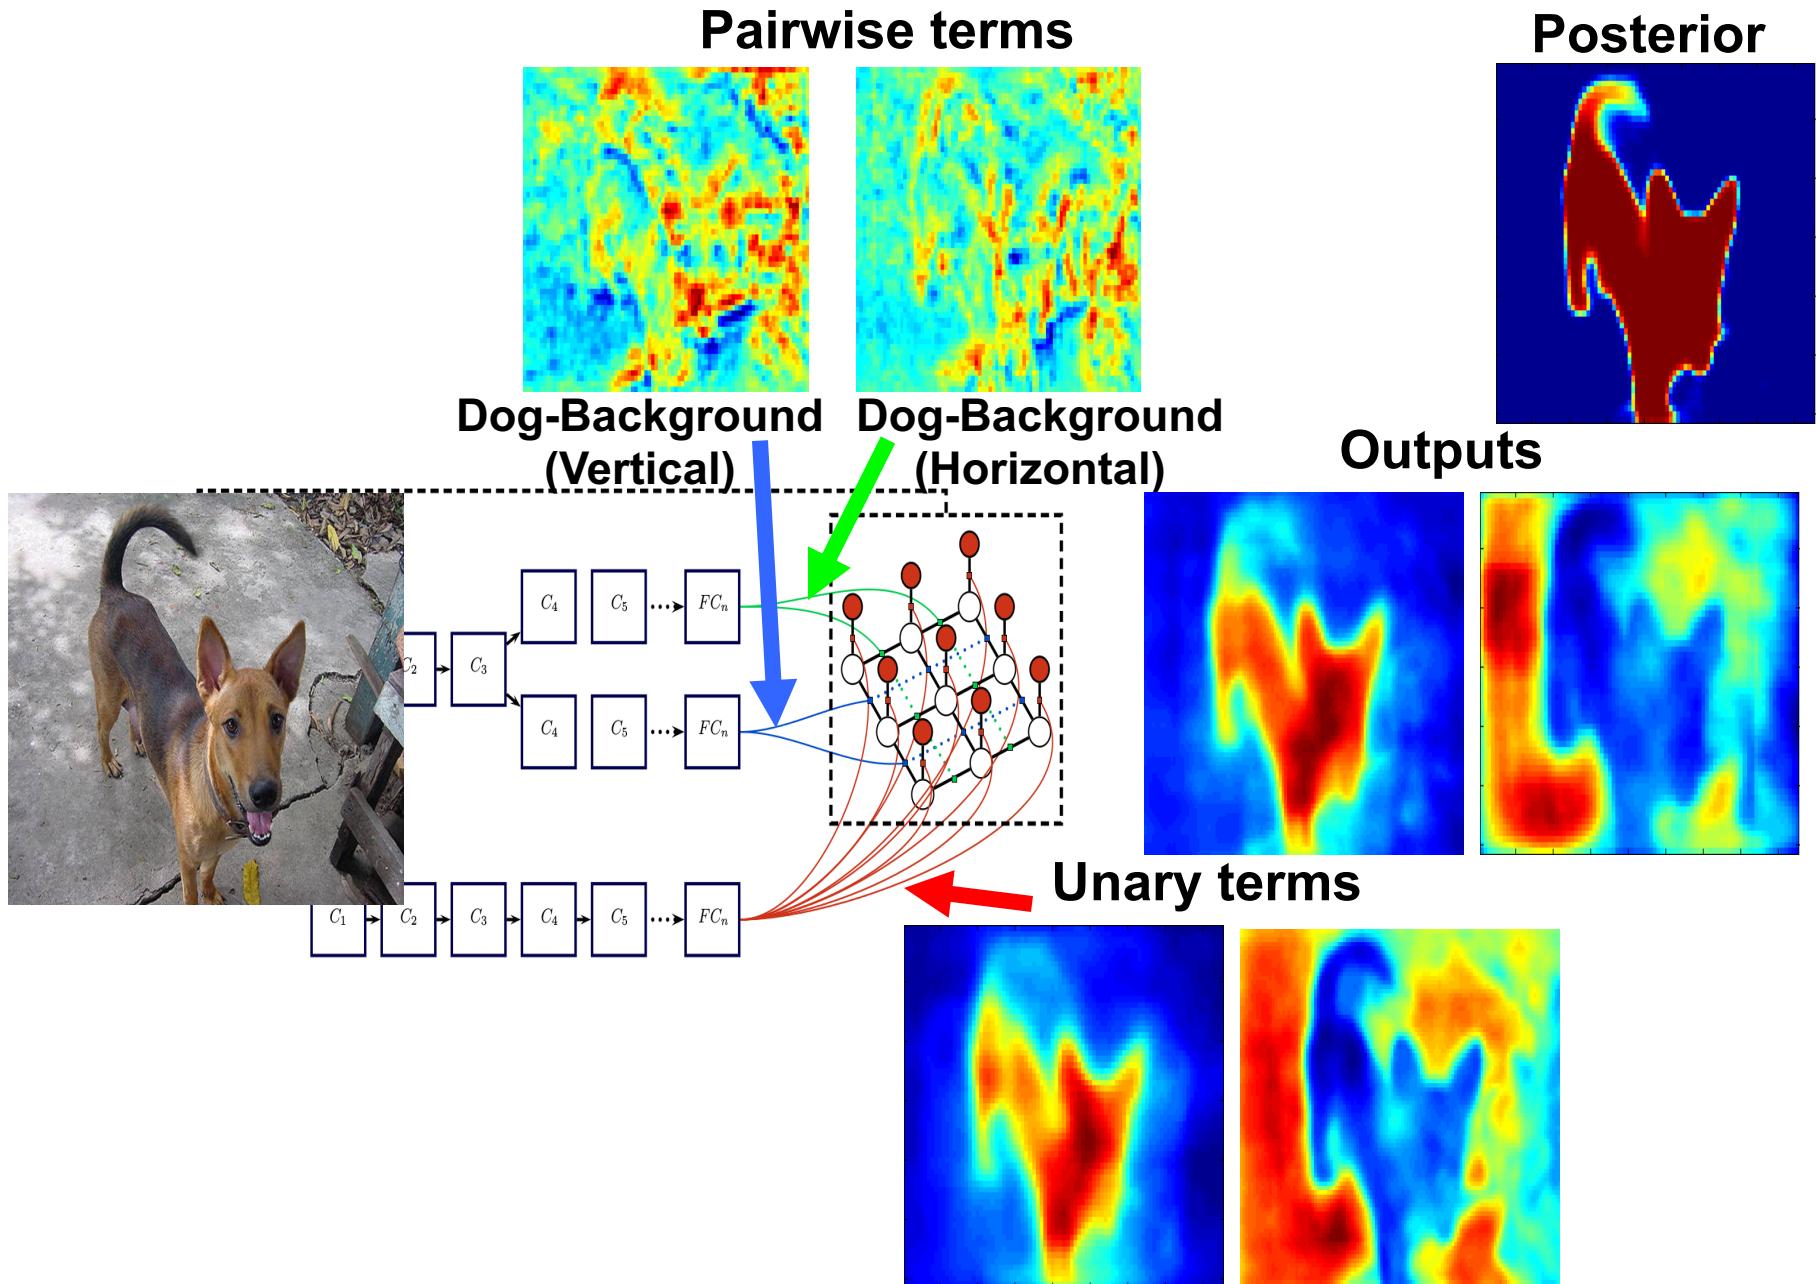
Gaussian MRF: blurry samples (hard to have outliers)

Gaussian CRF: image-based pairwise terms (e.g. discontinuity-preserving)

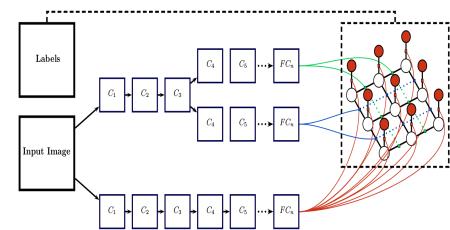
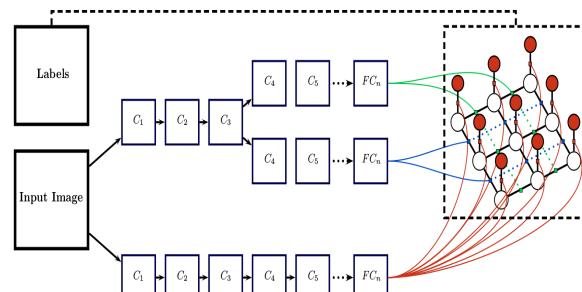
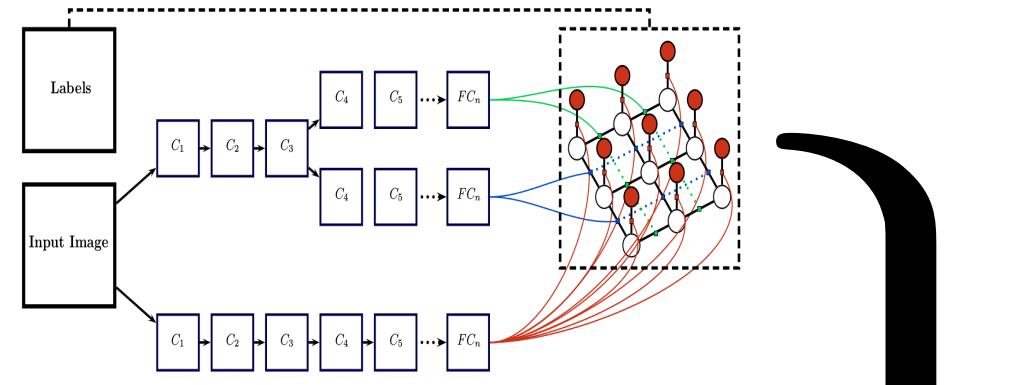
Jancsary, Nowozin, Sharp & Rother, Regression Tree Fields, CVPR12

Tappen, Liu, Adelson & Freeman, Learning Gaussian CRFs for low-level vision, CVPR07

# Deep Gaussian Conditional Random Field

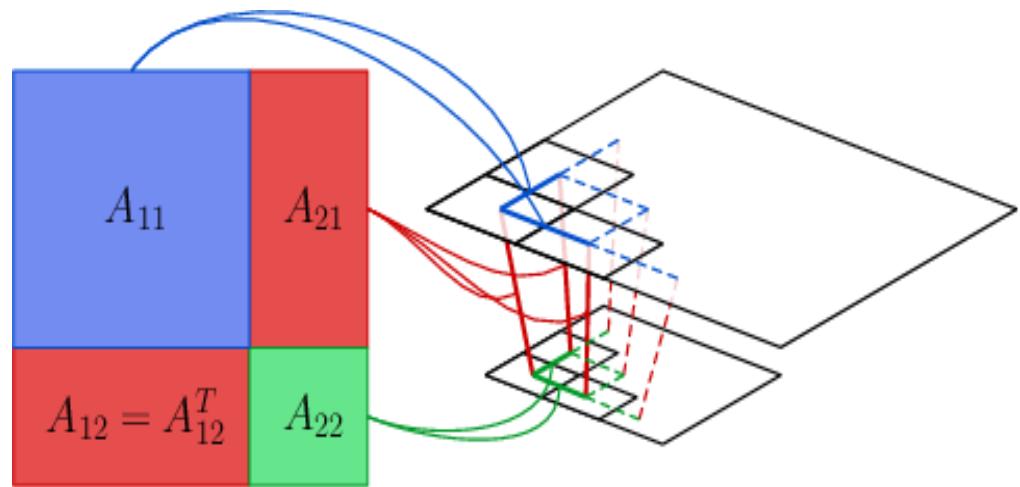


# Naïve Multi-Resolution Semantic Segmentation



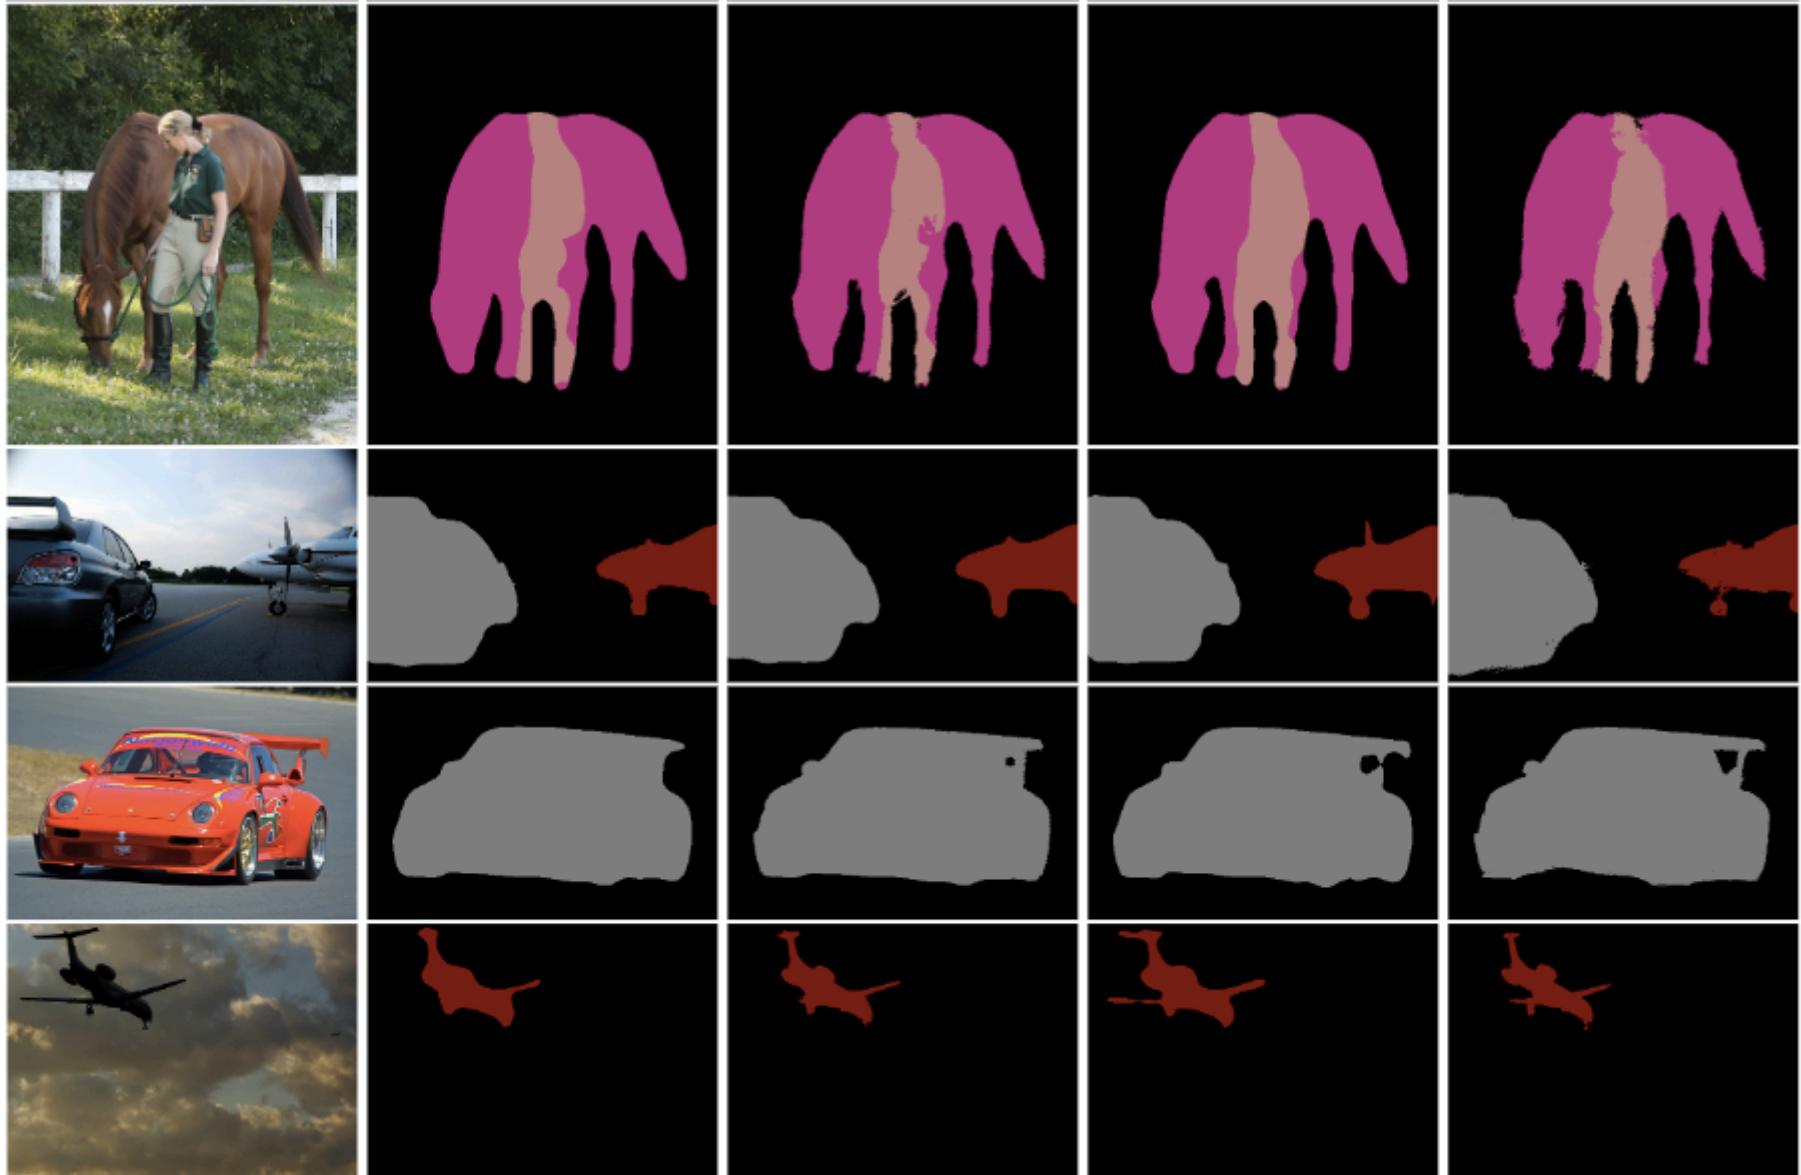
L.-C. Chen, Y. Yang, J. Wang, W. Xu and A. Yuille, ‘Attention to Scale: Scale-aware Semantic Image Segmentation, CVPR 2016  
I. Kokkinos, Pushing the Boundaries of Boundary Detection using Deep Learning, ICLR 2016

# Linear systems & Multi-resolution CRFs



**Learn to enforce coupling of different results  
Consistently better results than decoupled learning!**

# Improvements/Complementarity with DenseCRF



**FCNN**

**+DCRF**

**+GCRF**

**+DCRF + GCRF**

# Quantitative Results

Method	IoU	IoU after <i>dense CRF</i>
Basenet	72.72	73.78
QO <sub>4</sub>	73.41	75.13
QO <sub>4</sub> <sup>mres</sup>	73.86	75.46

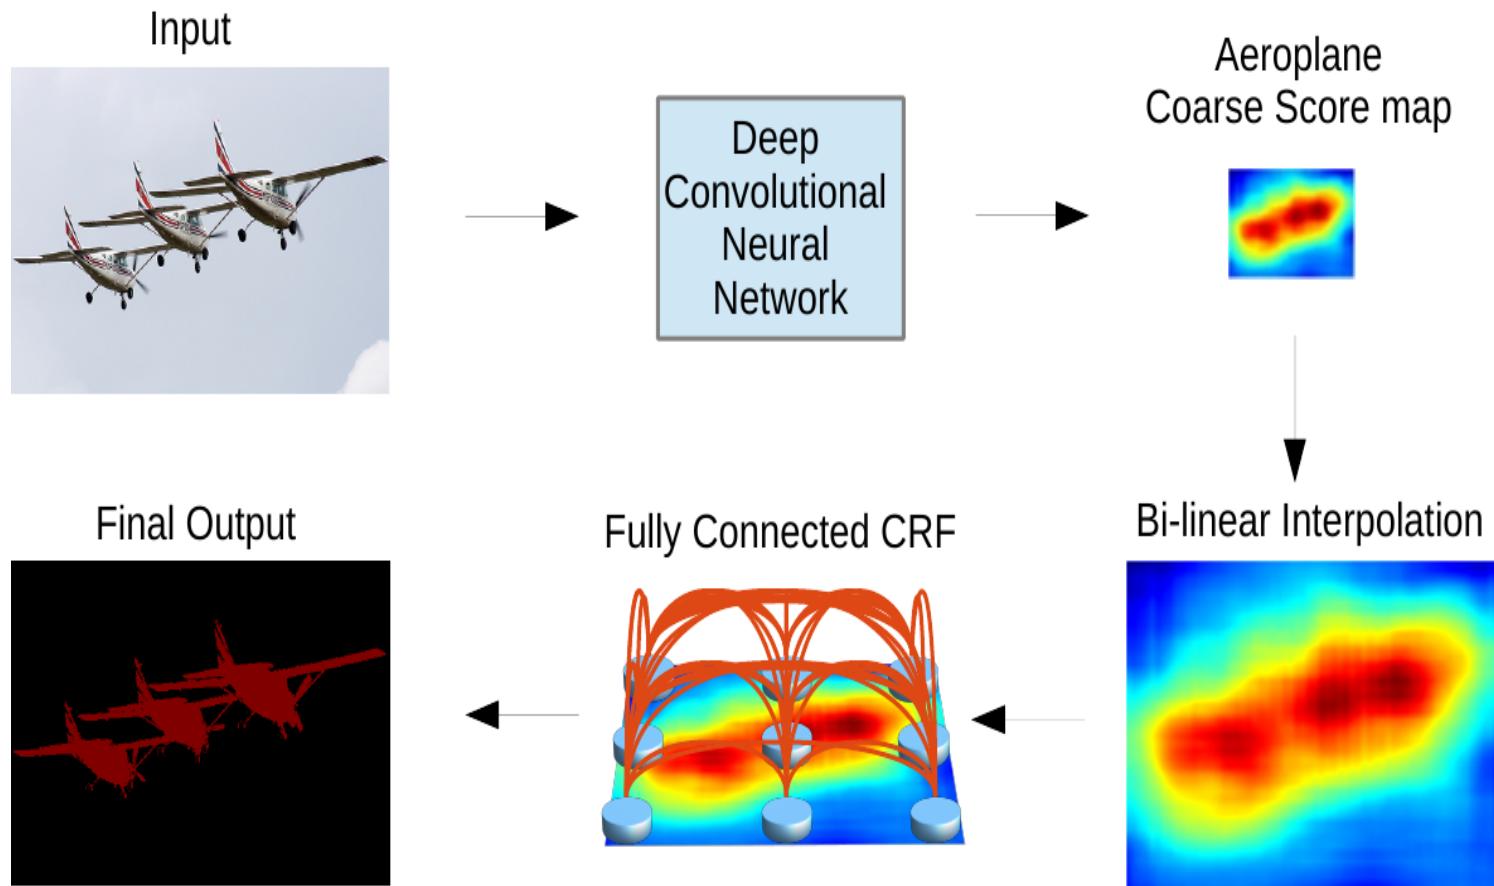
## Pascal Leaderboard:

<http://host.robots.ox.ac.uk:8080/leaderboard/displaylb.php?challengeid=11&compid=6>

	mean	aero plane	bicycle	bird	boat	bottle	bus	car	cat	chair	cow	dining table	dog	horse	motor bike	person	potted plant	sheep	sofa	train	tv/ monitor	submission date
RRR-ResNet152-COCO-MultiScale [?]	81.9	94.4	67.0	92.8	72.4	78.3	93.9	89.6	92.3	40.2	88.1	76.1	86.2	91.6	87.7	87.3	69.3	88.7	66.1	87.1	75.7	13-Sep-2016
SegModel [?]	81.8	93.5	60.2	92.6	60.1	75.4	86.3	88.2	95.5	37.0	90.8	73.3	91.1	94.2	88.6	88.6	64.8	90.1	63.7	87.3	78.7	23-Aug-2016
CentraleSupélec Deep G-CRF [?]	80.2	92.9	61.2	91.0	66.3	77.7	95.3	88.9	92.4	33.8	88.4	69.1	89.8	92.9	87.7	87.5	62.6	89.9	59.2	87.1	74.2	12-Aug-2016
CMT_FCN_BerNet_CRF [?]	80.0	92.5	55.3	92.2	66.0	76.9	95.1	88.6	93.9	35.1	87.6	71.6	89.3	92.8	87.9	88.0	62.0	88.0	59.7	86.1	75.7	02-Aug-2016
DeepLabv2-CRF [?]	79.7	92.6	60.4	91.6	63.4	76.3	95.0	88.4	92.6	32.7	88.5	67.6	89.6	92.1	87.0	87.4	63.3	88.3	60.0	86.8	74.5	06-Jun-2016
CASIA_SegResNet_CRF_LULU [?]	79.5	93.0	42.2	93.1	68.0	73.5	93.5	88.8	92.5	30.5	87.3	64.2	88.8	97.0	87.5	88.5	63.2	89.7	54.1	88.8	77.0	03-Jun-2016
LRR_4x_ResNet_COCO [?]	79.3	92.4	45.1	94.6	65.2	75.8	95.1	89.1	92.3	39.0	85.7	70.4	88.6	89.4	88.6	86.6	65.8	86.2	57.4	85.7	77.3	18-Jul-2016
Adelaide_VeryDeep_FCN_VOC [?]	79.1	91.9	48.1	93.4	69.3	75.5	94.2	87.5	92.8	36.7	86.9	65.2	89.1	90.2	86.5	87.2	64.6	90.1	59.7	85.5	72.7	13-May-2016
LRR_4x_COCO [?]	78.7	93.2	44.2	89.4	65.4	74.9	93.9	87.0	92.0	42.9	83.7	68.9	86.5	88.0	89.0	87.2	67.3	85.6	64.0	84.1	71.5	16-Jun-2016



# FCNN-DenseCRF: Accurate & Sharp



P. Krähenbühl and V. Koltun, Efficient Inference in Fully Connected CRFs with Gaussian Edge Potentials, NIPS 2011

L.-C. Chen, G. Papandreou, I. Kokkinos, K. Murphy and A. Yuille, Semantic Image Segmentation with Deep Convolutional Nets and Fully Connected CRFs, ICLR 2015

# Dense CRF: smart choice of pairwise term

$$\begin{aligned}\psi_{i,j}(l, l') &= \mu(l, l') \sum_{m=1}^M w_m k_m(\mathbf{f}_i, \mathbf{f}_j) \\ &= [l \neq l'] \left[ w_1 \exp \left( -\frac{\|p_i - p_j\|^2}{2\sigma_a^2} - \frac{\|I_i - I_j\|^2}{2\sigma_b^2} \right) + w_2 \exp \left( -\frac{\|p_i - p_j\|^2}{2\sigma_\gamma^2} \right) \right]\end{aligned}$$

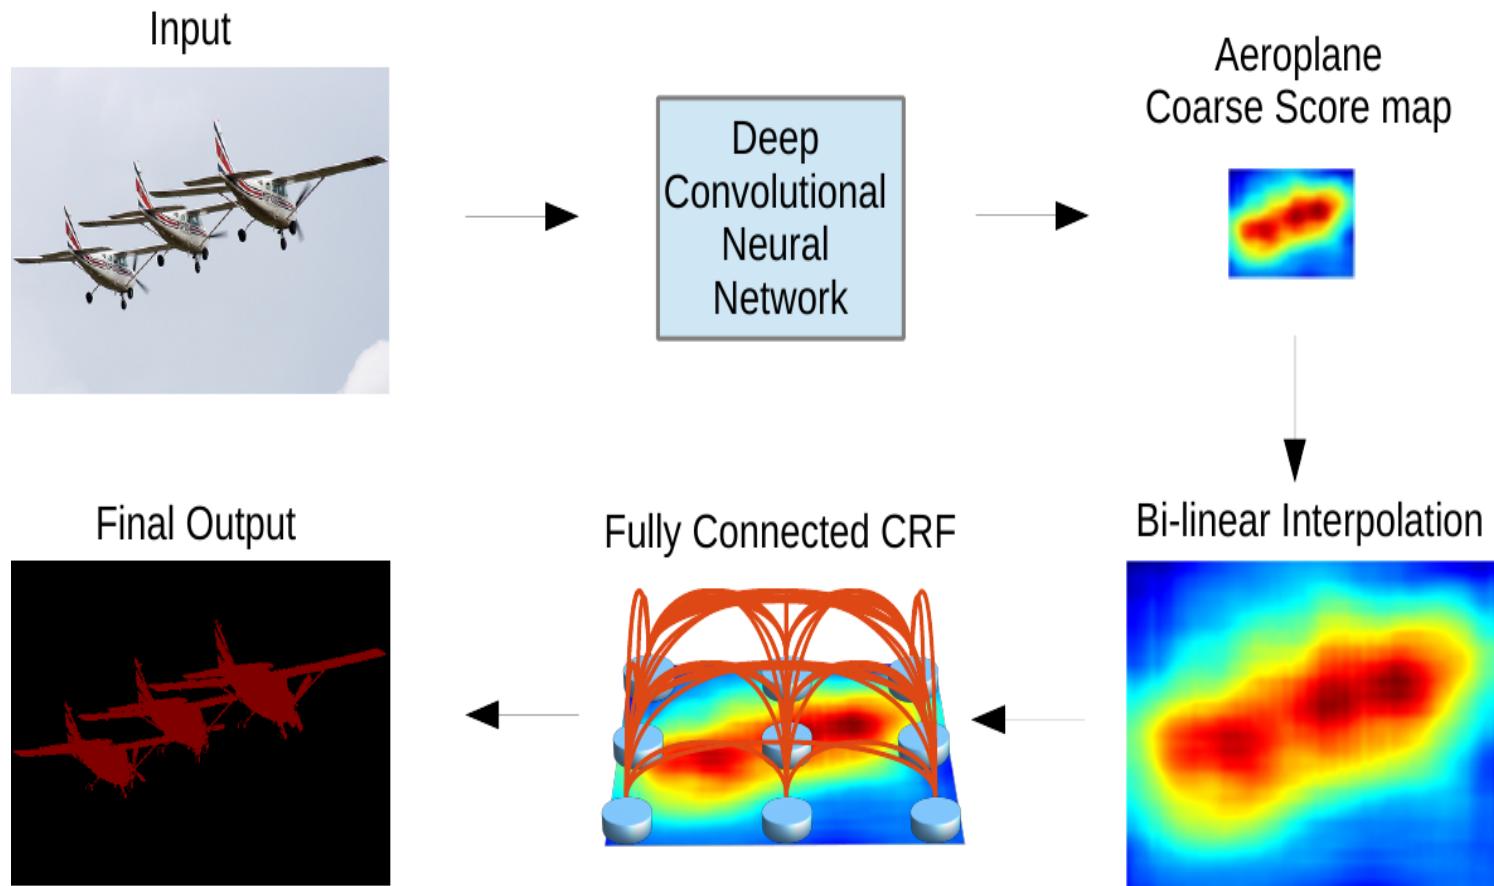
Potts model                  ‘Bilateral kernel’                  Spatial proximity

Mean Field Updates:

$$Q_i(l) = \frac{1}{Z_i} \exp \left\{ -\psi_i(l) - \sum_{l'} \mu(l, l') \sum_{m=1}^M w_m \sum_{j \in \mathcal{N}_i} k_m(\mathbf{f}_i, \mathbf{f}_j) Q_j(l') \right\}$$

Efficient high-dimensional convolutions using the Permutohedral Lattice

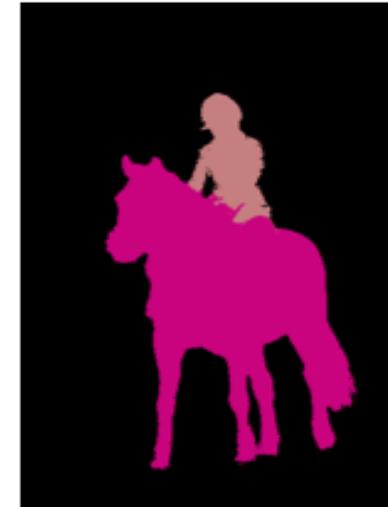
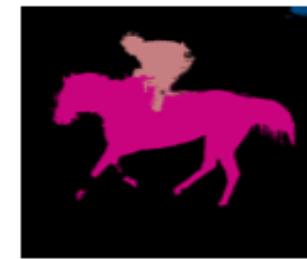
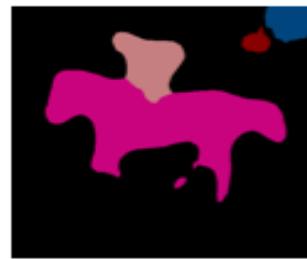
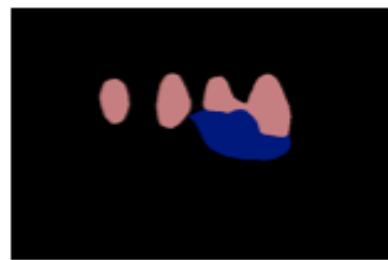
# FCNN-DenseCRF: Accurate & Sharp



P. Krähenbühl and V. Koltun, Efficient Inference in Fully Connected CRFs with Gaussian Edge Potentials, NIPS 2011

L.-C. Chen, G. Papandreou, I. Kokkinos, K. Murphy and A. Yuille, Semantic Image Segmentation with Deep Convolutional Nets and Fully Connected CRFs, ICLR 2015

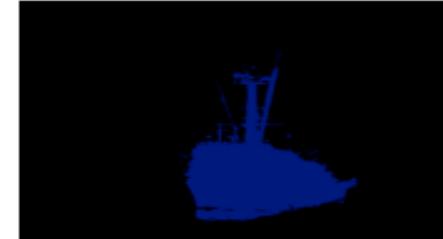
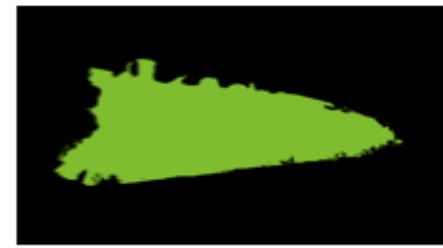
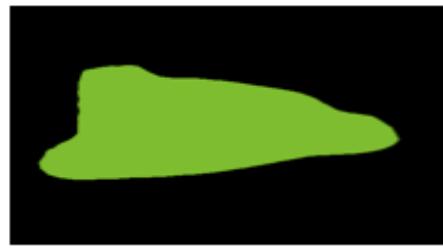
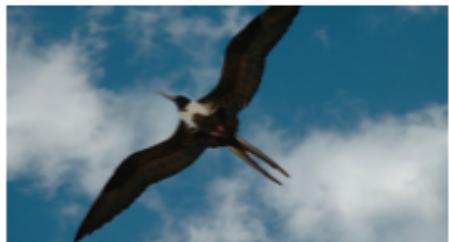
# Qualitative Results



FCNN

FCNN-DCRF

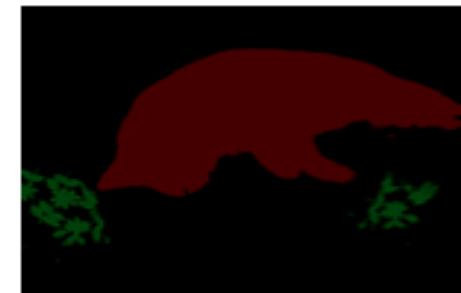
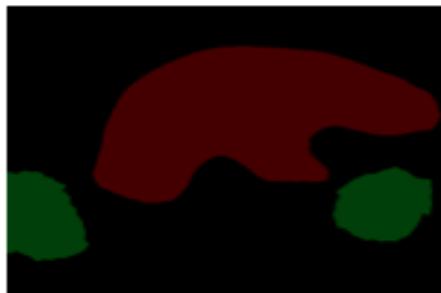
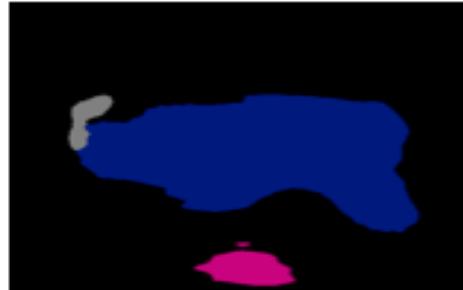
# Qualtiative Results



FCNN

FCNN-DCRF

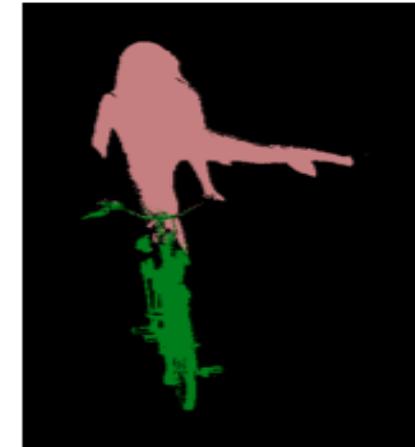
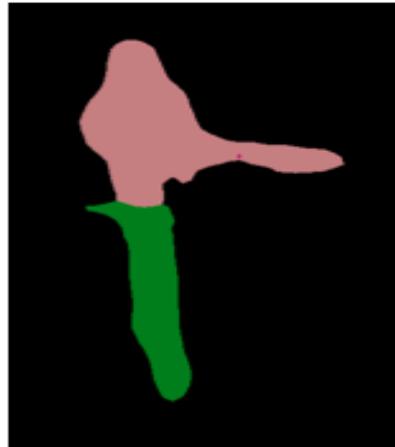
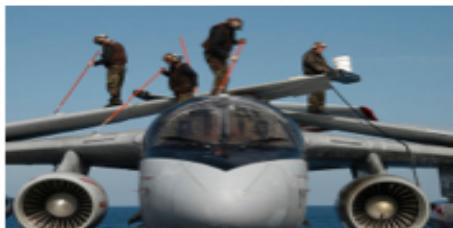
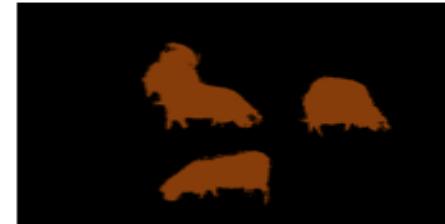
# Qualitative Results



FCNN

FCNN-DCRF

# Indicative Results

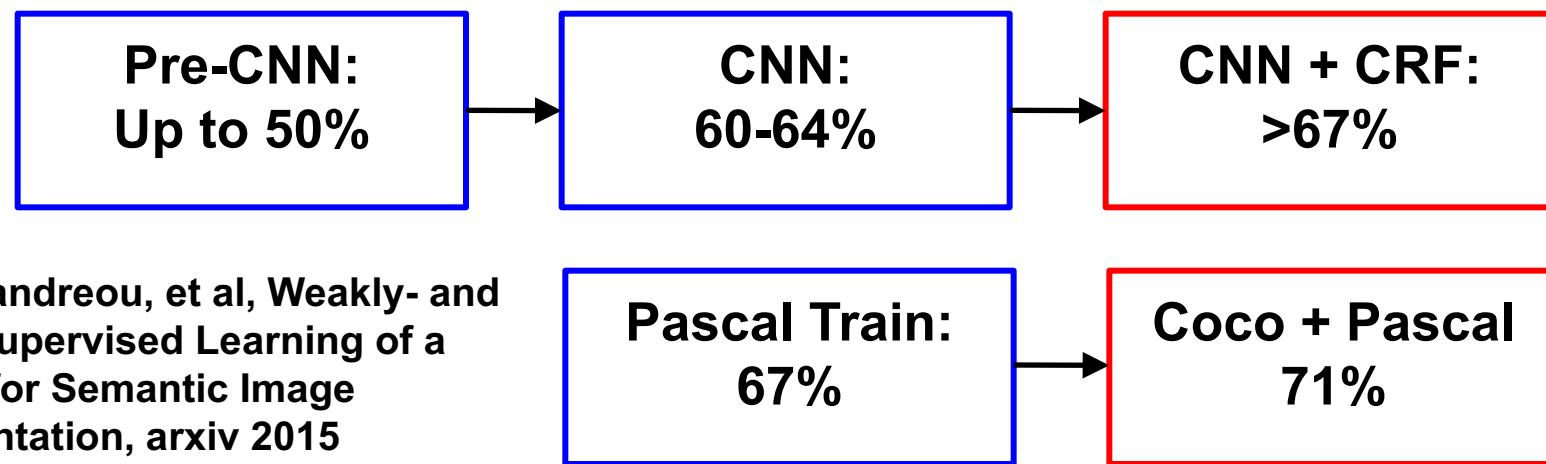


FCNN

FCNN-DCRF

# Comparison to state-of-the-art, 2014

Method	mean IOU (%)
MSRA-CFM	61.8
FCN-8s	62.2
TTI-Zoomout-16	64.4
DeepLab-CRF (our)	66.4
DeepLab-MSc-CRF (our)	67.1



G. Papandreou, et al, Weakly- and Semi-Supervised Learning of a DCNN for Semantic Image Segmentation, arxiv 2015

74.7 end-to-end S. Zheng, et al. CRFs as recurrent neural networks. In ICCV, 2015.

# Updated results (Deeplab v2)

Method	mIOU
DeepLab-CRF-LargeFOV-COCO [58]	72.7
MERL_DEEP_GCRF [88]	73.2
CRF-RNN [59]	74.7
POSTECH_DeconvNet_CRF_VOC [61]	74.8
BoxSup [60]	75.2
Context + CRF-RNN [74]	75.3
$QO_4^{mres}$ [66]	75.5
DeepLab-CRF-Attention [17]	75.7
CentraleSuperBoundaries++ [18]	76.0
DeepLab-CRF-Attention-DT [63]	76.3
H-ReNet + DenseCRF [89]	76.8
LRR_4x_COCO [90]	76.8
DPN [62]	77.5
Adelaide_Context [40]	77.8
Oxford_TVG_HO_CRF [87]	77.9
Context CRF + Guidance CRF [91]	78.1
Adelaide_VeryDeep_FCN_VOC [92]	79.1
DeepLab-CRF (ResNet-101)	79.3
Ensemble DeepLab-CRF (ResNet-101)	79.6

TABLE 5: Performance on PASCAL VOC 2012 *test* set. We have added some results from recent arXiv papers on top of the official leadearboard results.

# Updated results (Deeplab v2)

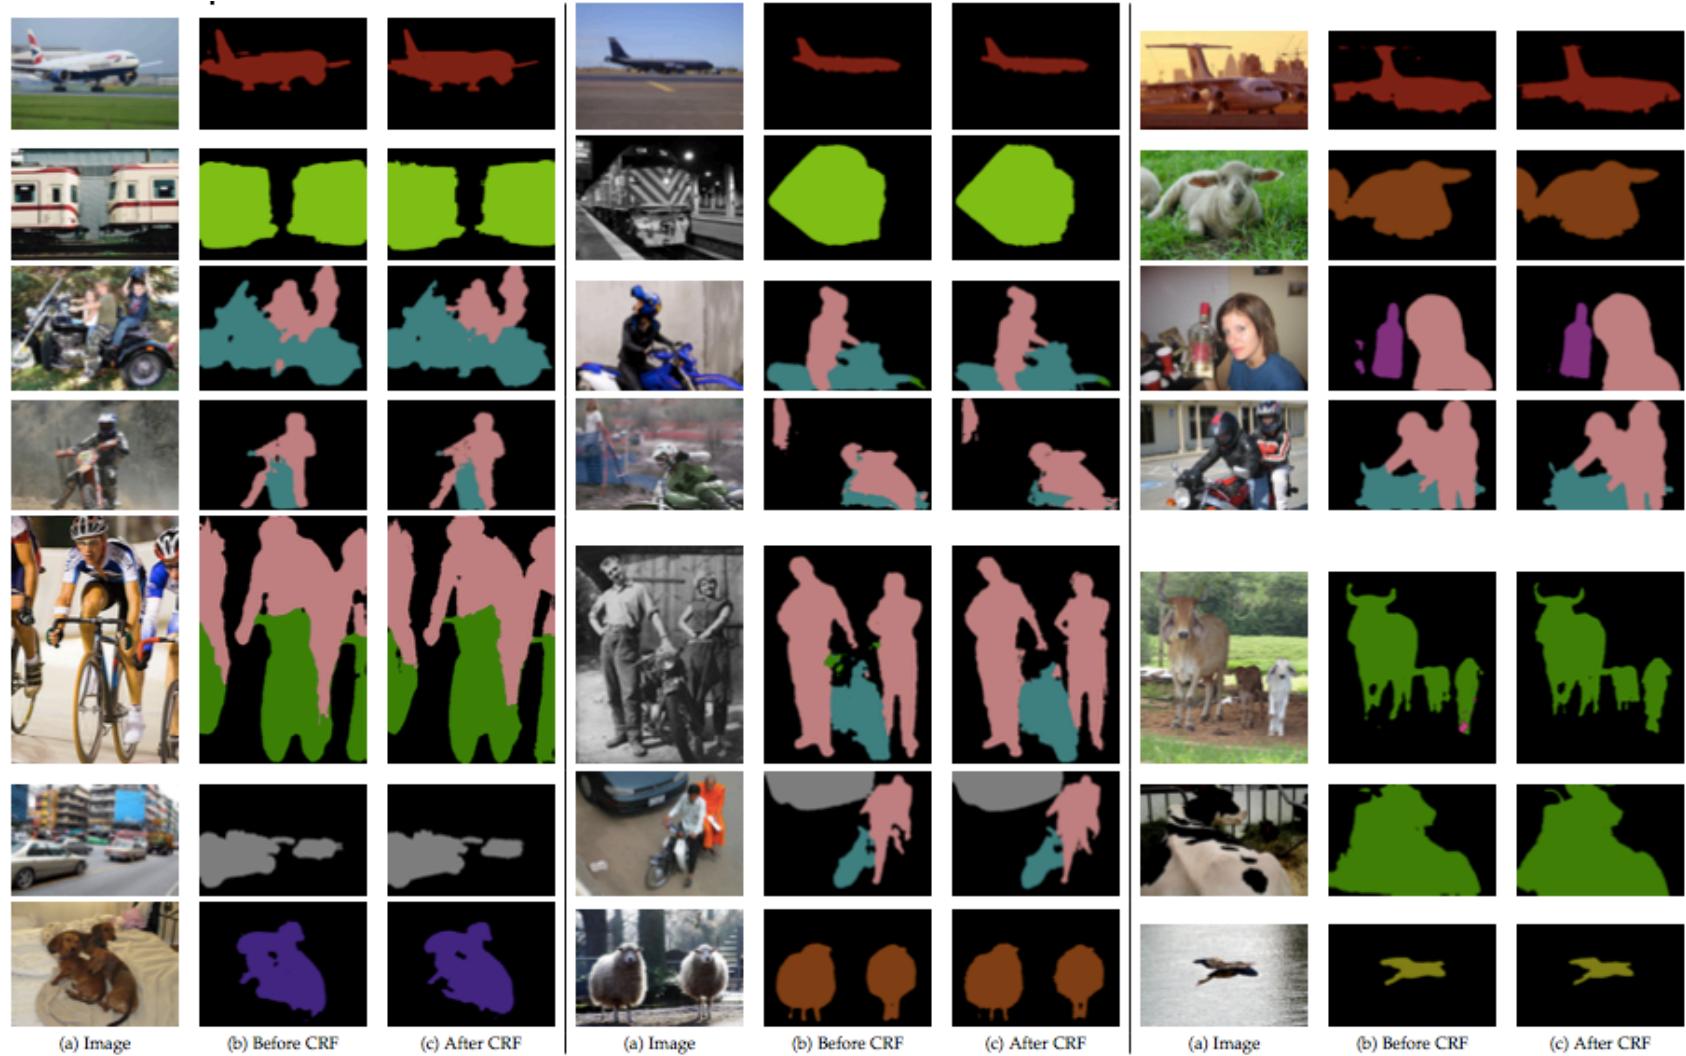
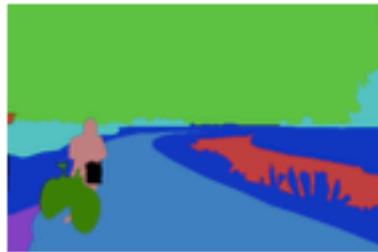


Fig. 7: Visualization of some VOC 2012 *val* results. For each row, we show the input image, the segmentation result before CRF, and the refined segmentation result after Fully Connected CRF (DeepLab-CRF).

# Updated results (Deeplab v2)

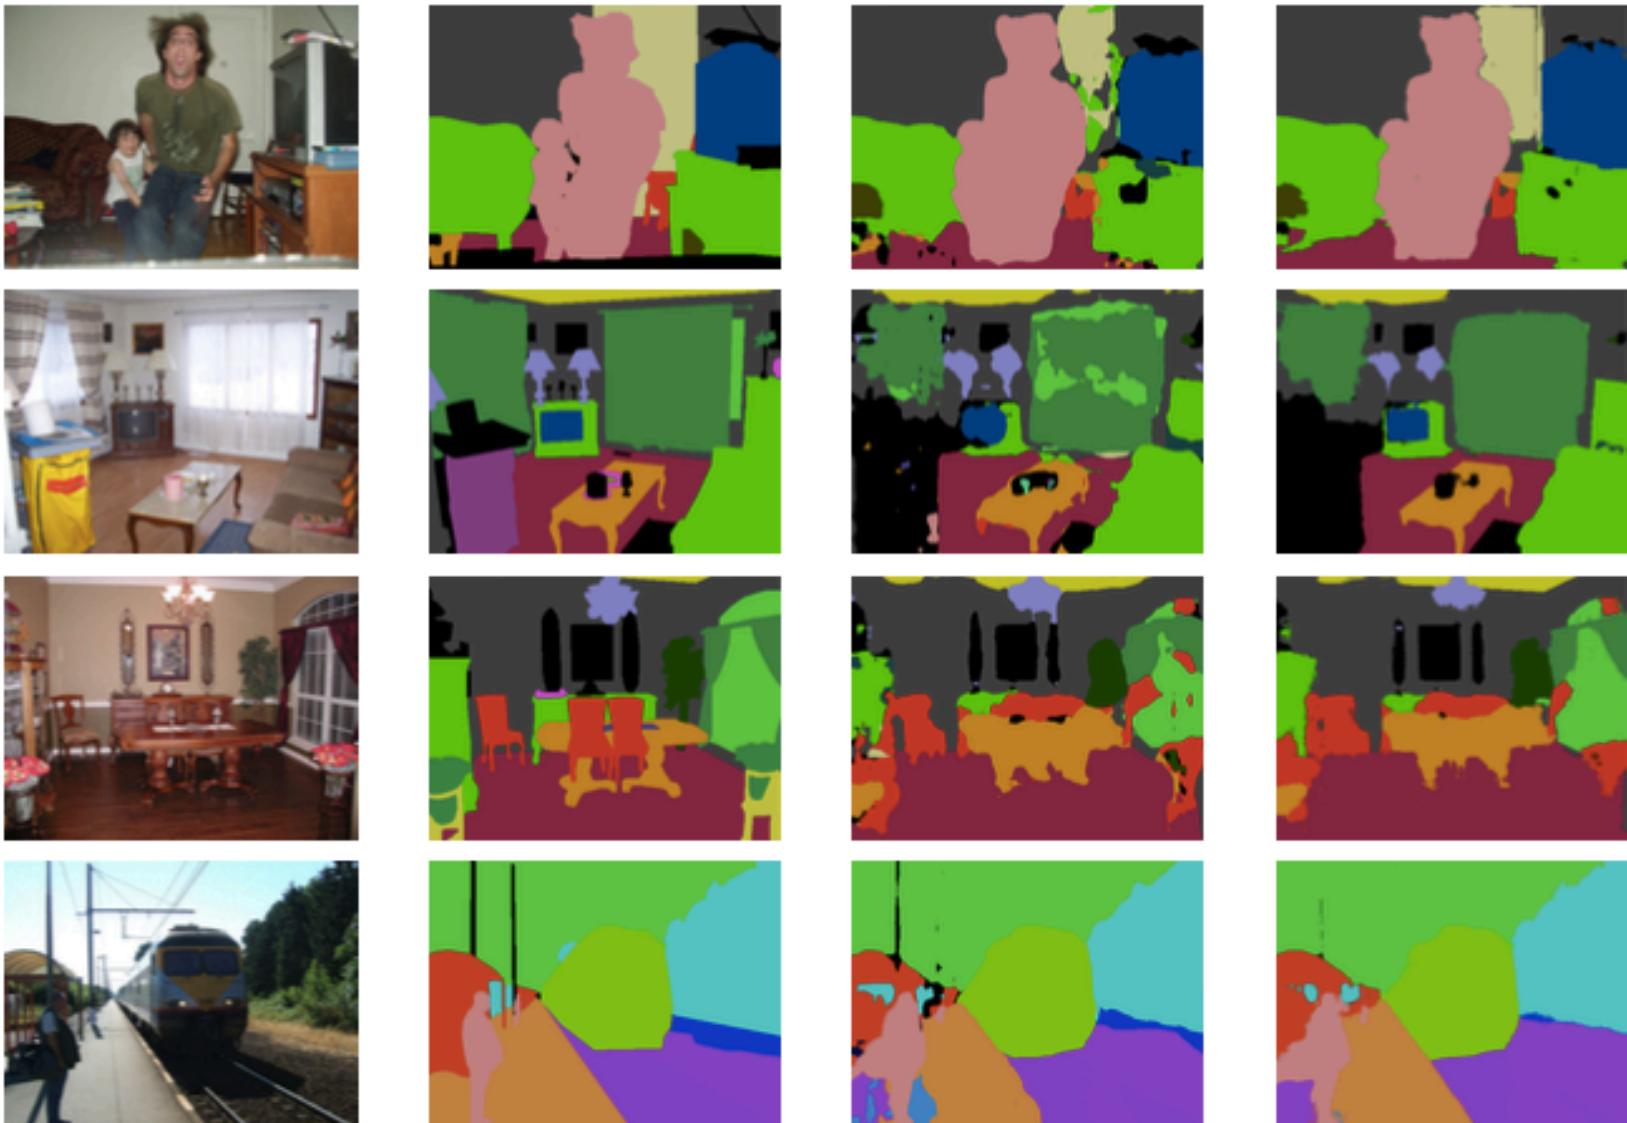


Ground truth

FCNN

FCNN-DCRF

# Updated results (Deeplab v2)



Ground truth

FCNN

FCNN-DCRF

# Updated results (Deeplab v2)

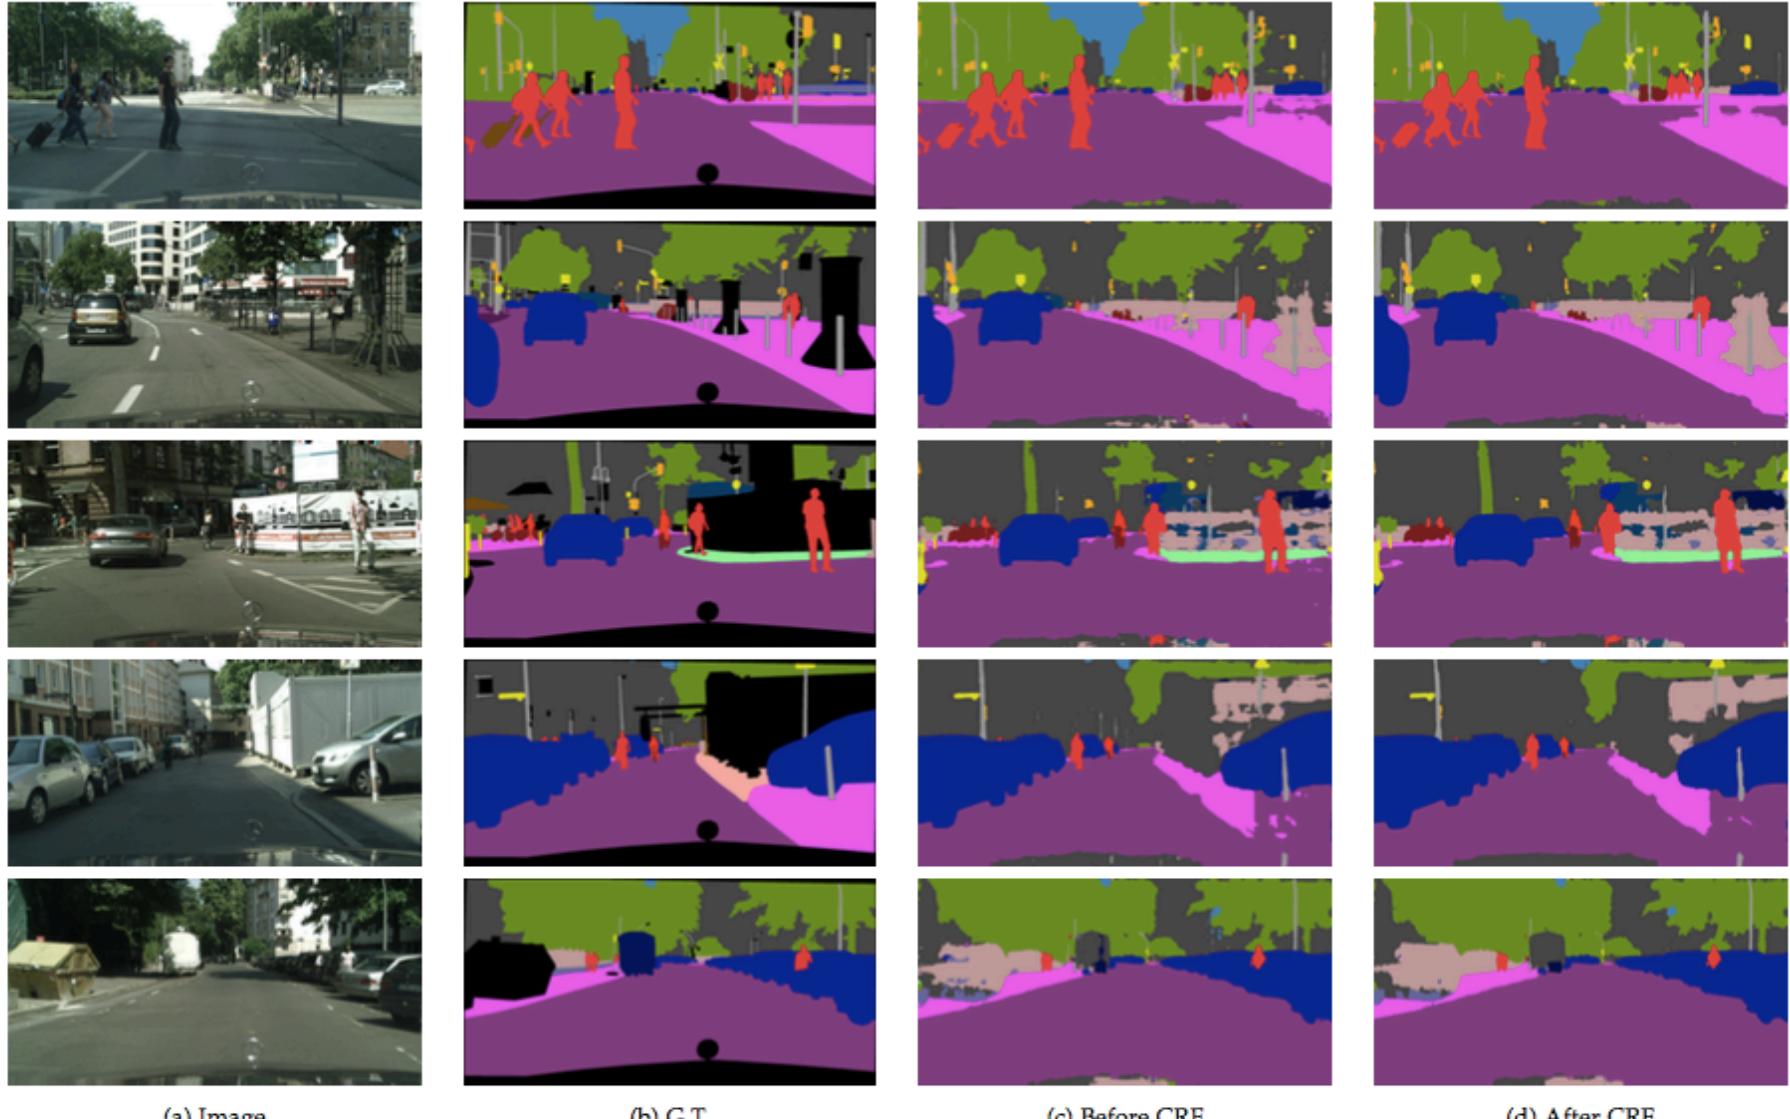
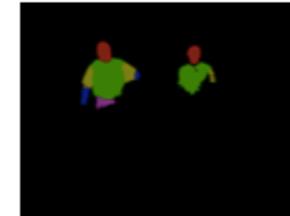
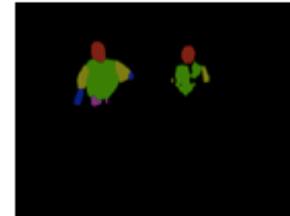
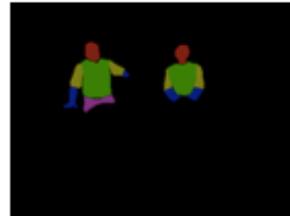


Fig. 11: Visualization results on Cityscapes. For each row, we show the input image, the ground-truth, and our DeepLab results before and after CRF.

# Updated results (Deeplab v2)



See also: S. Tsogkas, G. Papandreou, I. Kokkinos, and A. Vedaldi, Semantic Part Segmentation using high-level guidance, Arxiv, 2015





# Steps of Tracking

: predict next state of the object given past measurements

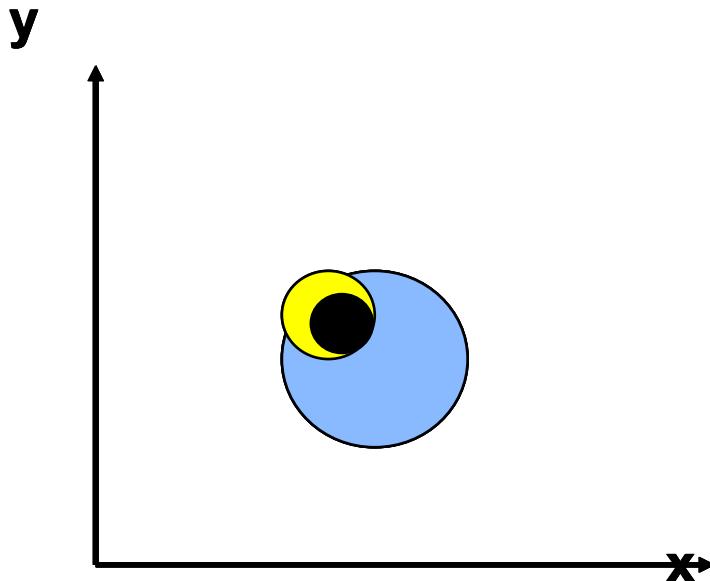
$$P(X_t | Y_0 = y_0, \dots, Y_{t-1} = y_{t-1})$$

: combine prediction and current measurement to yield new distribution.

$$P(X_t | Y_0 = y_0, \dots, Y_{t-1} = y_{t-1}, Y_t = y_t)$$

# Tracking: from t to t+1

**Initial position**  
**Prediction**  
**Measurement**  
**Update**



# Induction Step: Prediction

- Prediction involves representing given  $P(X_{t-1} | y_0, \dots, y_{t-1})$

$$P(X_t | y_0, \dots, y_{t-1})$$

$$P(X_t | y_0, \dots, y_{t-1})$$

$$= \int P(X_t, X_{t-1} | y_0, \dots, y_{t-1}) dX_{t-1}$$

**Law of total probability**

$$P(A) = \int P(A, B) dB$$

# Induction Step: Prediction

- Prediction involves representing given  $P(X_{t-1} | y_0, \dots, y_{t-1})$

$$P(X_t | y_0, \dots, y_{t-1})$$

$$P(X_t | y_0, \dots, y_{t-1})$$

$$= \int P(X_t, X_{t-1} | y_0, \dots, y_{t-1}) dX_{t-1}$$

$$= \int P(X_t | X_{t-1}, y_0, \dots, y_{t-1}) P(X_{t-1} | y_0, \dots, y_{t-1}) dX_{t-1}$$

**Conditioning on  $X_{t-1}$**

$$P(A, B) = P(A | B) P(B)$$

# Induction Step: Prediction

- Prediction involves representing given  $P(X_{t-1} | y_0, \dots, y_{t-1})$

$$P(X_t | y_0, \dots, y_{t-1})$$

$$P(X_t | y_0, \dots, y_{t-1})$$

$$= \int P(X_t, X_{t-1} | y_0, \dots, y_{t-1}) dX_{t-1}$$

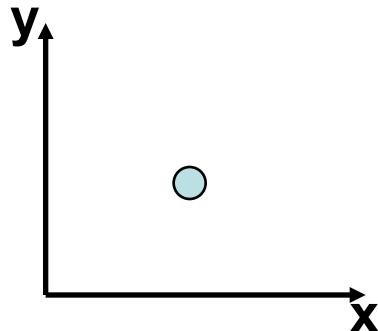
$$= \int P(X_t | X_{t-1}, y_0, \dots, y_{t-1}) P(X_{t-1} | y_0, \dots, y_{t-1}) dX_{t-1}$$

$$= \int P(X_t | X_{t-1}) P(X_{t-1} | y_0, \dots, y_{t-1}) dX_{t-1}$$

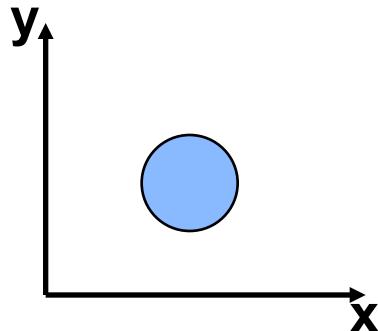
**Independence assumption**

# What we just described:

**Initial position**



**Prediction**



# Induction Step: Correction

- Correction involves computing given predicted value

$$\frac{P(X_t | y_0, \dots, y_t)}{P(X_t | y_0, \dots, y_{t-1})}$$

$$P(X_t | y_0, \dots, y_t)$$

$$= \frac{P(y_t | X_t, y_0, \dots, y_{t-1}) P(X_t | y_0, \dots, y_{t-1})}{P(y_t | y_0, \dots, y_{t-1})}$$

**Bayes rule**

$$P(A | B) = \frac{P(B | A) P(A)}{P(B)}$$

# Induction Step: Correction

- Correction involves computing given predicted value

$$\frac{P(X_t | y_0, \dots, y_t)}{P(X_t | y_0, \dots, y_{t-1})}$$

$$P(X_t | y_0, \dots, y_t)$$

$$= \frac{P(y_t | X_t, y_0, \dots, y_{t-1}) P(X_t | y_0, \dots, y_{t-1})}{P(y_t | y_0, \dots, y_{t-1})}$$

$$= \frac{P(y_t | X_t) P(X_t | y_0, \dots, y_{t-1})}{P(y_t | y_0, \dots, y_{t-1})}$$

Independence assumption  
(observation  $y_t$  depends only on state  $X_t$ )  
102

# Induction Step: Correction

- Correction involves computing given predicted value

$$\frac{P(X_t | y_0, \dots, y_t)}{P(X_t | y_0, \dots, y_{t-1})}$$

$$P(X_t | y_0, \dots, y_t)$$

$$= \frac{P(y_t | X_t, y_0, \dots, y_{t-1}) P(X_t | y_0, \dots, y_{t-1})}{P(y_t | y_0, \dots, y_{t-1})}$$

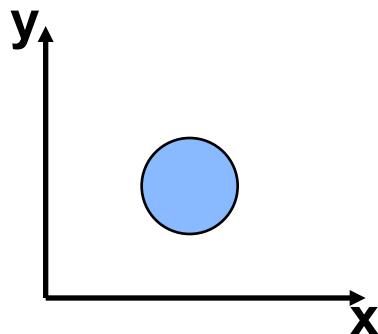
$$= \frac{P(y_t | X_t) P(X_t | y_0, \dots, y_{t-1})}{P(y_t | y_0, \dots, y_{t-1})}$$

$$= \frac{P(y_t | X_t) P(X_t | y_0, \dots, y_{t-1})}{\int P(y_t | X_t) P(X_t | y_0, \dots, y_{t-1}) dX_t}$$

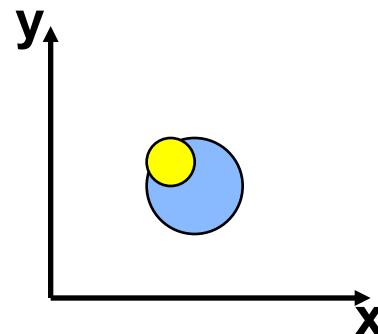
**Conditioning on  $X_t$**

# What we just described

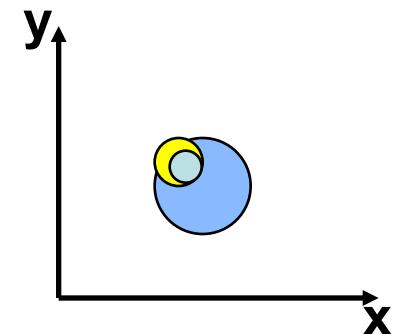
Prediction



Measurement



Update



# Summary: Prediction and Correction

- Prediction:

$$P(X_t | y_0, \dots, y_{t-1}) = \int \underbrace{P(X_t | X_{t-1})}_{\text{Dynamics model}} \underbrace{P(X_{t-1} | y_0, \dots, y_{t-1})}_{\text{Corrected estimate from previous step}} dX_{t-1}$$

# Summary: Prediction and Correction

- Prediction:

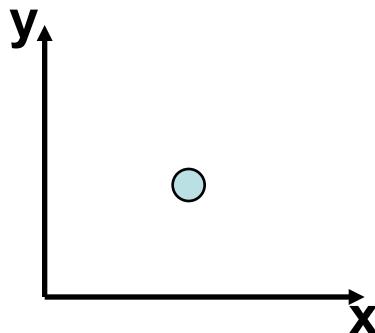
$$P(X_t | y_0, \dots, y_{t-1}) = \int \underbrace{P(X_t | X_{t-1})}_{\text{Dynamics model}} \underbrace{P(X_{t-1} | y_0, \dots, y_{t-1})}_{\text{Corrected estimate from previous step}} dX_{t-1}$$

- Correction:

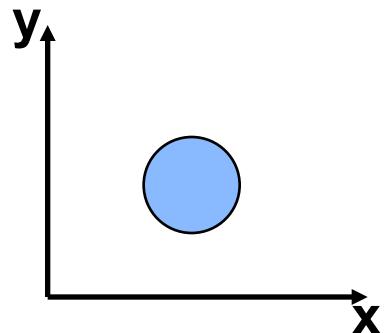
$$P(X_t | y_0, \dots, y_t) = \frac{\underbrace{P(y_t | X_t)}_{\text{Observation model}} \underbrace{P(X_t | y_0, \dots, y_{t-1})}_{\text{Predicted estimate}}}{\int P(y_t | X_t) P(X_t | y_0, \dots, y_{t-1}) dX_t}$$

# Predict-measure-update

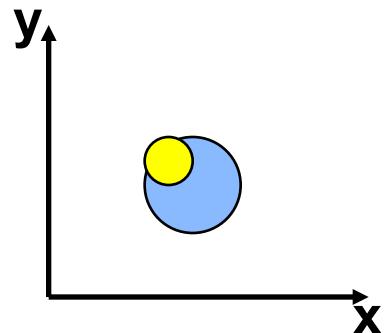
Initial position



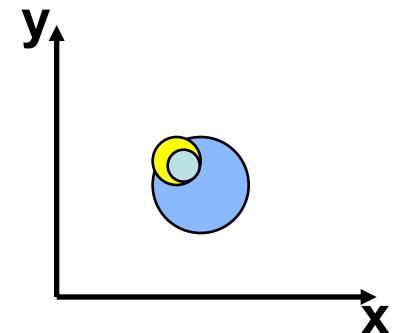
Prediction



Measurement

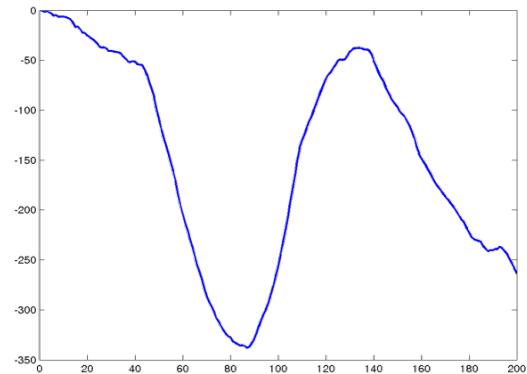


Update



# This lecture

- Tracking
  - Linear Dynamical Models



$$\begin{bmatrix} x_{t+1} \\ \delta_{t+1} \end{bmatrix} = \begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} x_t \\ \delta_t \end{bmatrix} + w_t$$

# Linear Dynamical Models

- Linear Dynamics Assumption

$$x_t = \underbrace{D_t}_{N \times N} x_{t-1} + e_d, \quad e_d \sim N(0, \Sigma_d)$$

- Linear Appearance Assumption

$$y_t = \underbrace{M_t}_{M \times N} x_t + e_m, \quad e_m \sim N(0, \Sigma_m)$$

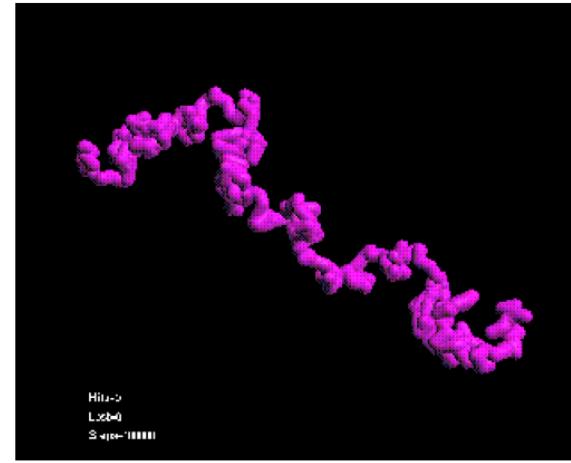
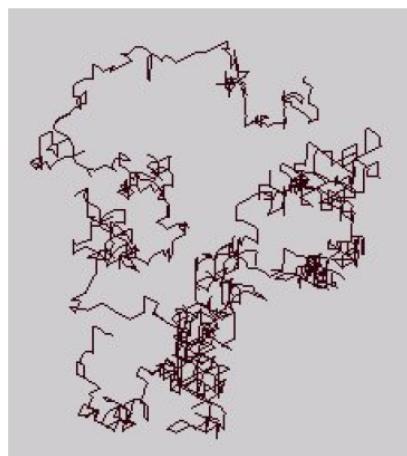
# Example: Randomly Drifting Points

- State: position of object
  - Only motion: random noise term.

$$x_t = p_t \quad p_t = p_{t-1} + \varepsilon$$

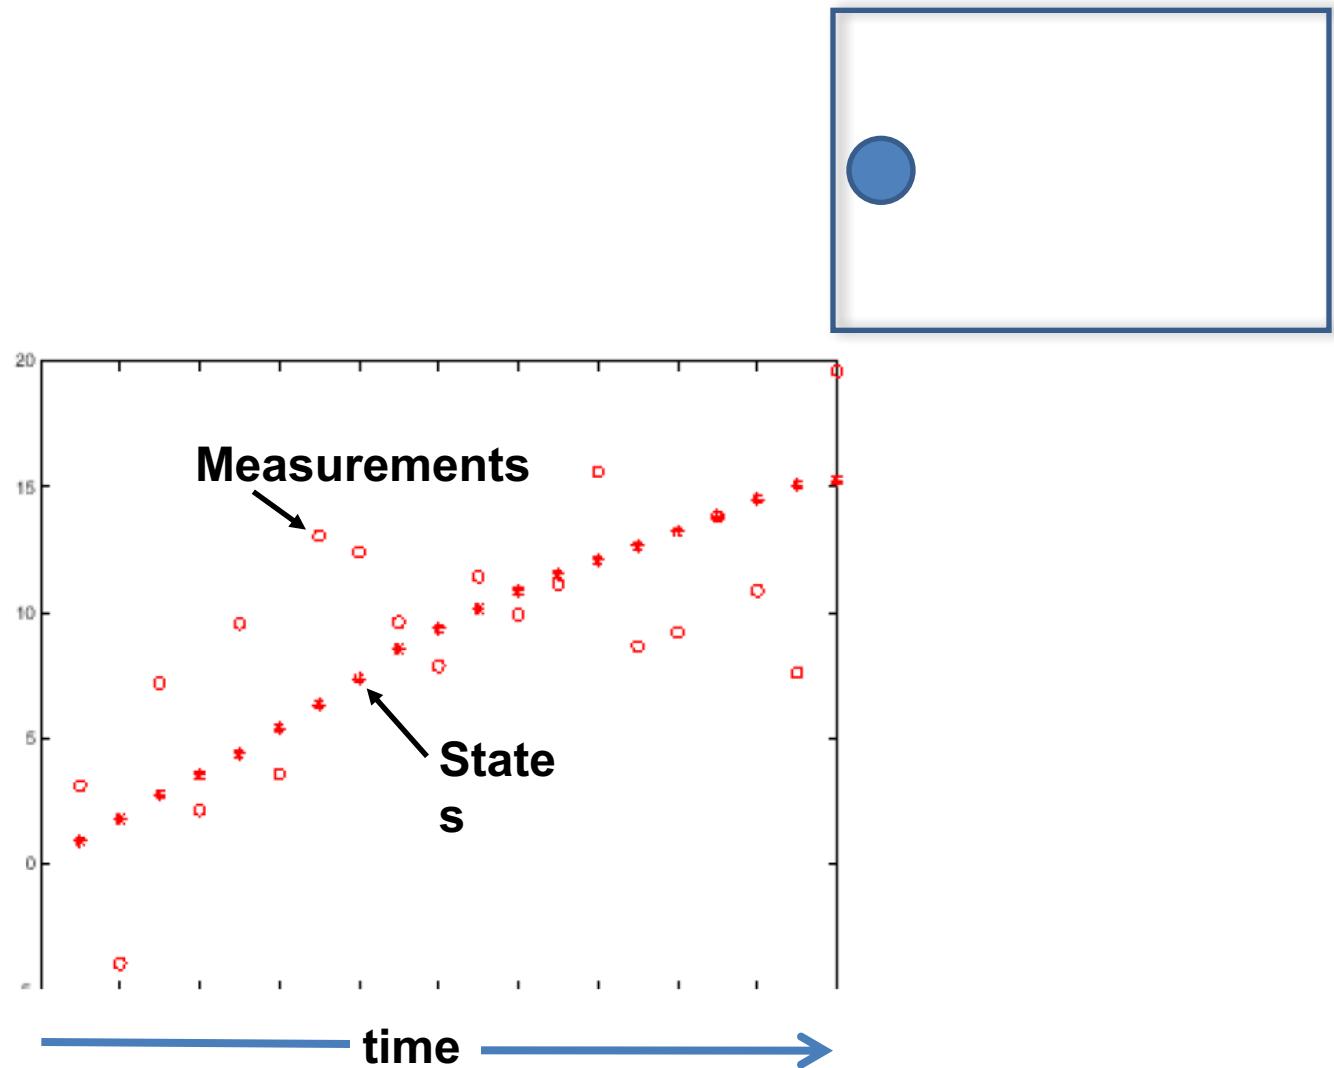
⇒ State evolution is described by identity matrix D=I

$$x_t = D_t x_{t-1} + \text{noise} = I p_{t-1} + \text{noise}$$



cic.nist.gov/lipman/sciviz/images/random3.gif  
<http://www.grunch.net/synergetics/images/random3.jpg>

# Example: Constant Velocity (1D Points)



# Example: Constant Velocity (1D Points)

- State vector: position  $p$  and velocity  $v$

$$x_t = \begin{bmatrix} p_t \\ v_t \end{bmatrix} \quad p_t =$$

(greek letters  
denote noise  
terms)

$$x_t = D_t x_{t-1} + noise =$$

- Measurement is position only

$$y_t = M x_t + noise =$$

# Example: Constant Velocity (1D Points)

- State vector: position  $p$  and velocity  $v$

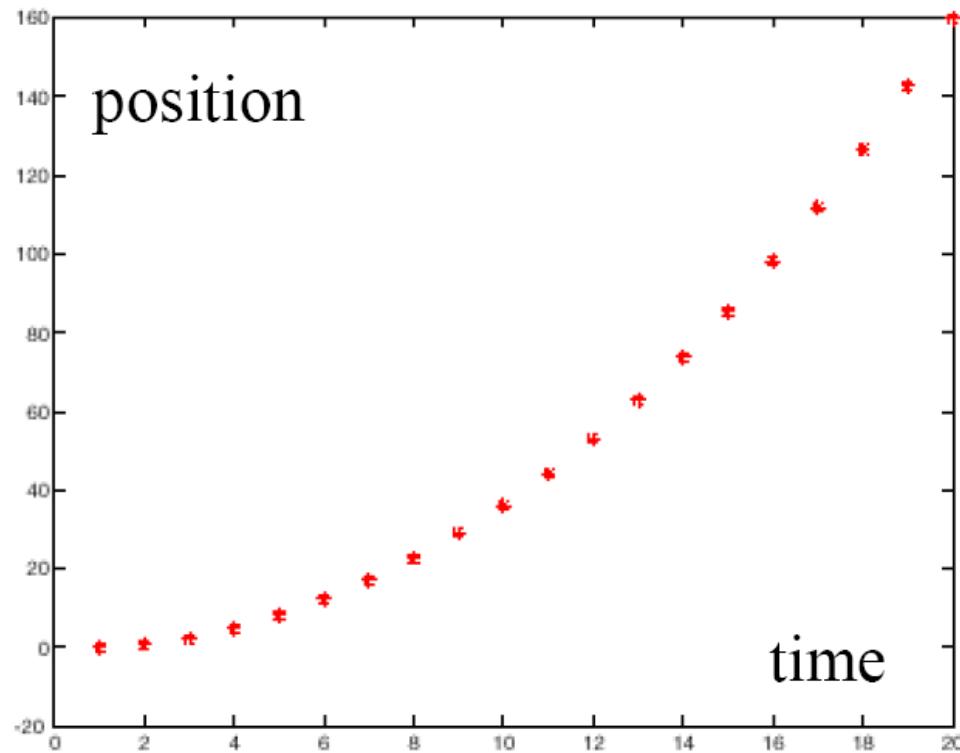
$$x_t = \begin{bmatrix} p_t \\ v_t \end{bmatrix} \quad p_t = p_{t-1} + (\Delta t)v_{t-1} + \varepsilon \quad (\text{greek letters denote noise terms})$$
$$v_t = v_{t-1} + \xi$$

$$x_t = D_t x_{t-1} + \text{noise} = \begin{bmatrix} 1 & \Delta t \\ 0 & 1 \end{bmatrix} \begin{bmatrix} p_{t-1} \\ v_{t-1} \end{bmatrix} + \text{noise}$$

- Measurement is position only

$$y_t = M x_t + \text{noise} = \begin{bmatrix} 1 & 0 \end{bmatrix} \begin{bmatrix} p_t \\ v_t \end{bmatrix} + \text{noise}$$

# Example: Constant Acceleration (1D Points)



# Example: Constant Acceleration (1D Points)

- State vector: position  $p$ , velocity  $v$ , and acceleration  $a$ .

$$x_t = \begin{bmatrix} p_t \\ v_t \\ a_t \end{bmatrix}$$

$$p_t = p_{t-1} + (\Delta t)v_{t-1} + \varepsilon$$

$$v_t =$$

$$a_t =$$

(greek letters  
denote noise  
terms)

$$x_t = D_t x_{t-1} + noise =$$

- Measurement is position

$$y_t = M x_t + noise =$$

# Example: Constant Acceleration (1D Points)

- State vector: position  $p$ , velocity  $v$ , and acceleration  $a$ .

$$x_t = \begin{bmatrix} p_t \\ v_t \\ a_t \end{bmatrix} \quad \begin{aligned} p_t &= p_{t-1} + (\Delta t)v_{t-1} + \varepsilon \\ v_t &= v_{t-1} + (\Delta t)a_{t-1} + \xi \\ a_t &= a_{t-1} + \zeta \end{aligned} \quad \text{(greek letters denote noise terms)}$$

$$x_t = D_t x_{t-1} + \text{noise} = \begin{bmatrix} 1 & \Delta t & 0 \\ 0 & 1 & \Delta t \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} p_{t-1} \\ v_{t-1} \\ a_{t-1} \end{bmatrix} + \text{noise}$$

- Measurement is position only

$$y_t = M x_t + \text{noise} = \begin{bmatrix} 1 & 0 & 0 \end{bmatrix} \begin{bmatrix} p_t \\ v_t \\ a_t \end{bmatrix} + \text{noise}$$

# General Motion Models

- Assuming we have differential equations for the motion
  - E.g. for (undamped) periodic motion of a pendulum

$$\frac{d^2 p}{dt^2} = -p$$

- Substitute variables to transform this into linear system

$$p_1 = p \quad p_2 = \frac{dp}{dt} \quad p_3 = \frac{d^2 p}{dt^2}$$

‘With a large enough state-space, the universe is Markov’ U. Grenander

# The Kalman Filter

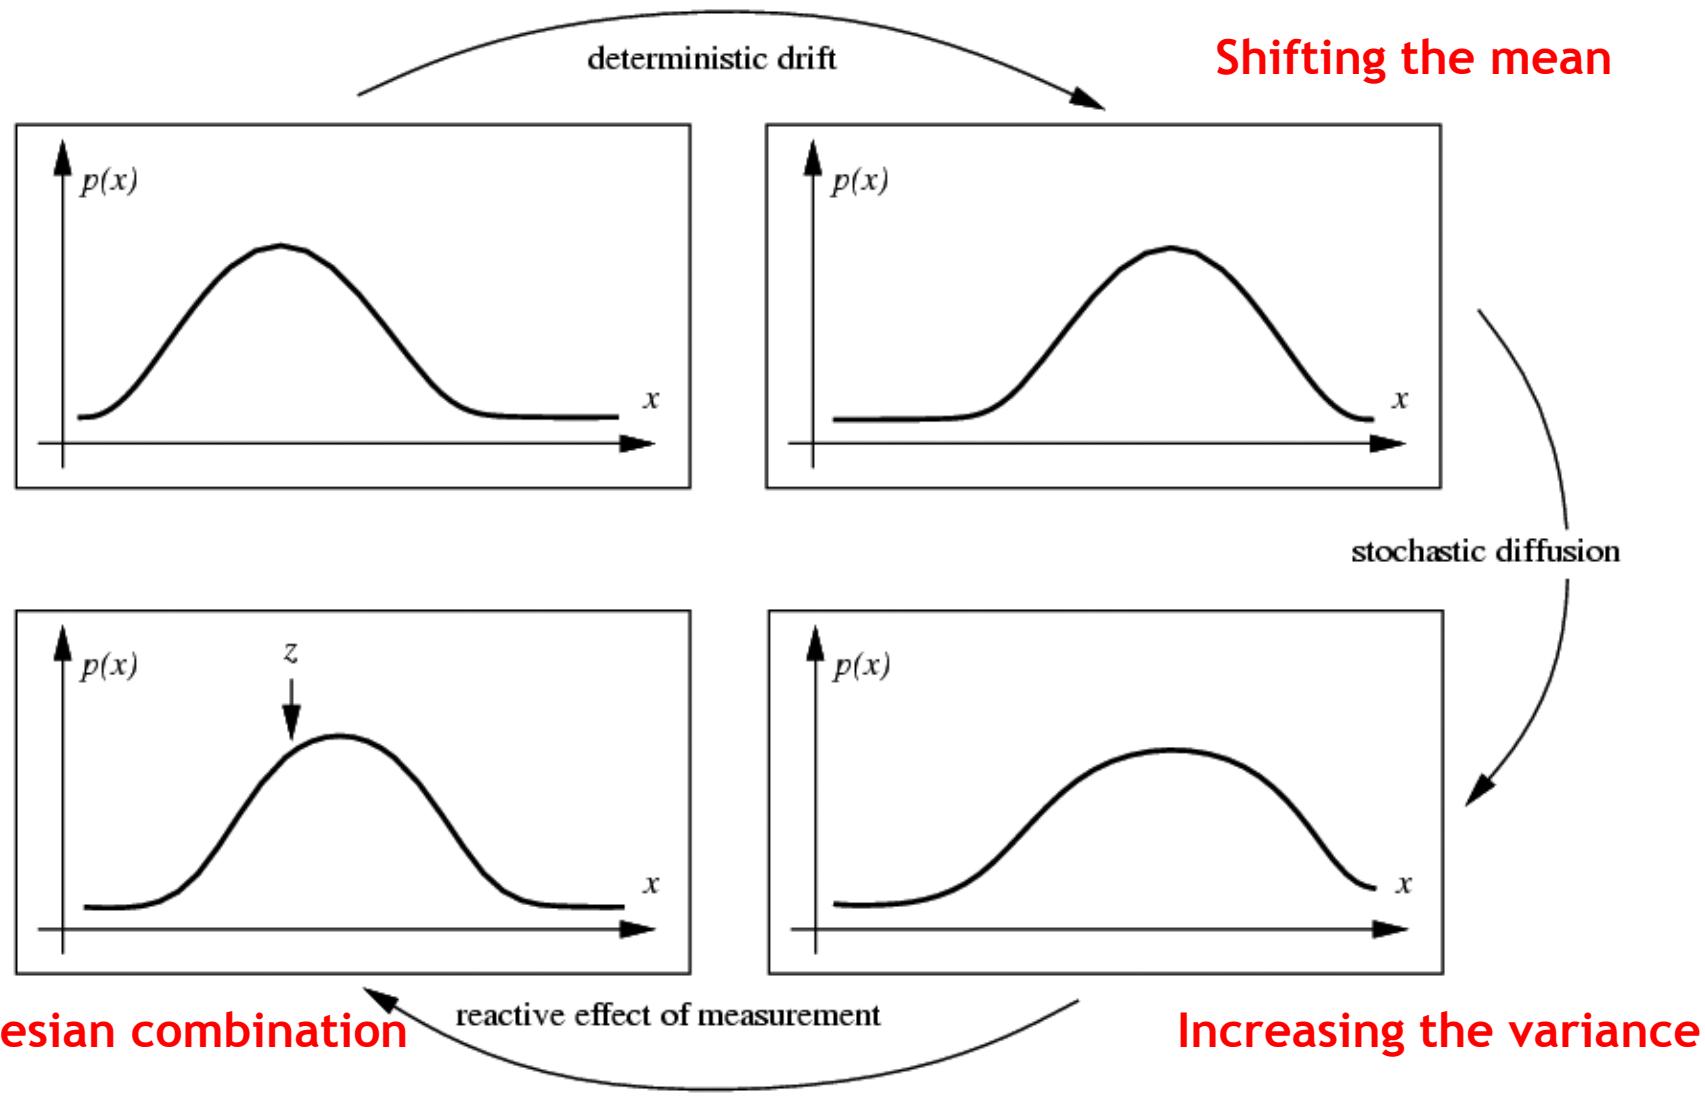
- Linear dynamical models + Gaussian noise
  - Linear transformation of Gaussian r.v.: Gaussian r.v.

$$\boldsymbol{x}_t \sim N(D_t \boldsymbol{x}_{t-1}, \Sigma_d)$$

$$y_t \sim N(M_t \boldsymbol{x}_t, \Sigma_m)$$

- The predicted/corrected state distributions are Gaussian
  - Analytical expressions for Gaussian parameters
- Applications: finance, aerospace, robotics .... vision
  - ‘Top 10 algorithms of the century’

# Propagation of Gaussian densities



# Kalman Filter for 1D State

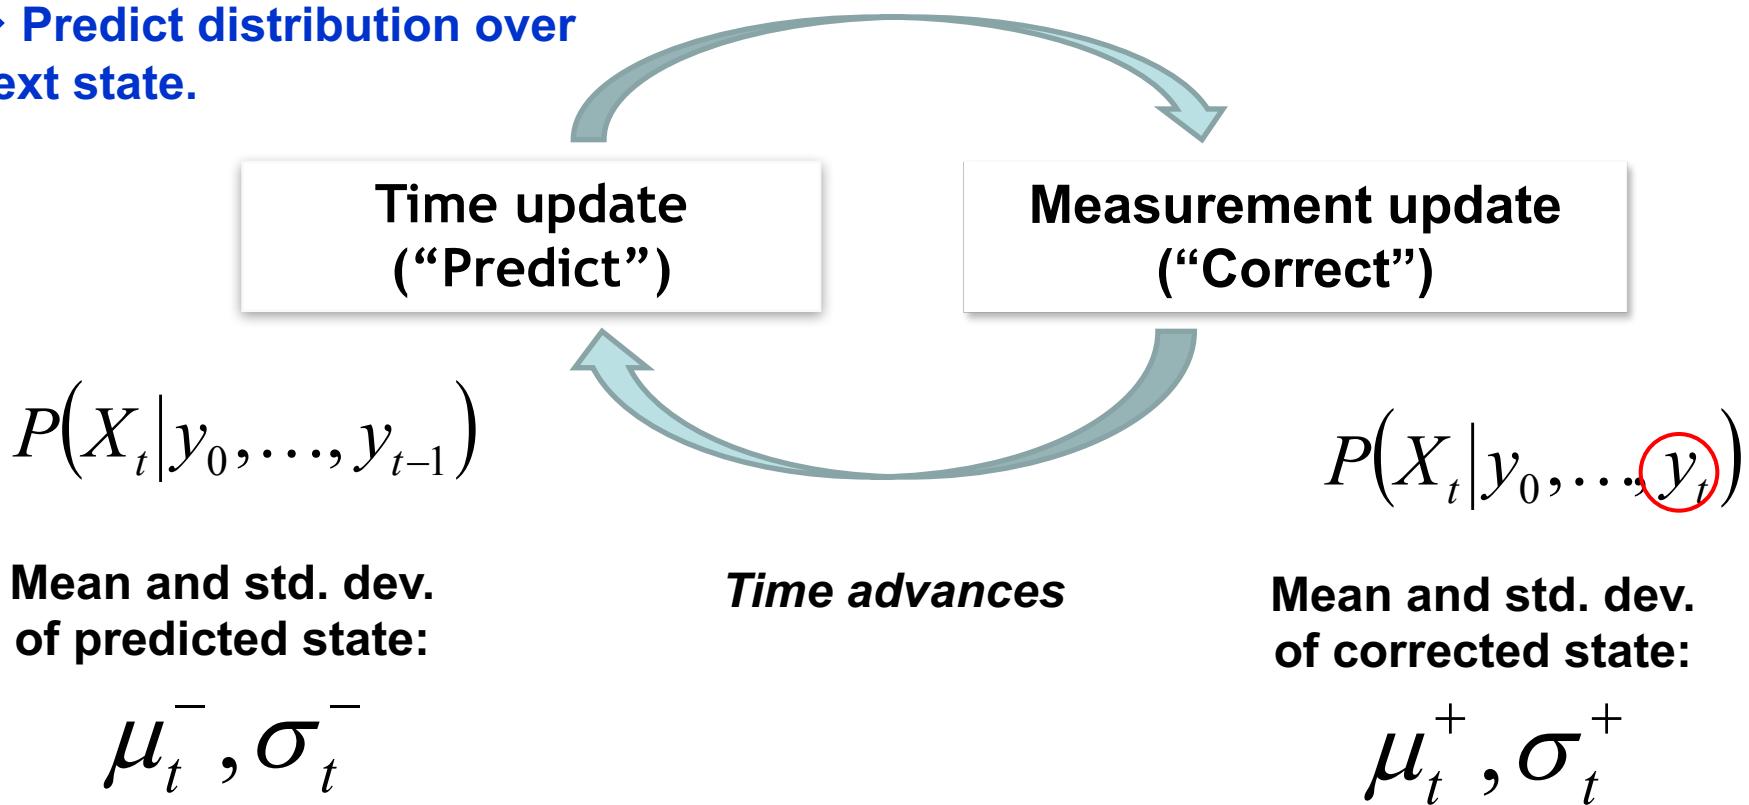
**Want to  
represent  
and update**

$$P(x_t | y_0, \dots, y_{t-1}) = N(\mu_t^-, (\sigma_t^-)^2)$$

$$P(x_t | y_0, \dots, y_t) = N(\mu_t^+, (\sigma_t^+)^2)$$

# The Kalman Filter

Know corrected state from previous time step, and all measurements up to the current one  
→ Predict distribution over next state.



# 1D Kalman Filter: Prediction

- Given:

- Linear dynamic model

$$x_t \sim N(dx_{t-1}, \sigma_d)$$

- Estimated distribution for previous state

$$P(x_{t-1} | y_1, \dots, y_{t-1}) = N(\mu_{t-1}^+, \sigma_{t-1}^+)$$

- Wanted: predicted distribution for next state

$$P(x_t | y_1, \dots, y_{t-1}) = N(\mu_t^-, \sigma_t^-)$$

- New parameters

$$\mu_t^- = d\mu_{t-1}^+ \quad (\sigma_t^-)^2 = \sigma_d^2 + (d\sigma_{t-1}^+)^2$$

# 1D Kalman Filter: Correction

- Given
  - Prediction

$$P(x_t | y_1, \dots, y_{t-1}) = N(\mu_t^-, \sigma_t^-)$$

- Measurement model

$$y_t \sim N(mx_t, \sigma_m)$$

- Wanted

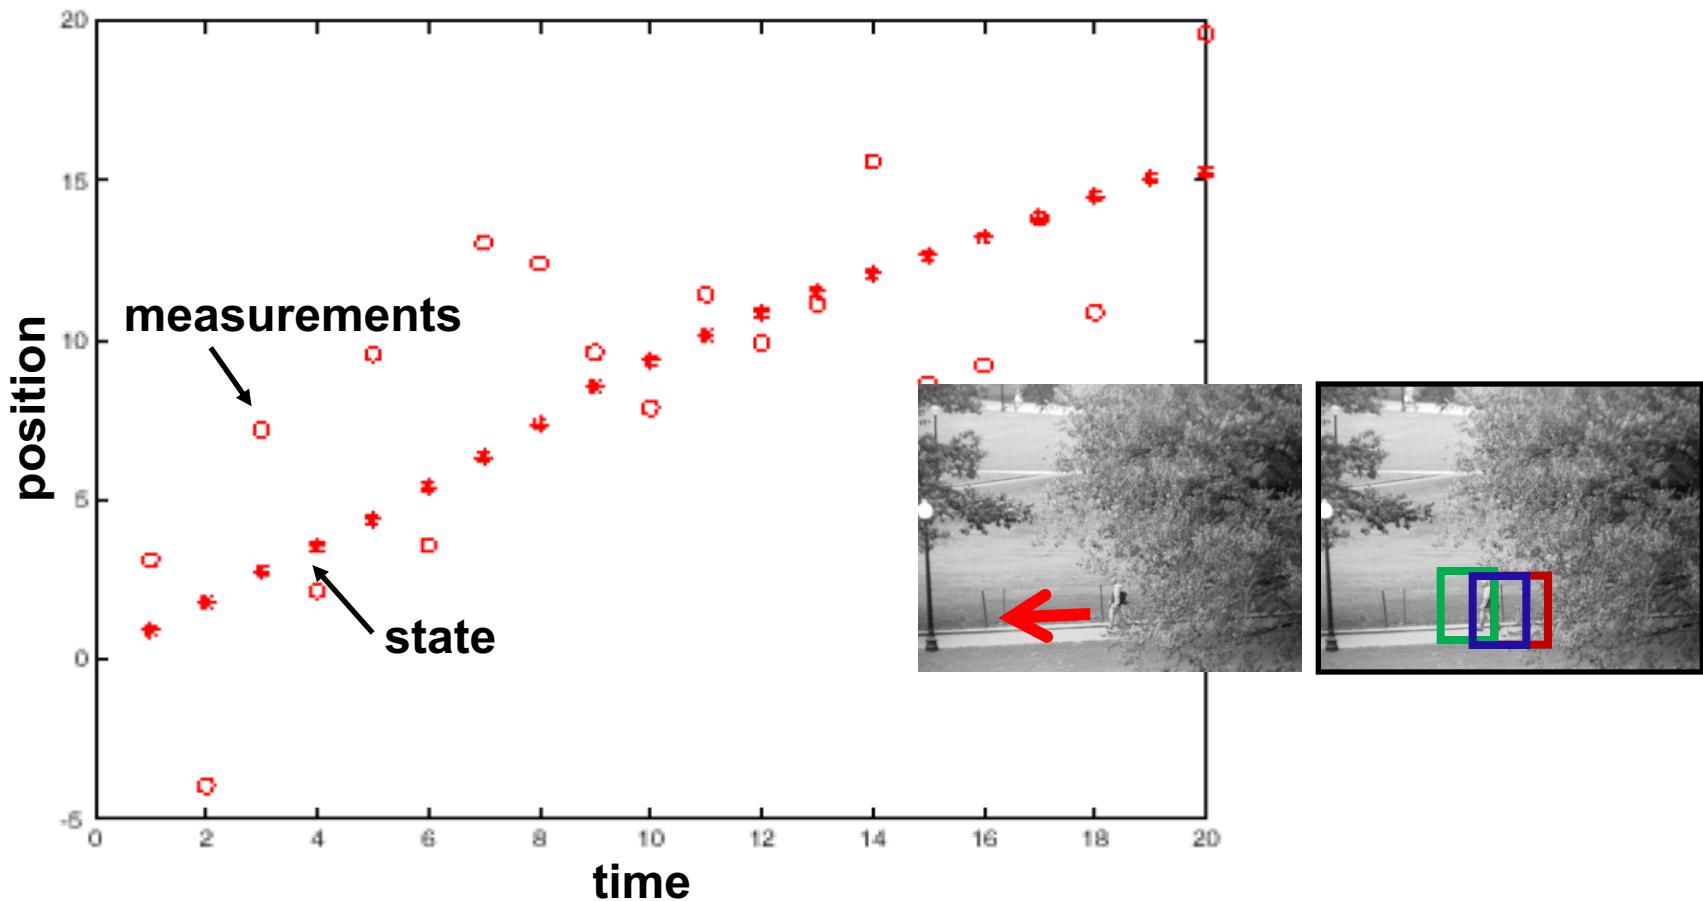
$$P(x_t | y_1, \dots, y_t) \propto P(y_t | x_t) P(x_t | y_1, \dots, y_{t-1})$$

- Gaussian

$$N(\mu_t^+, \sigma_t^+)$$

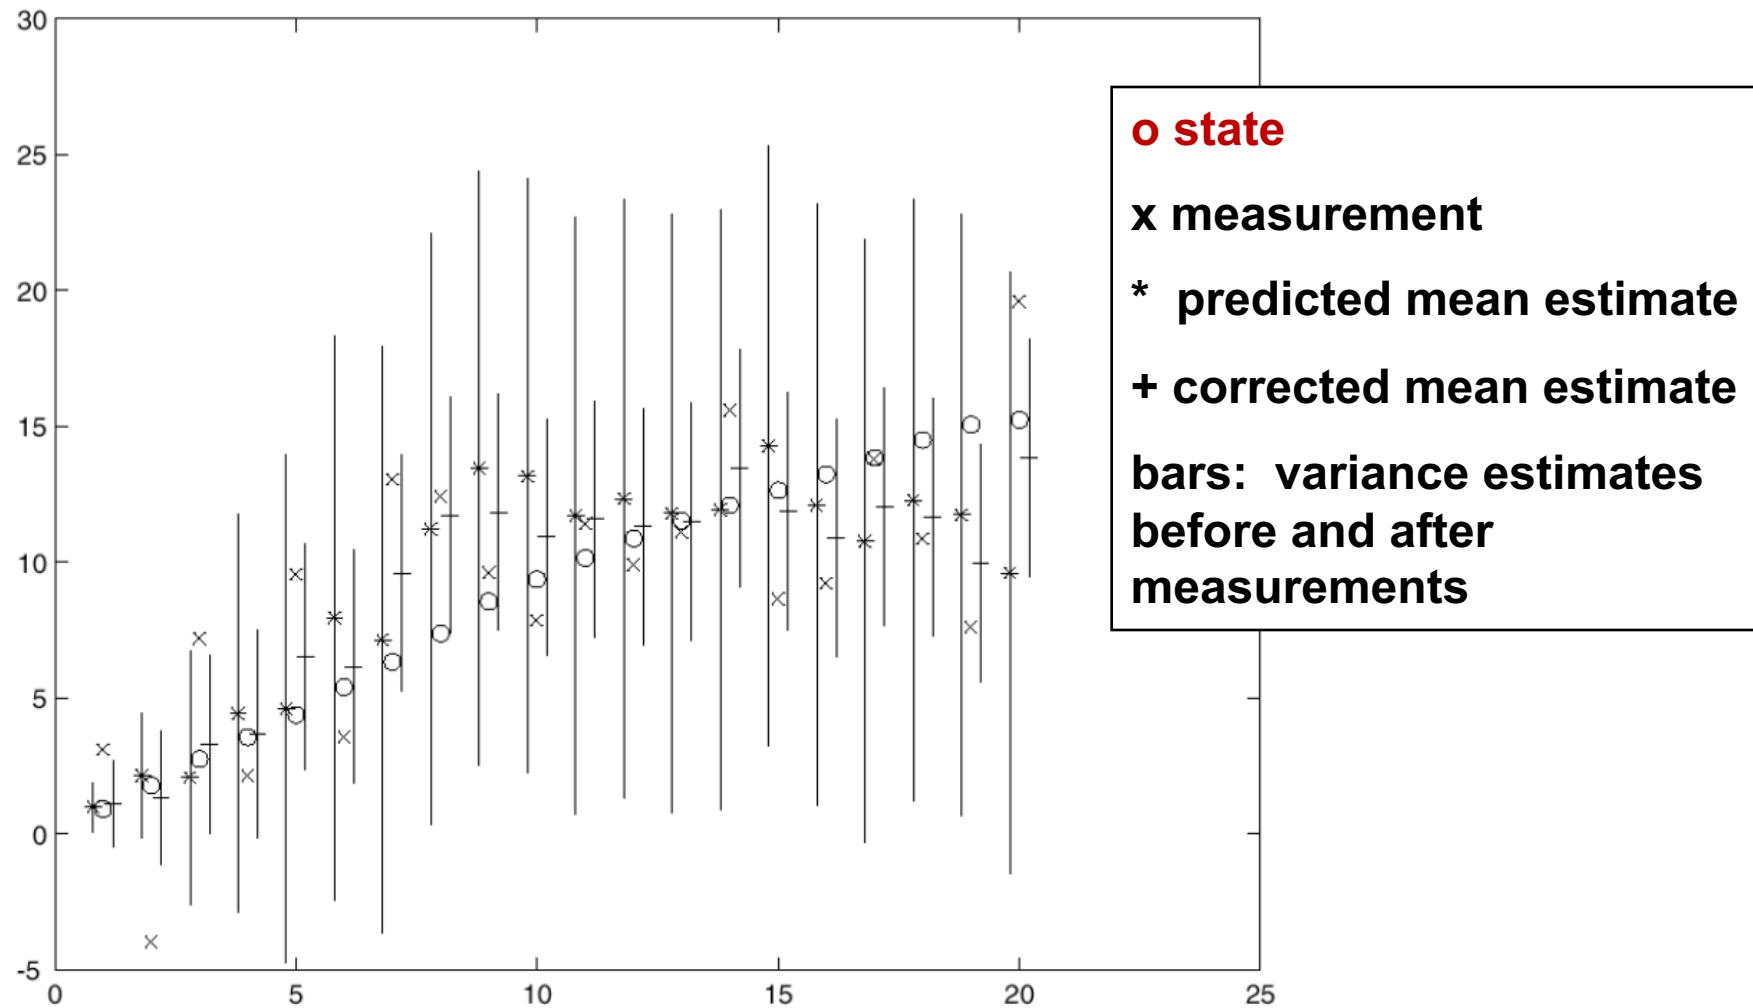
$$\mu_t^+ = \frac{\mu_t^- \sigma_m^2 + my_t (\sigma_t^-)^2}{\sigma_m^2 + (m\sigma_t^-)^2} \quad \frac{1}{\sigma_t^+} = \frac{1}{(\sigma_t^-)^2} + \frac{1}{\sigma_m^2}$$

# Recall: Constant Velocity Example

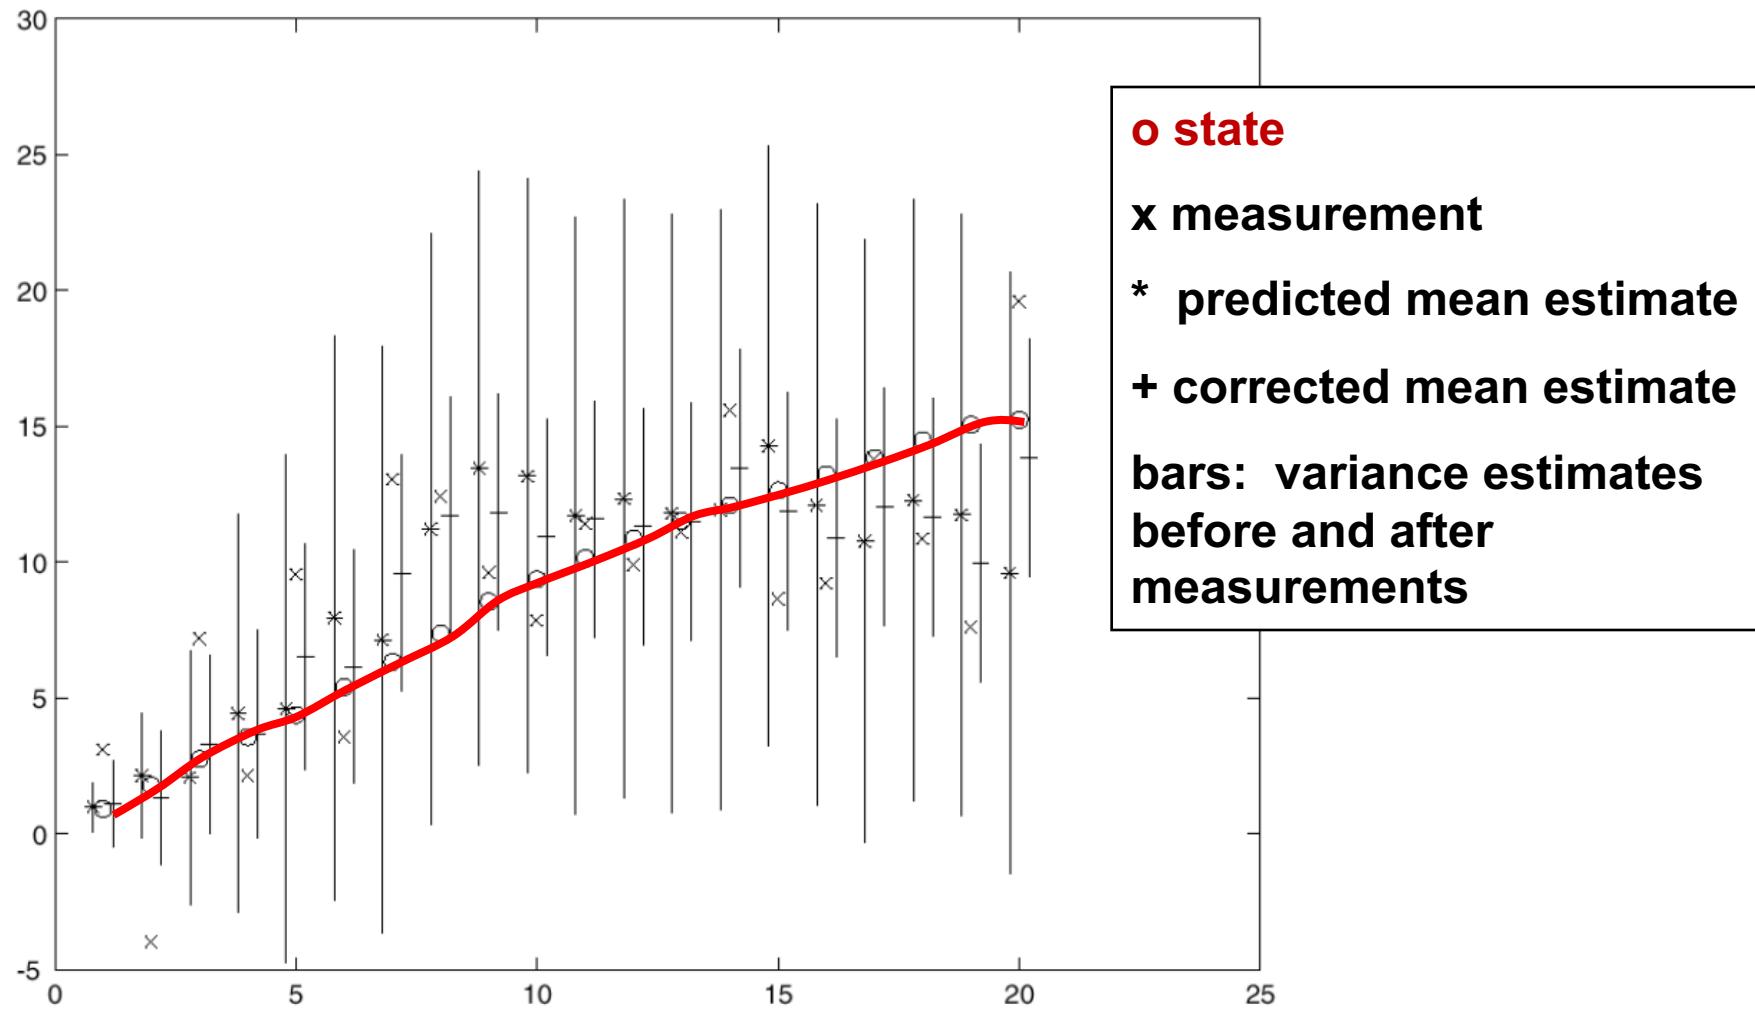


State is 2D: position + velocity  
Measurement is 1D: position

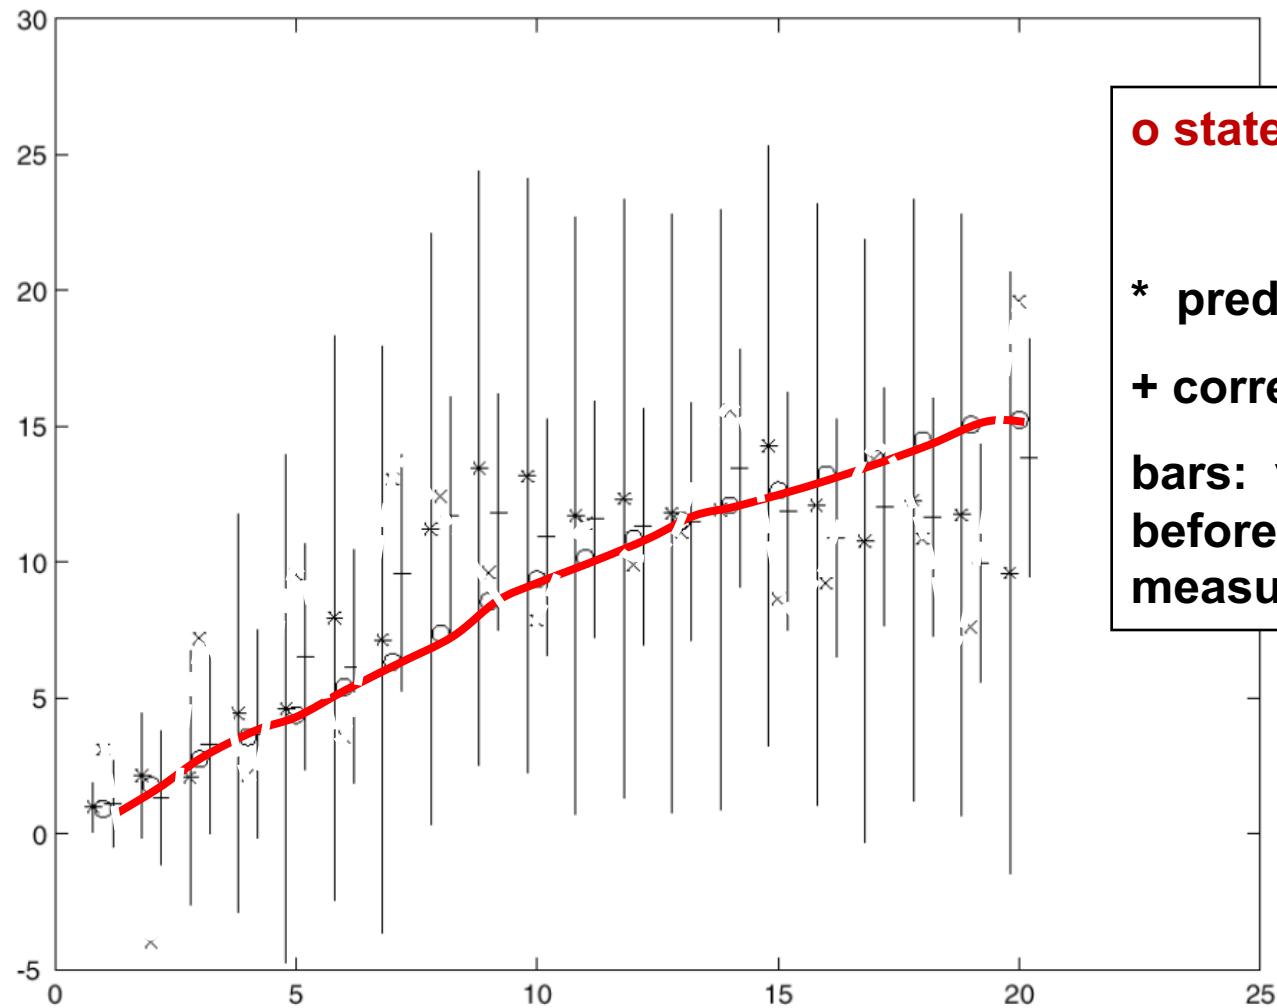
# Constant Velocity Model



# Constant Velocity Model



# Constant Velocity Model



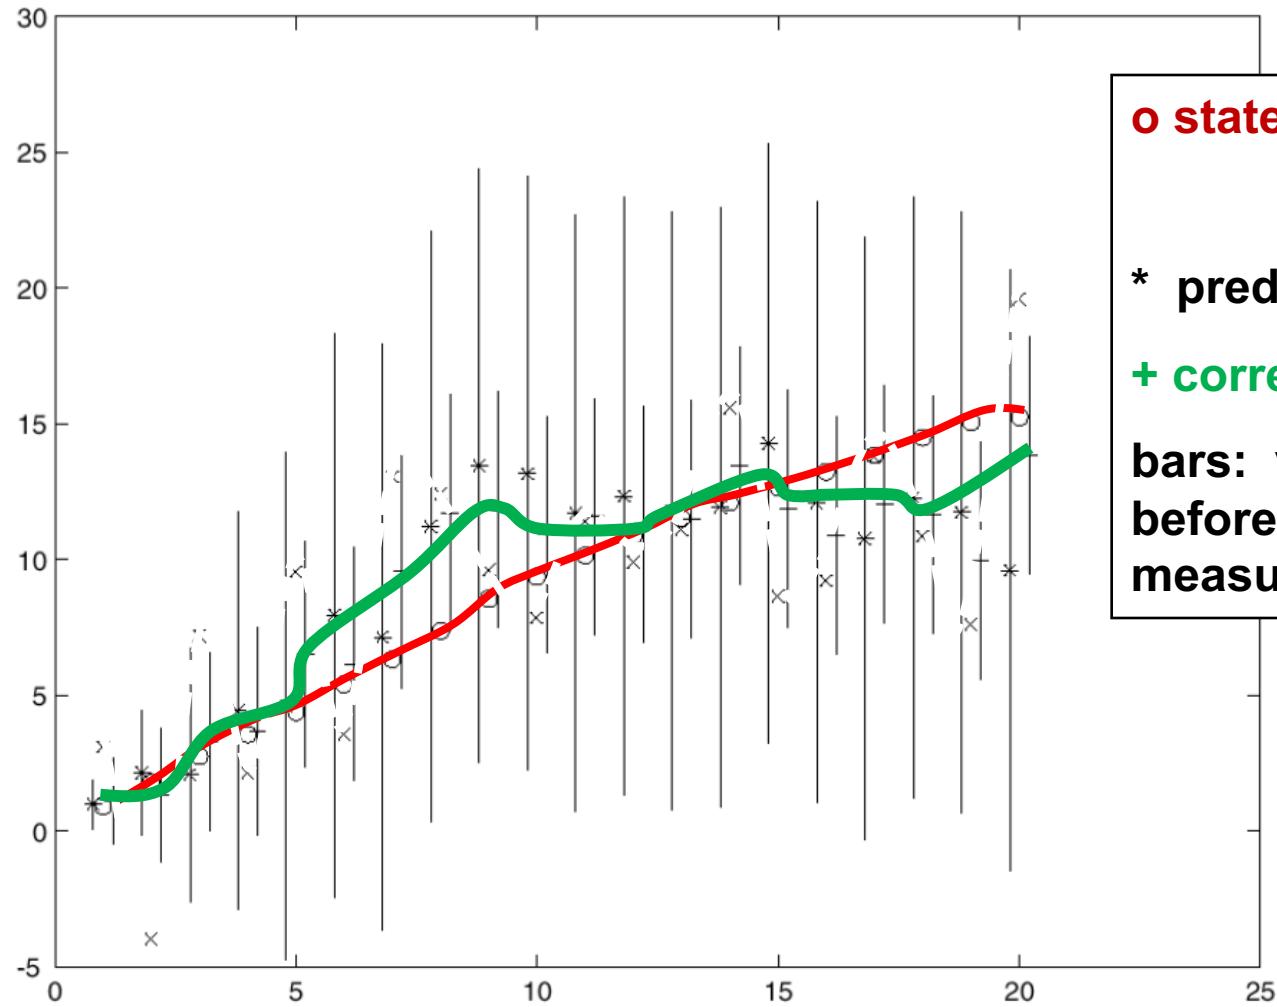
**o state**

\* predicted mean estimate

+ corrected mean estimate

bars: variance estimates  
before and after  
measurements

# Constant Velocity Model



o state

\* predicted mean estimate

+ corrected mean estimate

bars: variance estimates  
before and after  
measurements

# Kalman Filter: General Case (>1dim)

- What if state vectors have more than one dimension?

PREDICT

$$x_t^- = D_t x_{t-1}^+$$

$$\Sigma_t^- = D_t \Sigma_{t-1}^+ D_t^T + \Sigma_{d_t}$$

CORRECT

$$K_t = \Sigma_t^- M_t^T \left( M_t \Sigma_t^- M_t^T + \Sigma_{m_t} \right)^{-1}$$

$$x_t^+ = x_t^- + K_t \left( y_t - M_t x_t^- \right) \text{ "residual"}$$

$$\Sigma_t^+ = (I - K_t M_t) \Sigma_t^-$$

More weight on residual  
when measurement error  
covariance approaches 0.

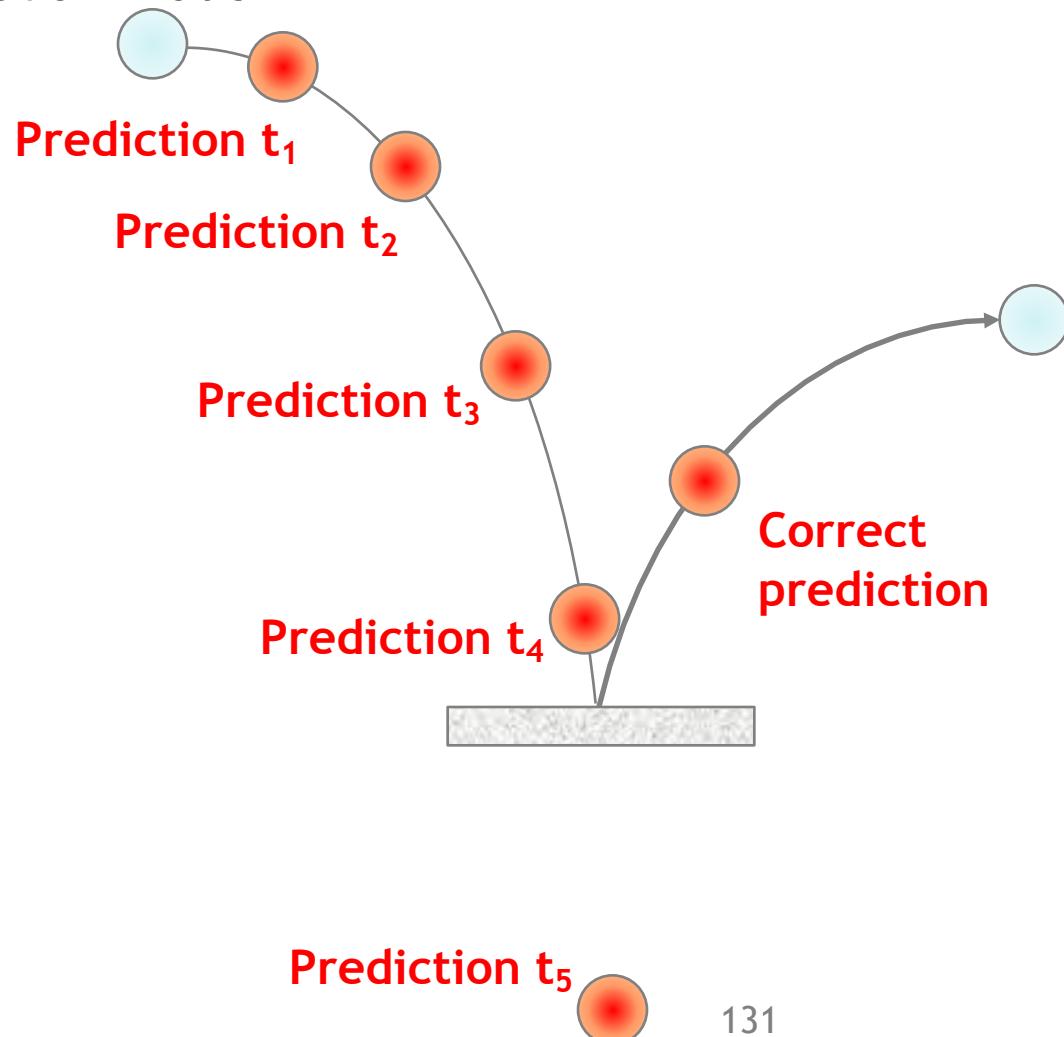
Less weight on residual as  
a priori estimate error  
covariance approaches 0.

# This lecture

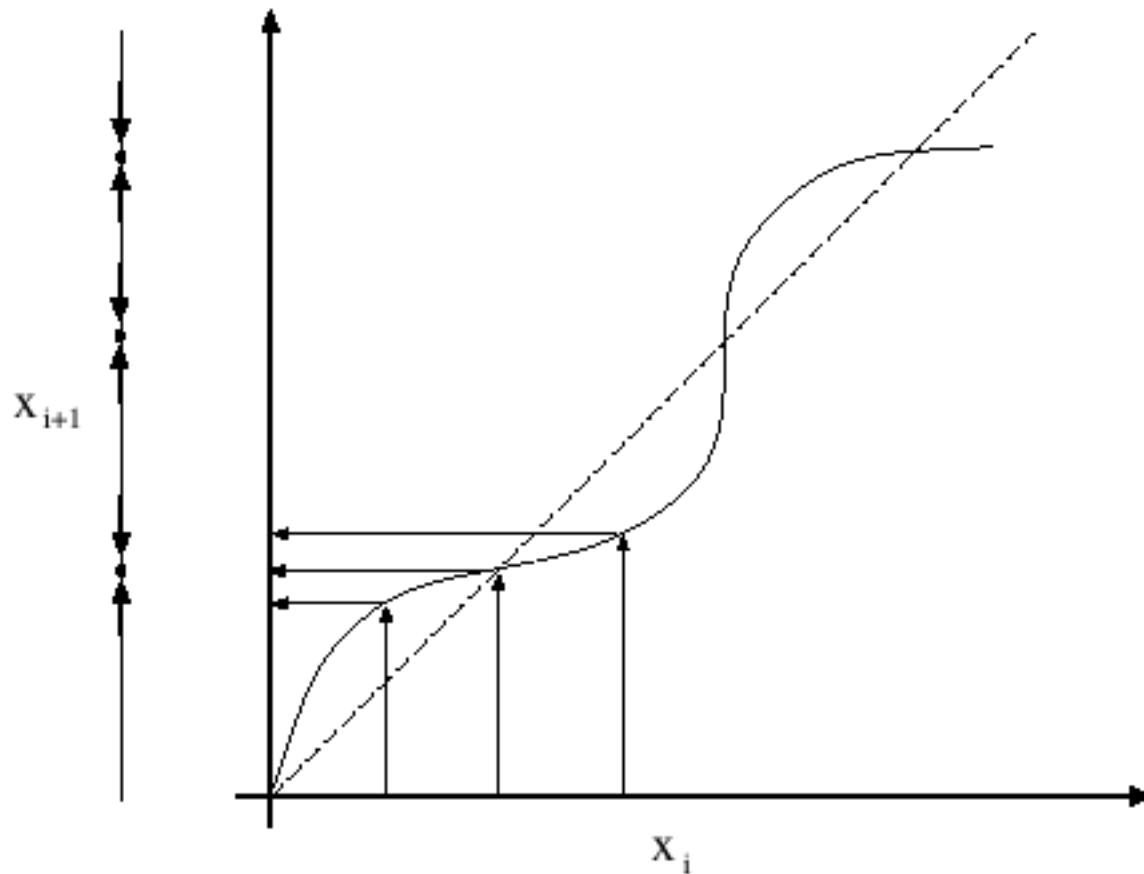
- Tracking
  - Nonlinear models

# Ball Example: What Goes Wrong Here?

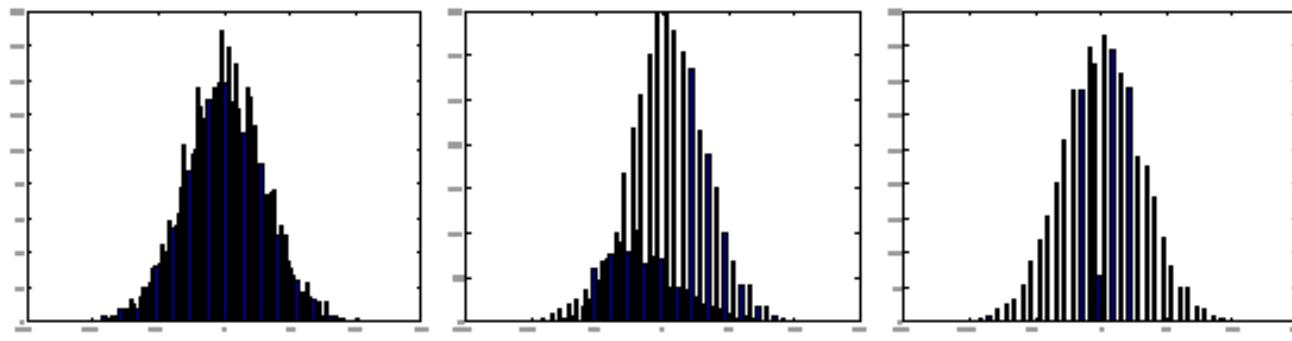
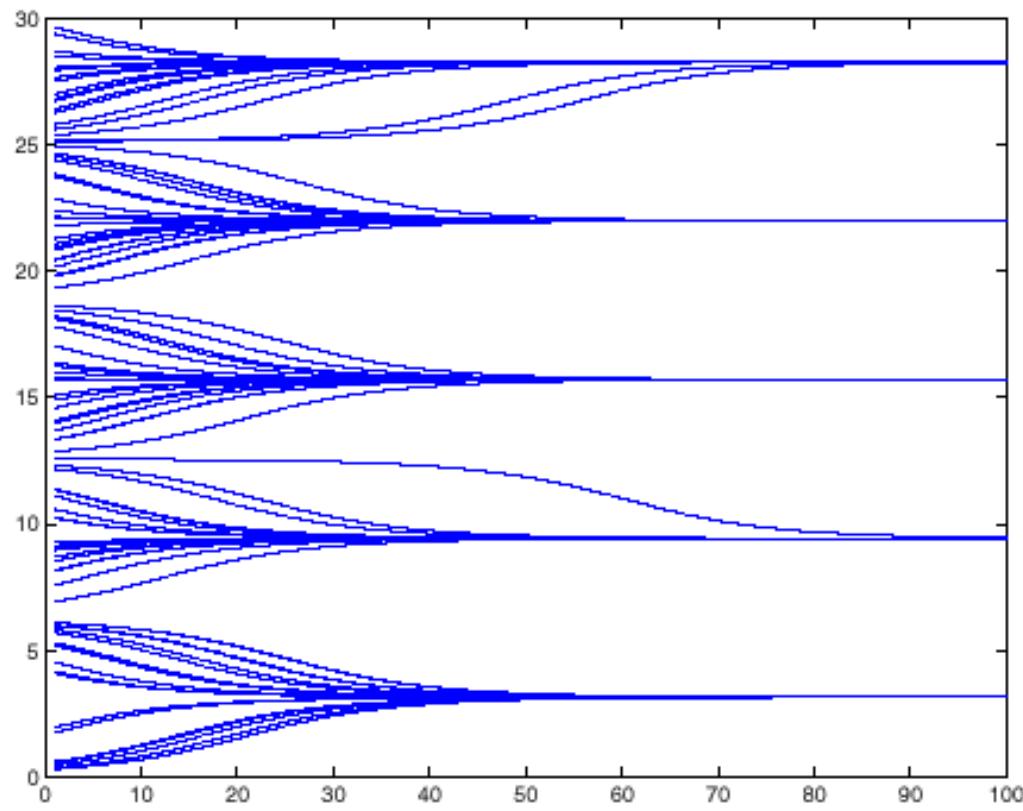
- Assuming constant acceleration model



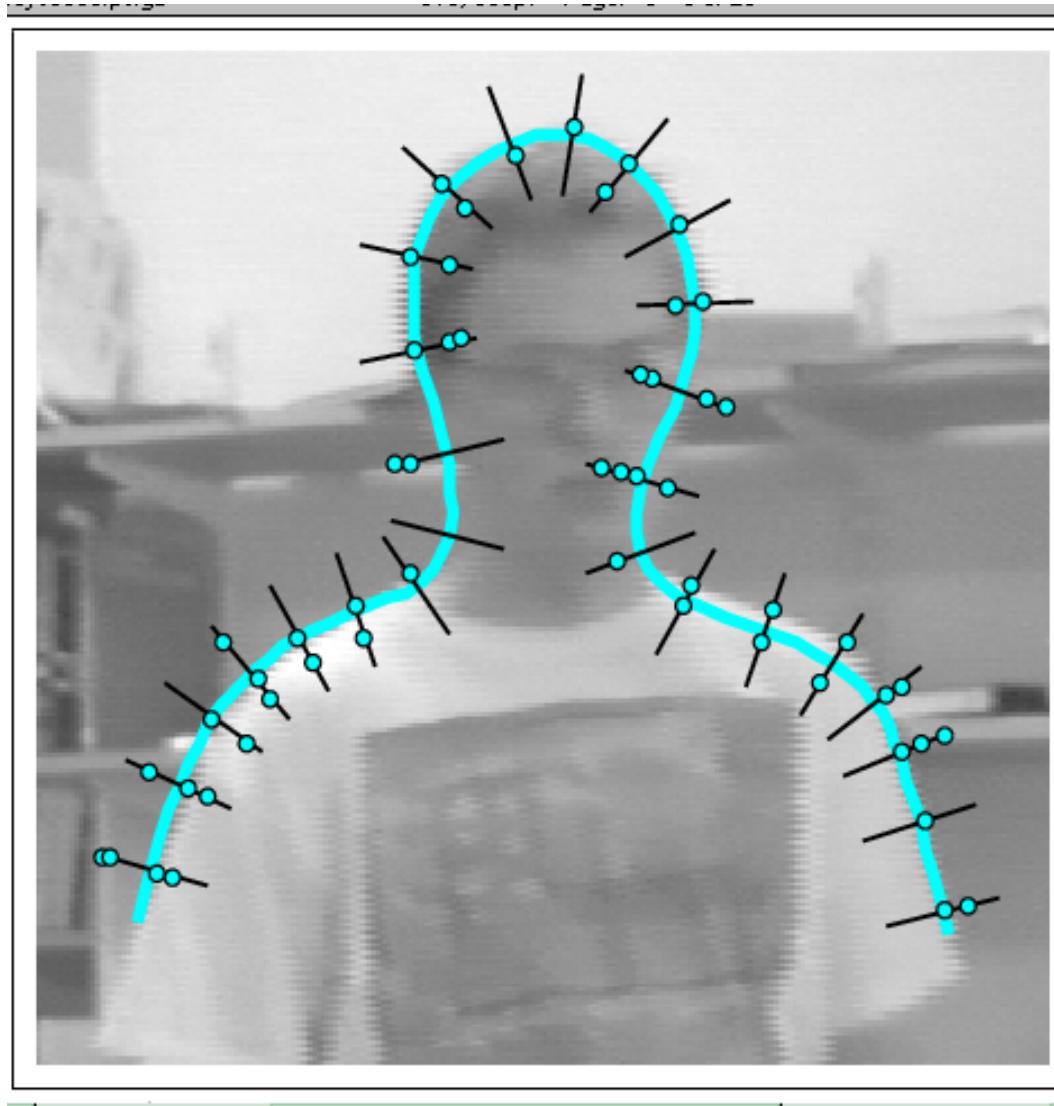
- Prediction is too far from true position to compensate...



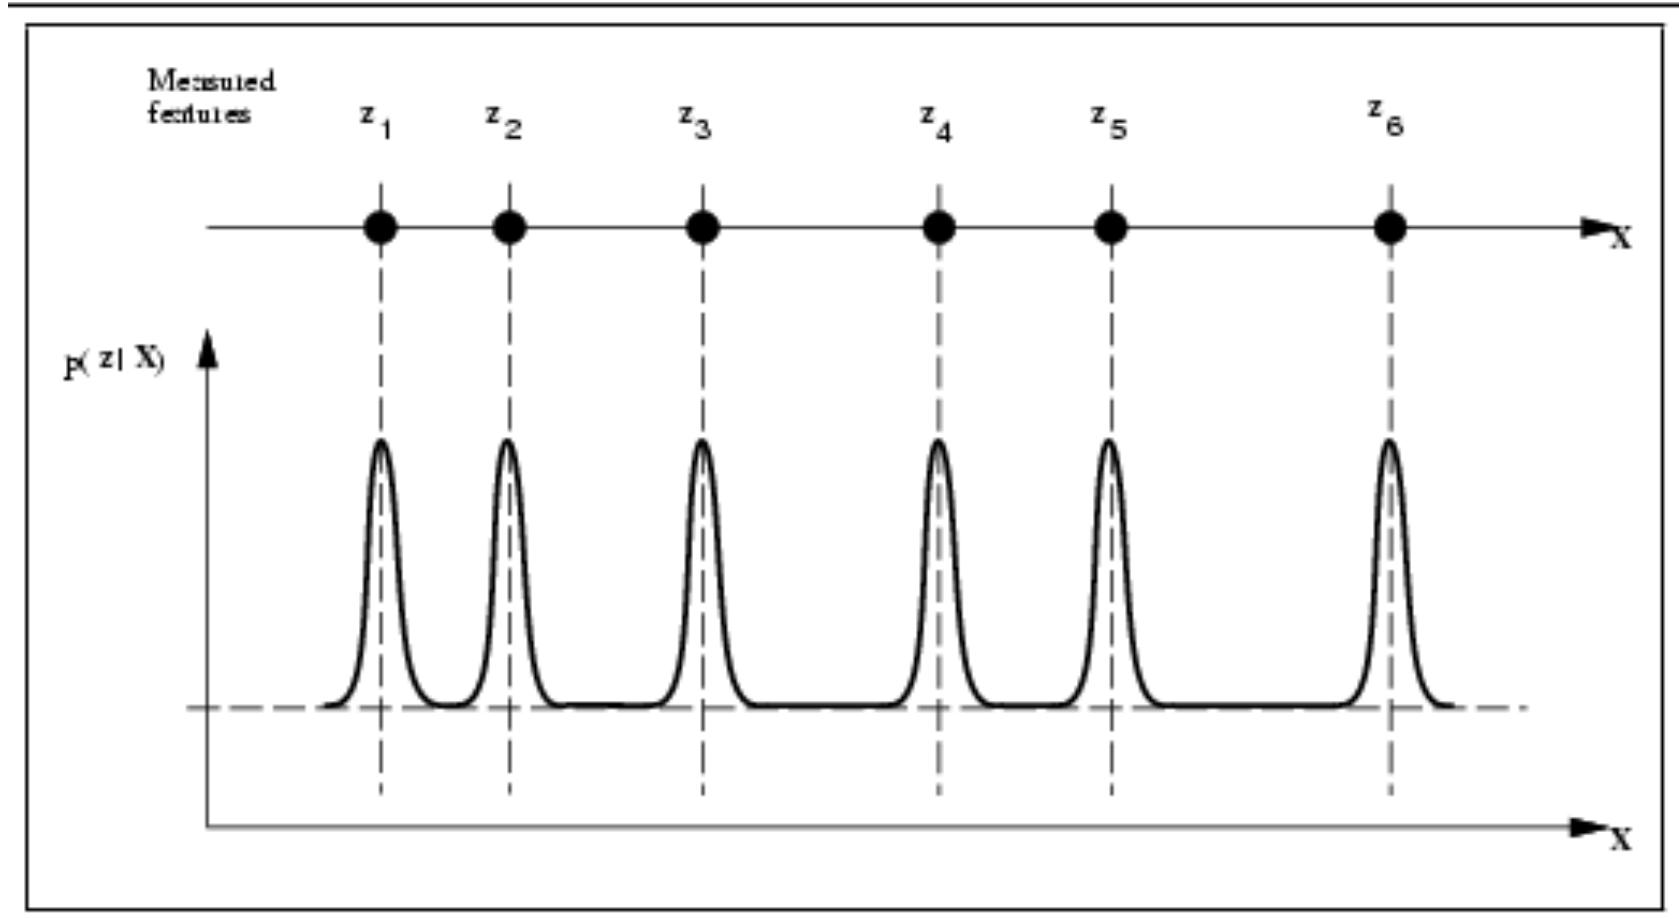
$$x_{i+1} = x_i + \epsilon \sin(x_i)$$



# Nonlinear Observation Likelihoods

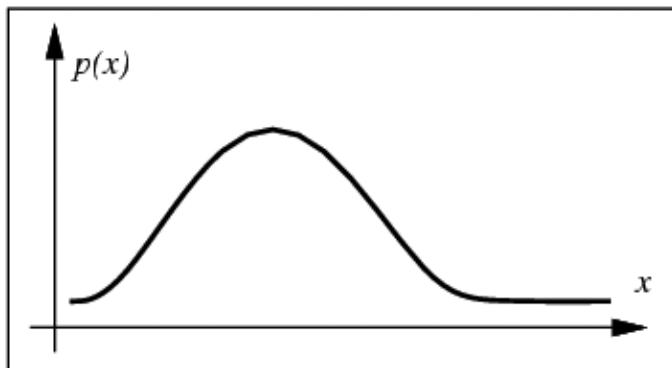


# Visual clutter → observational nonlinearity

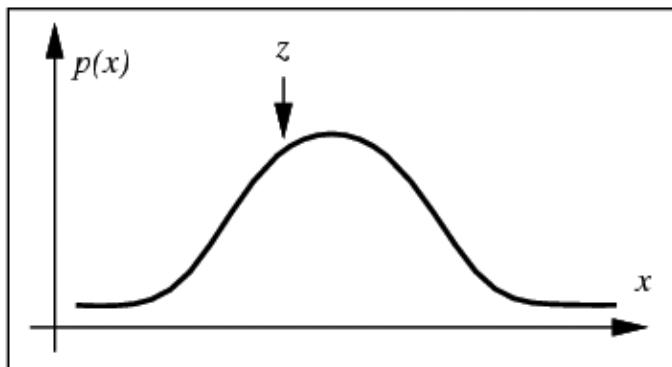
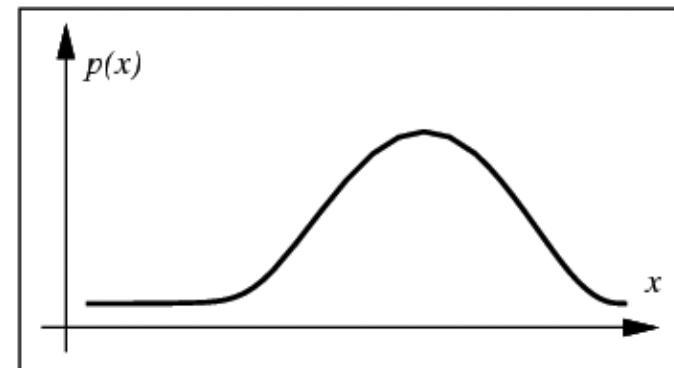


# Propagation of Gaussian densities

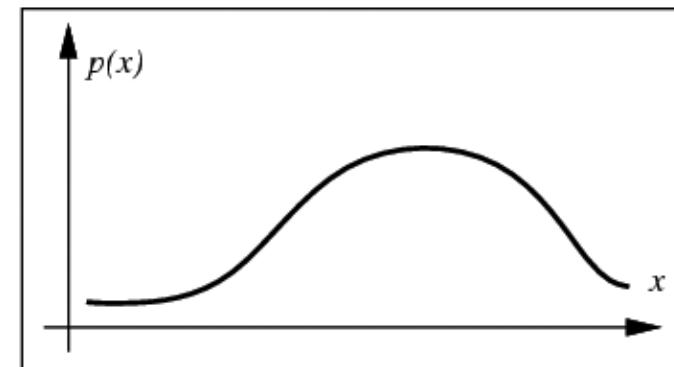
Valid only for linear dynamics



deterministic drift

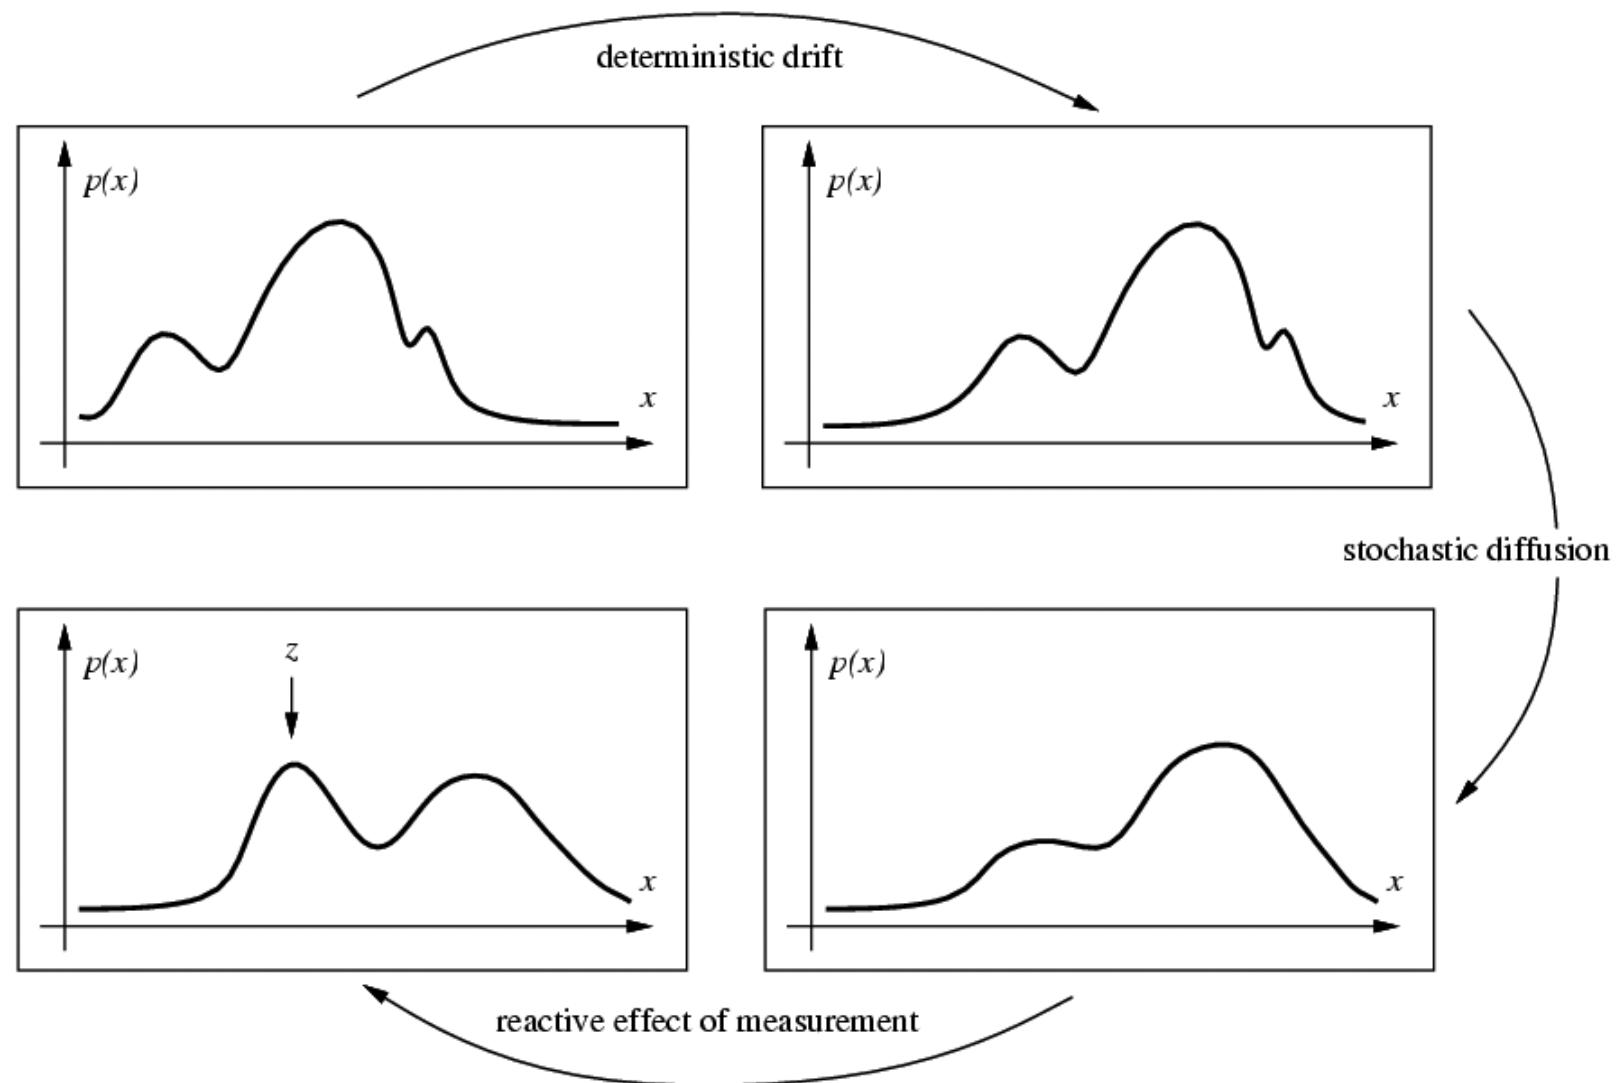


stochastic diffusion



reactive effect of measurement

# Propagation of General Densities



# Dealing with multi-modal distributions

- Mixture modelling?
  - How many mixtures?
  - Iterative fitting (EM)
- Nonparametric density estimate

$$\hat{f}(x) = \frac{1}{N} \sum_{i=1}^N K(|x - x^i|)$$

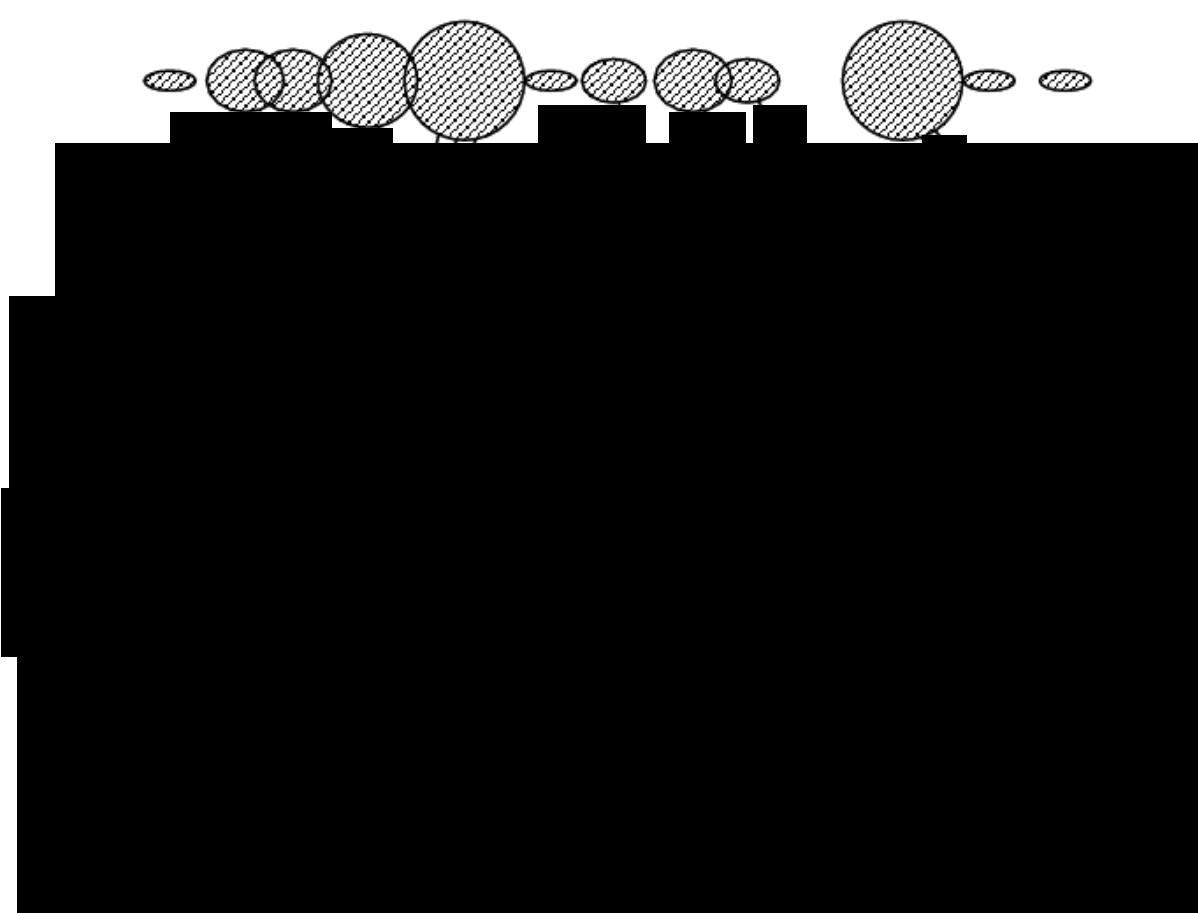


e. g.  $K(|x - x^i|) = \frac{1}{(\sqrt{2\pi}\sigma)^D} \exp\left(-\frac{|x - x^i|^2}{2\sigma^2}\right)$

# Particle Filtering

- Use sampling to propagate densities over time (i.e., across frames in a video sequence).
- At each time step, represent posterior  $P(X_t|Y_t)$  with weighted sample set.
- Previous time step's sample set  $P(X_{t-1}|Y_{t-1})$  is passed to next time step as the effective prior.

# Particle Filtering



Start with weighted samples from previous time step

Sample and shift according to dynamics model

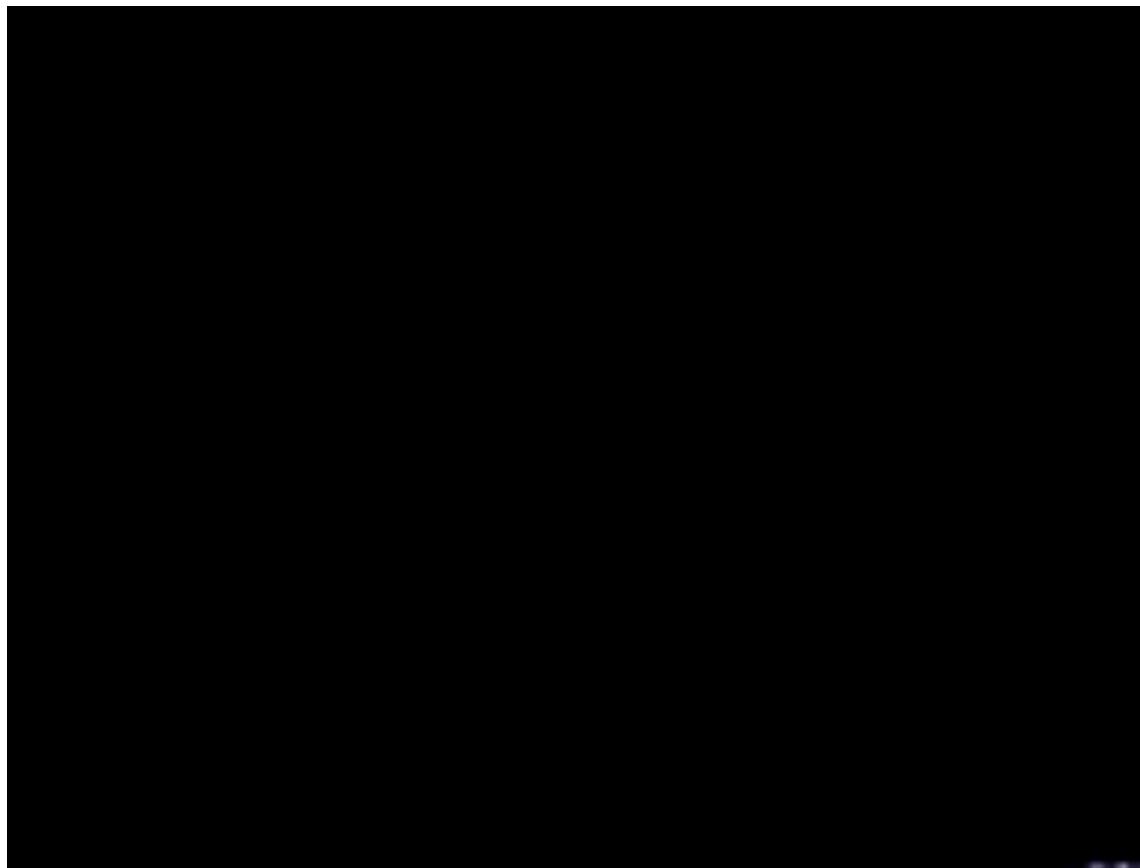
Spread due to randomness; this is predicted density  $P(X_t|Y_{t-1})$

Weight the samples according to observation density

Arrive at corrected density estimate  $P(X_t|Y_t)$

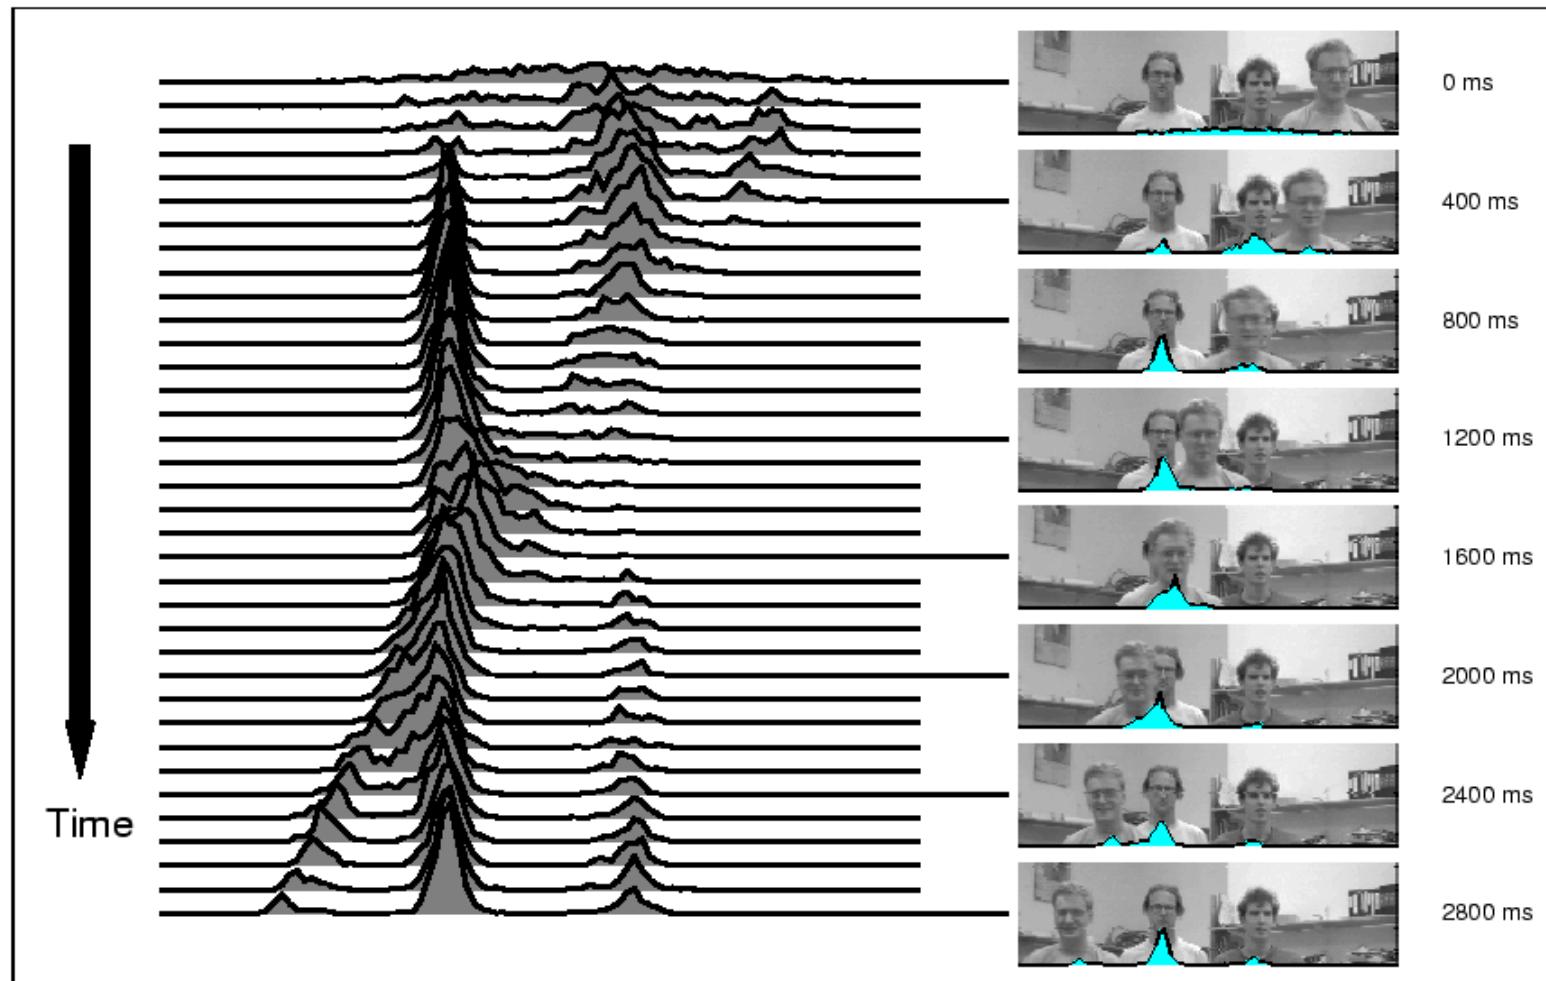
M. Isard and A. Blake, [CONDENSATION -- conditional density propagation for visual tracking](#), IJCV 29(1):5-28, 1998

# Particle Filtering – Visualization



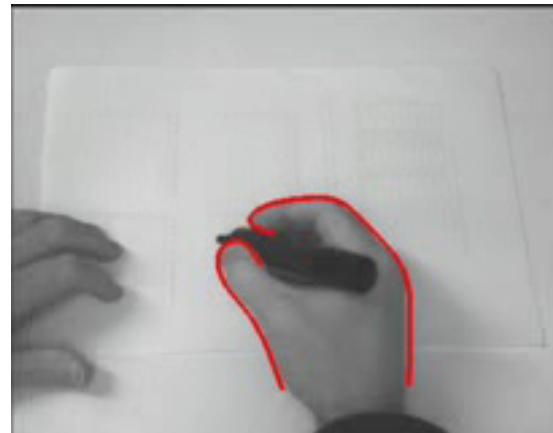
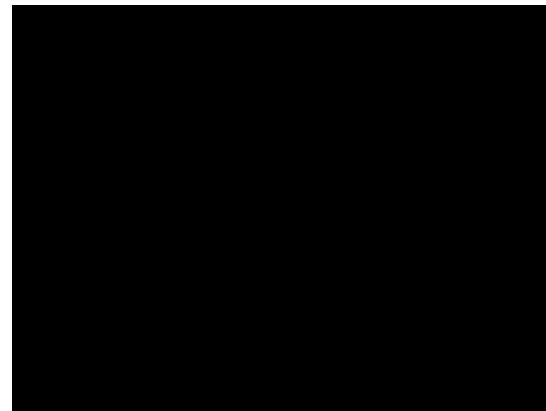
Code and video available from  
<http://www.robots.ox.ac.uk/~misard/condensation.html>

# Particle Filtering Results



<http://www.robots.ox.ac.uk/~misard/condensation.html>

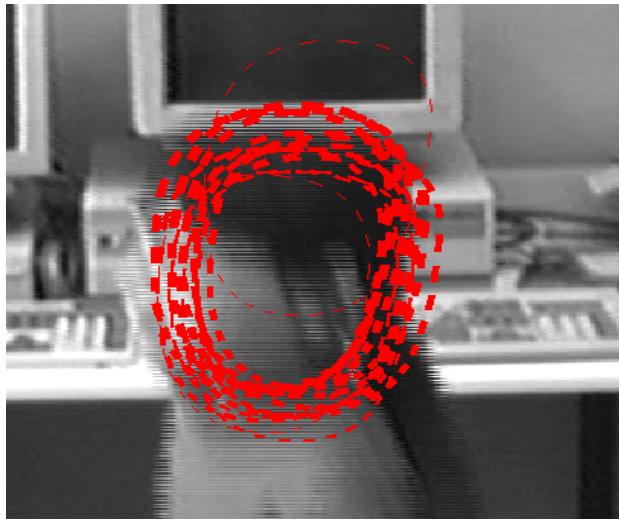
# Particle Filtering Results



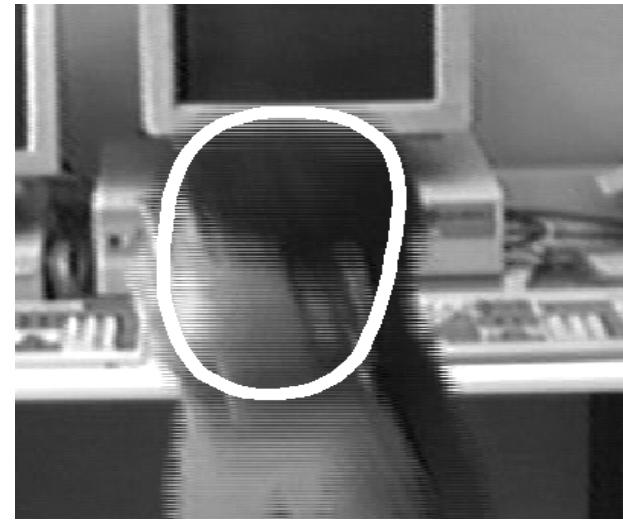
<http://www.robots.ox.ac.uk/~misard/condensation.html>

# Obtaining a State Estimate

- Note that there's no explicit state estimate maintained—just a “cloud” of particles
- Can obtain an estimate at a particular time by querying the current particle set, e.g. by computing the average



**State samples**  
(thickness proportional to weight)



**Mean of weighted  
state samples** 144

From Isard & Blake, 1998

# Extensions

- Inference on Markov Random Fields
  - Main problem: high-dimensional state space

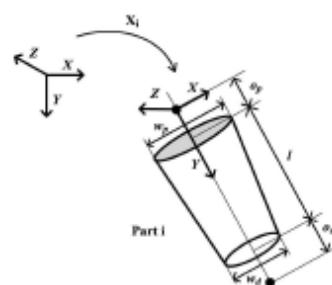
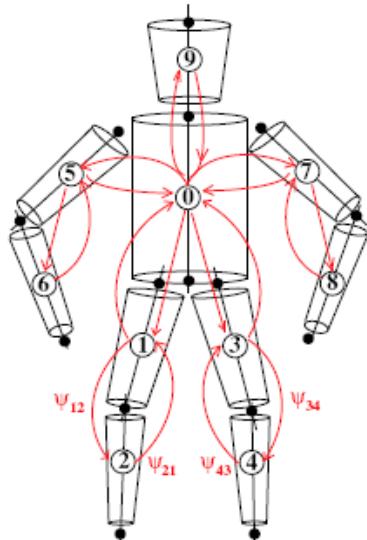


Figure 2: Parameterization of part  $i$ .

$$m_{i,j}(c_j) = \sum_{c_i} \Phi_i(c_i) \Psi_{i,j}(c_i, c_j) \prod_{k \in \{\mathcal{N}(i) \setminus j\}} m_{k,i}(c_i)$$

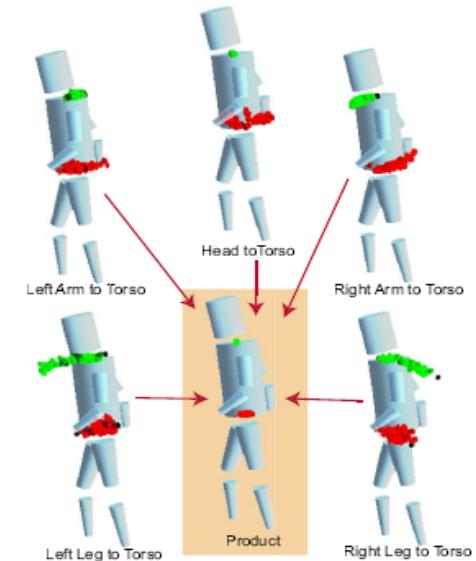


Figure 4: Message Product. The head, upper arms, and upper legs send messages to the torso. Samples from these messages are illustrated by showing the predicted torso location with green balls. The distribution over the orientation of the torso is illustrated by showing a red ball at the distal end of the torso for each sample. While any single message represents uncertain information about the torso pose, the product of these messages tightly constrains the torso position and orientation.

# Applications

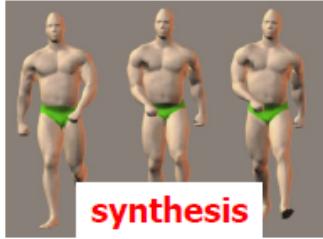
- Human Motion Analysis



surveillance



sport analysis



synthesis

Human  
Motion  
Analysis



biomechanics



motion  
capture



game interfaces



sign language

# Current state-of-the-art

- 'Tracking People by Learning their Appearance' Ramanan, Forsyth & Zisserman, PAMI 2007
  - Tracking-by-detection



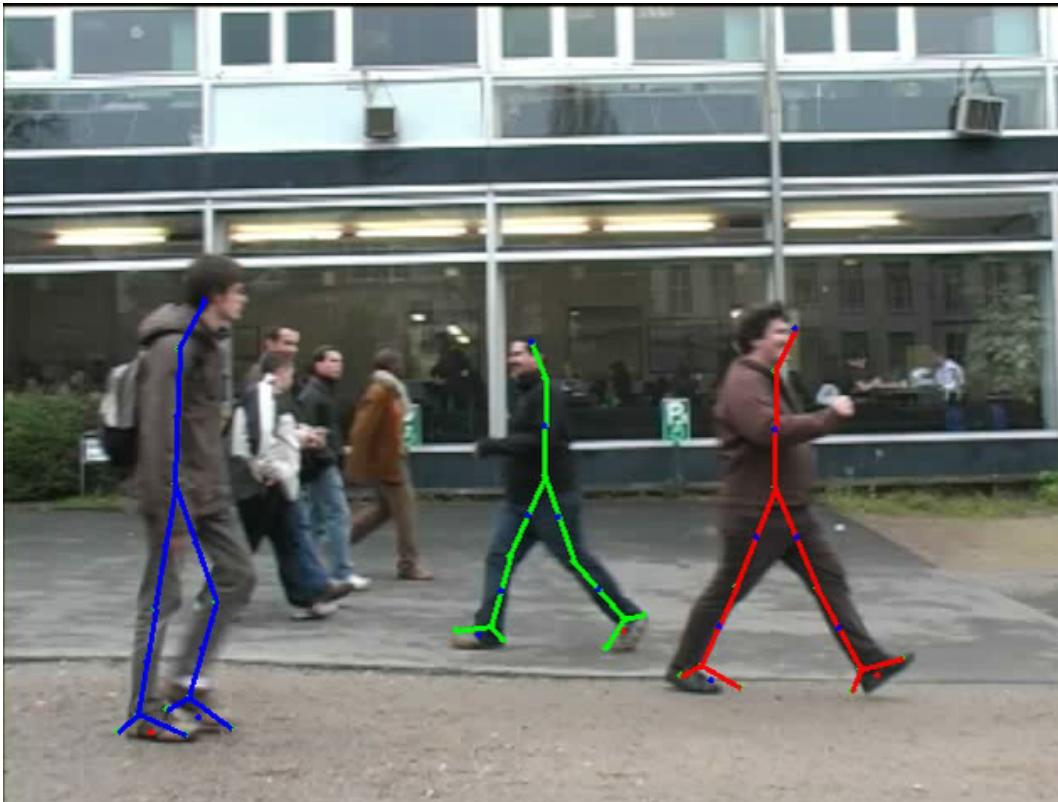
# Current state-of-the-art

- 'Hierarchical Models for 3D Visual Inference', CVPR 2007, Atul Kanaujia, Cristian Sminchisescu and Dimitris Metaxas
  - 3D pose estimation



# Current state-of-the-art

- Multiple people tracking & pose estimation
  - ‘People tracking by detection and people detection by tracking’, CVPR 2008, M. Adriluka, S. Roth and B. Schiele



# Current state-of-the-art

- Multiple people tracking & pose estimation

